

Computer Architecture: MIPS Pipeline Datapath

Hossein Asadi (asadi@sharif.edu)

Department of Computer Engineering

Sharif University of Technology

Spring 2024



Copyright Notice

- Some Parts (text & figures) of this Lecture adopted from following:
 - D.A. Patterson and J.L. Hennessy, “[Computer Organization and Design: the Hardware/Software Interface](#)” (MIPS), 6th Edition, 2020.
 - J.L. Hennessy and D.A. Patterson, “[Computer Architecture: A Quantitative Approach](#)”, 6th Edition, Nov. 2017.
 - “Intro to Computer Architecture” handouts, by Prof. Hoe, CMU, Spring 2009.
 - “Computer Architecture & Engineering” handouts, by Prof. Kubiawicz, UC Berkeley, Spring 2004.
 - “Intro to Computer Architecture” handouts, by Prof. Hoe, UWisc, Spring 2021.
 - “Computer Arch I” handouts, by Prof. Garzarán, UIUC, Spring 2009.




Quick Reminder from Previous Lecture



Single-Cycle CPU Clock Cycle Time

	I-cache	Decode, R-Read	ALU	PC update	D-cache	R-Write	Total
R-type	1	1	.9	-	-	.8	3.7
Load	1	1	.9	-	1	.8	4.7
Store	1	1	.9	-	1	-	3.9
beq	1	1	.9	.1	-	-	3.0



Clock cycle time
= 4.7 + setup + hold

Load on critical path

Setup time?

Hold time?

Critical path?



Multicycle Implementation

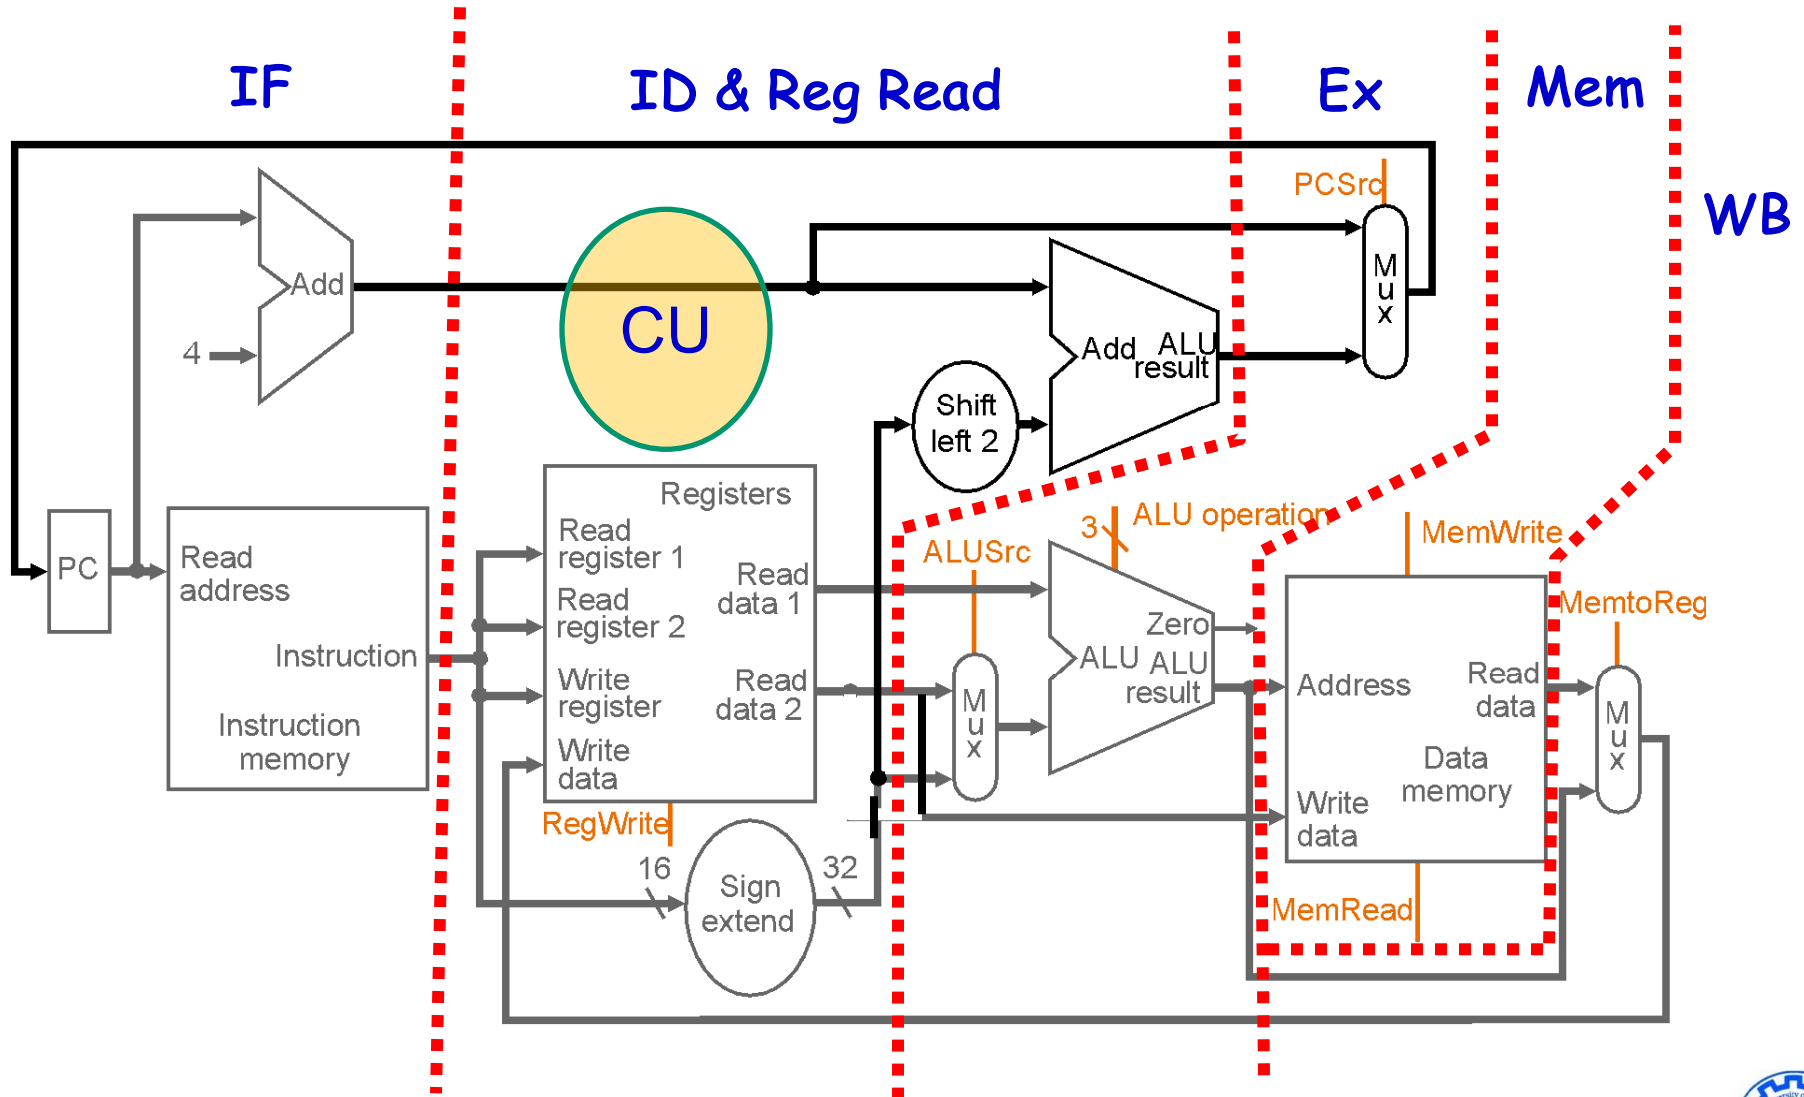
Goal: Balance amount of work done each cycle

	I cache	Decode, R-Read	ALU	PC update	D cache	R- Write	Total
R-type	1	1	.9	-	-	.8	3.7
Load	1	1	.9	-	1	.8	4.7
Store	1	1	.9	-	1	-	3.9
beq	1	1	.9	.1	-	-	3.0


- Load needs 5 cycles
- Store and R-type need 4
- beq needs 3



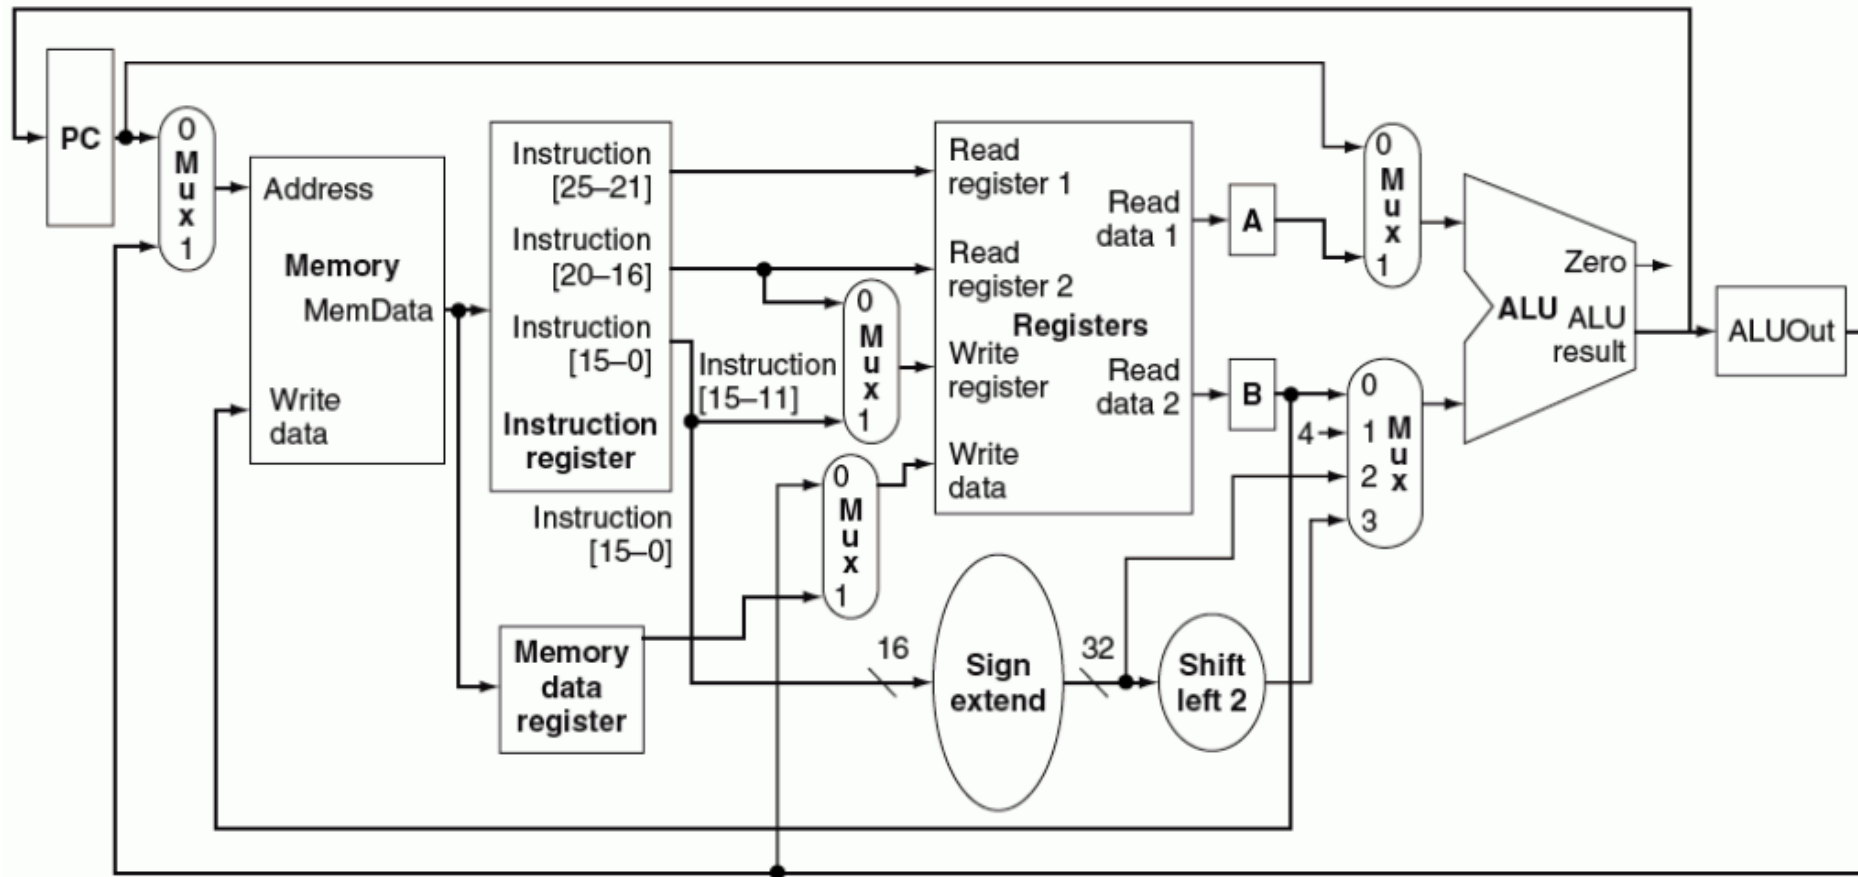
Partitioning Single-Cycle Design



Where to Add Registers?



Multicycle Datapath (cont.)



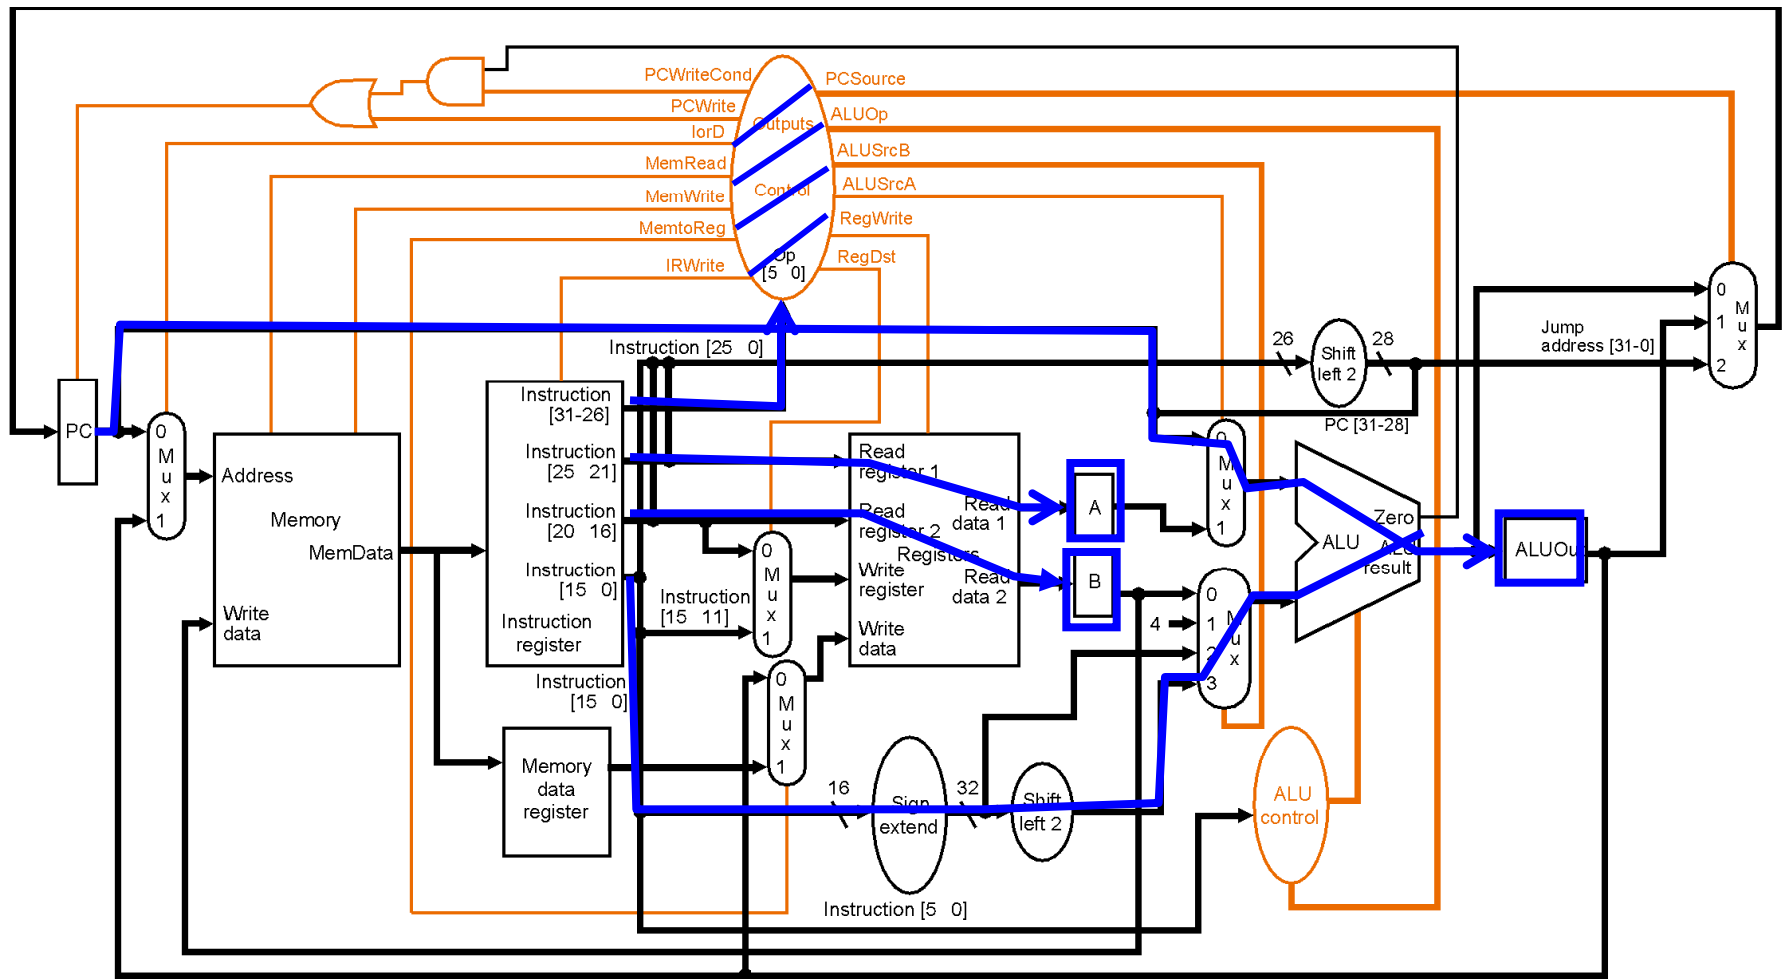
The diagram illustrates the internal components and data flow of a MIPS processor. Key elements include:

- Control Signals:** A set of control signals (PCWriteCond, PCWrite, IorD, MemRead, MemWrite, MemtoReg, IRWrite, Op, ALUOp, ALUSrcB, ALUSrcA, RegWrite, RegDst) is shown at the top, originating from the Instruction Register and branching to various units.
- PC (Program Counter):** Receives the 'Jump address [31-0]' and outputs the 'PC [31-28]' to the 'Shift left 2' block.
- Instruction Register:** Receives the 'Instruction [31-26]' and outputs various fields to the Memory, Registers, and ALU control.
- Memory:** Receives the 'Address' and outputs 'MemData' to the 'Memory data register'.
- Registers:** Contains 'Read register 1', 'Read register 2', and 'Write register'. It receives 'Read data 1', 'Read data 2', and 'Write data' from the Instruction Register and outputs to the ALU.
- ALU (Arithmetic Logic Unit):** Receives 'ALU control' and 'ALUSrcA' from the Instruction Register, and 'ALUSrcB' from the Registers. It outputs the 'ALUOut' and the 'Zero' flag.
- Multiplexers:** Several multiplexers (labeled 'Mux 1', 'Mux 2', 'Mux 3') are used to select between different data paths, such as the 'PC [31-28]', 'ALUOut', and 'Jump address [31-0]'.
- Shifters:** 'Shift left 2' blocks are used to shift the 'PC [31-28]' and the 'Sign extend' output of the 'Memory data register'.
- Sign Extension:** A 'Sign extend' block takes the 'Memory data register' output and shifts it left by 2 bits to produce a 32-bit result.

```
IR <= Mem[PC]
PC <= PC + 4
```

lorD=0, MemRead=1, MemWrite=0,
IRwrite=1, ALUsrcA=0, ALUsrcB=01,
PCWrite=1, ALUOp=00, PCsource=00

Cycle 2: ID & RF Cycle



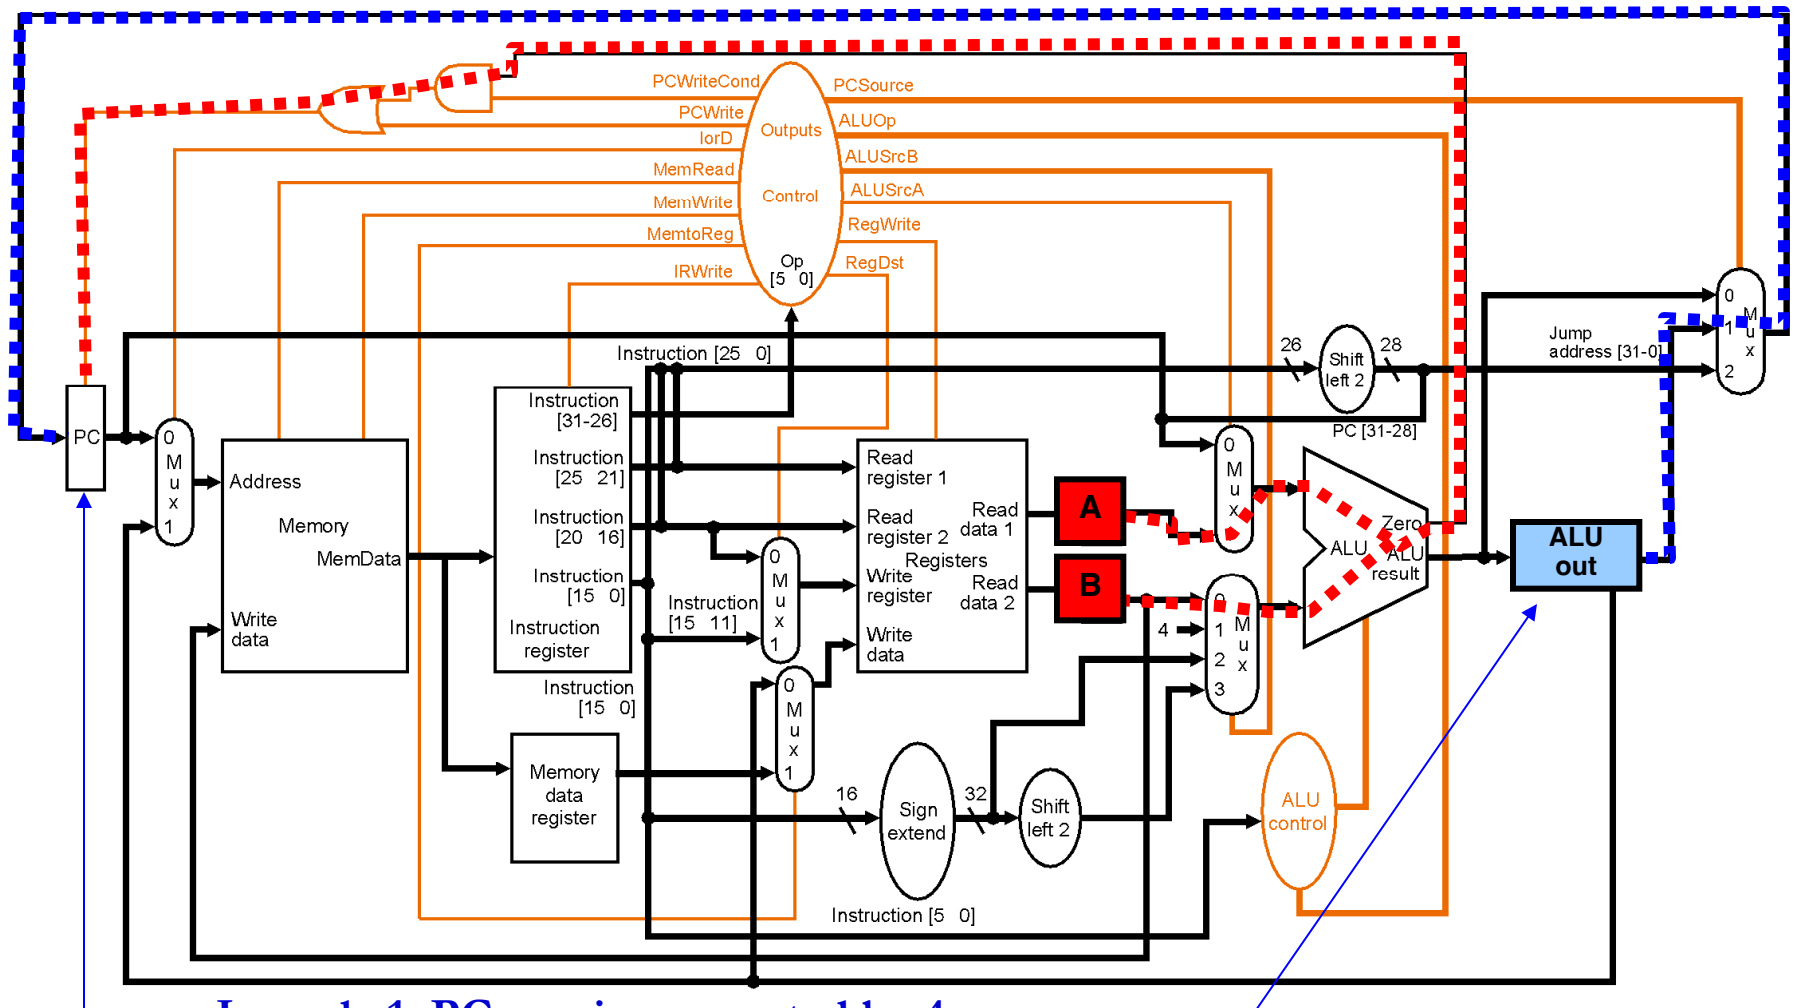
$A \leftarrow \text{GPR}[\text{IR}[25-21]]$

$B \leftarrow \text{GPR}[\text{IR}[20-16]]$

$\text{ALUOut} \leftarrow \text{PC} + (\text{sign-extend}(\text{IR}[15-0]) \ll 2)$

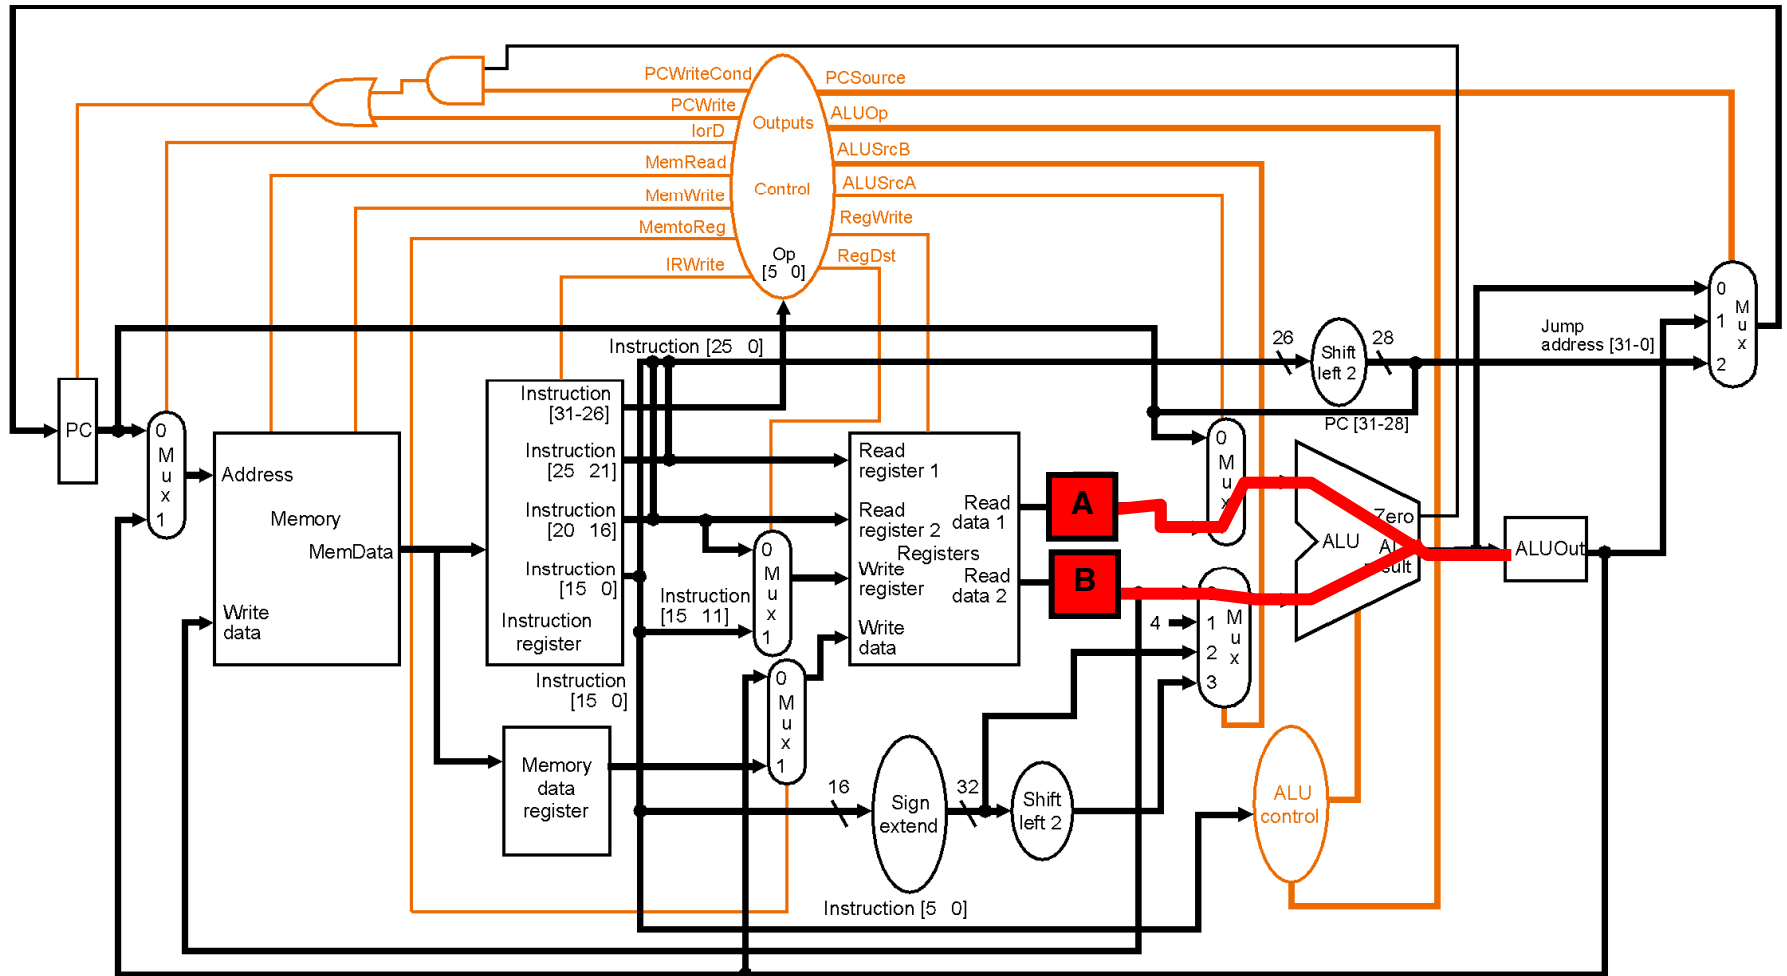


Cycle 3 for beq: Execute



- In cycle 1, PC was incremented by 4
- In cycle 2, ALUout was set to branch target
- This cycle, we conditionally update PC: if (A==B) PC=ALUout

R-Type Execution

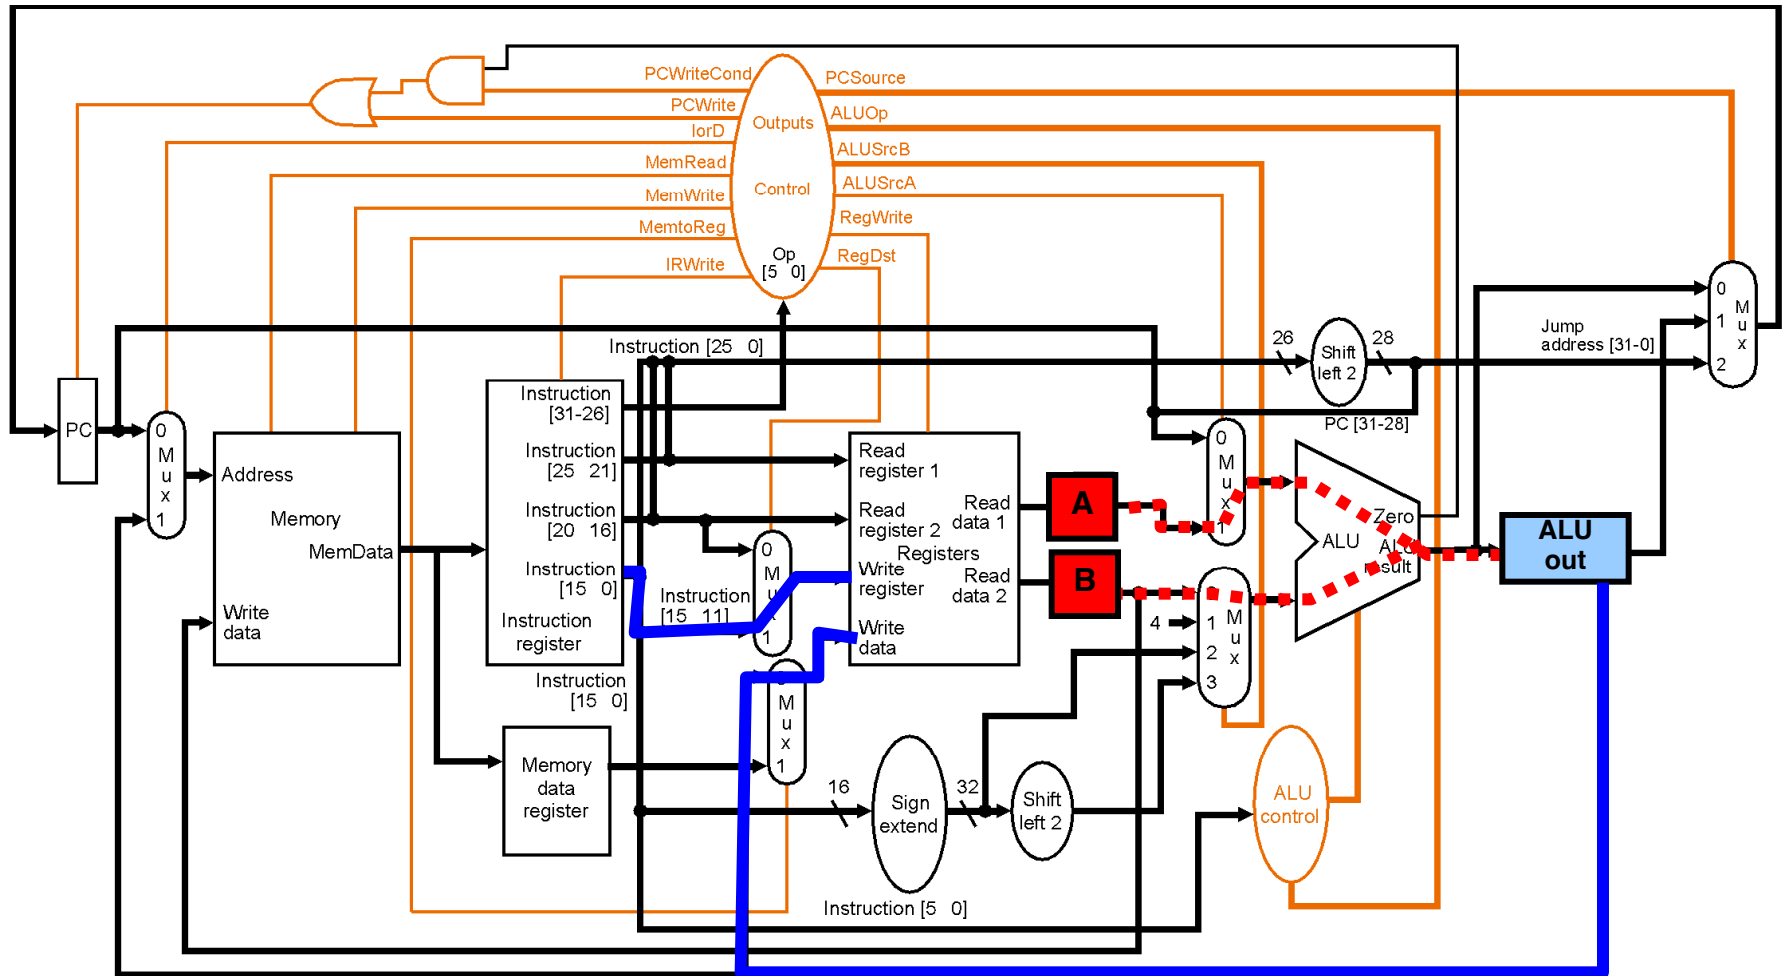


Cycle 3: **ALUout = A op B**


Cycle 4: `GPR[IR[15-11]] = ALUout`



R-Type Execution & WB




Cycle 3 for lw/sw: Address Computation



$$\text{ALUout} = A + \text{sign-extend}(\text{IR}[15-0])$$



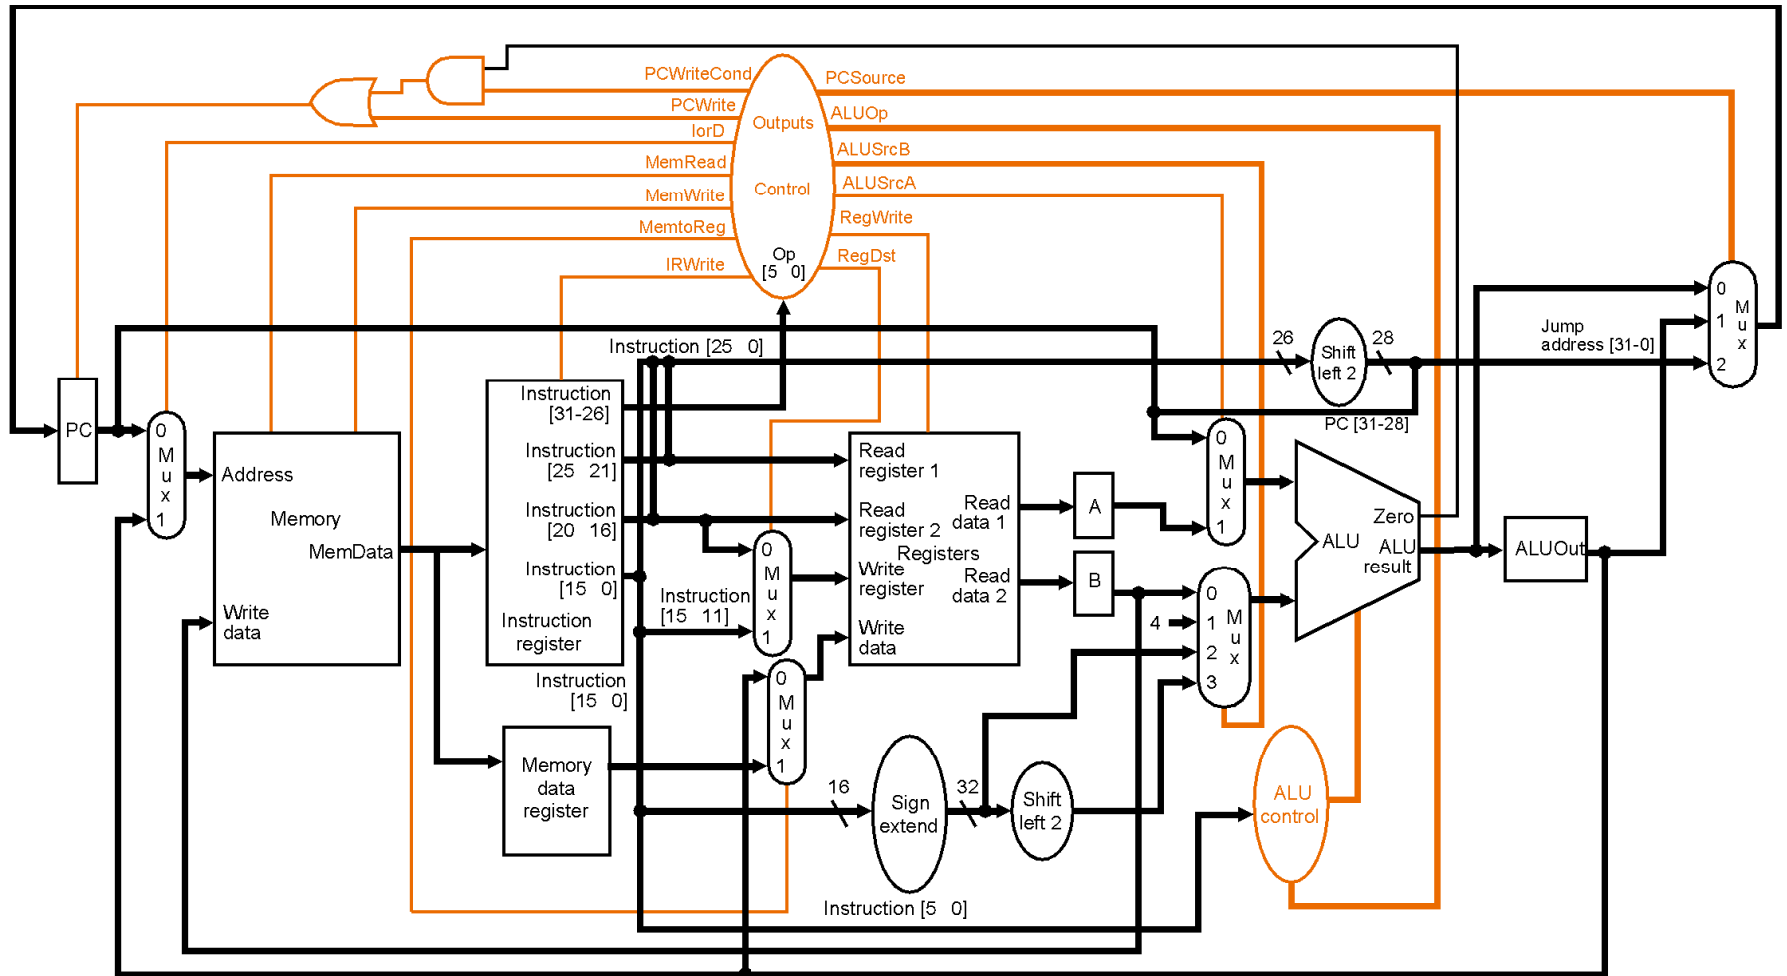
Cycle 4 for Store: Memory Access



Memory[ALUout] = B




Cycle 4 for Load: Memory Access



MDR = Memory[ALUOut]



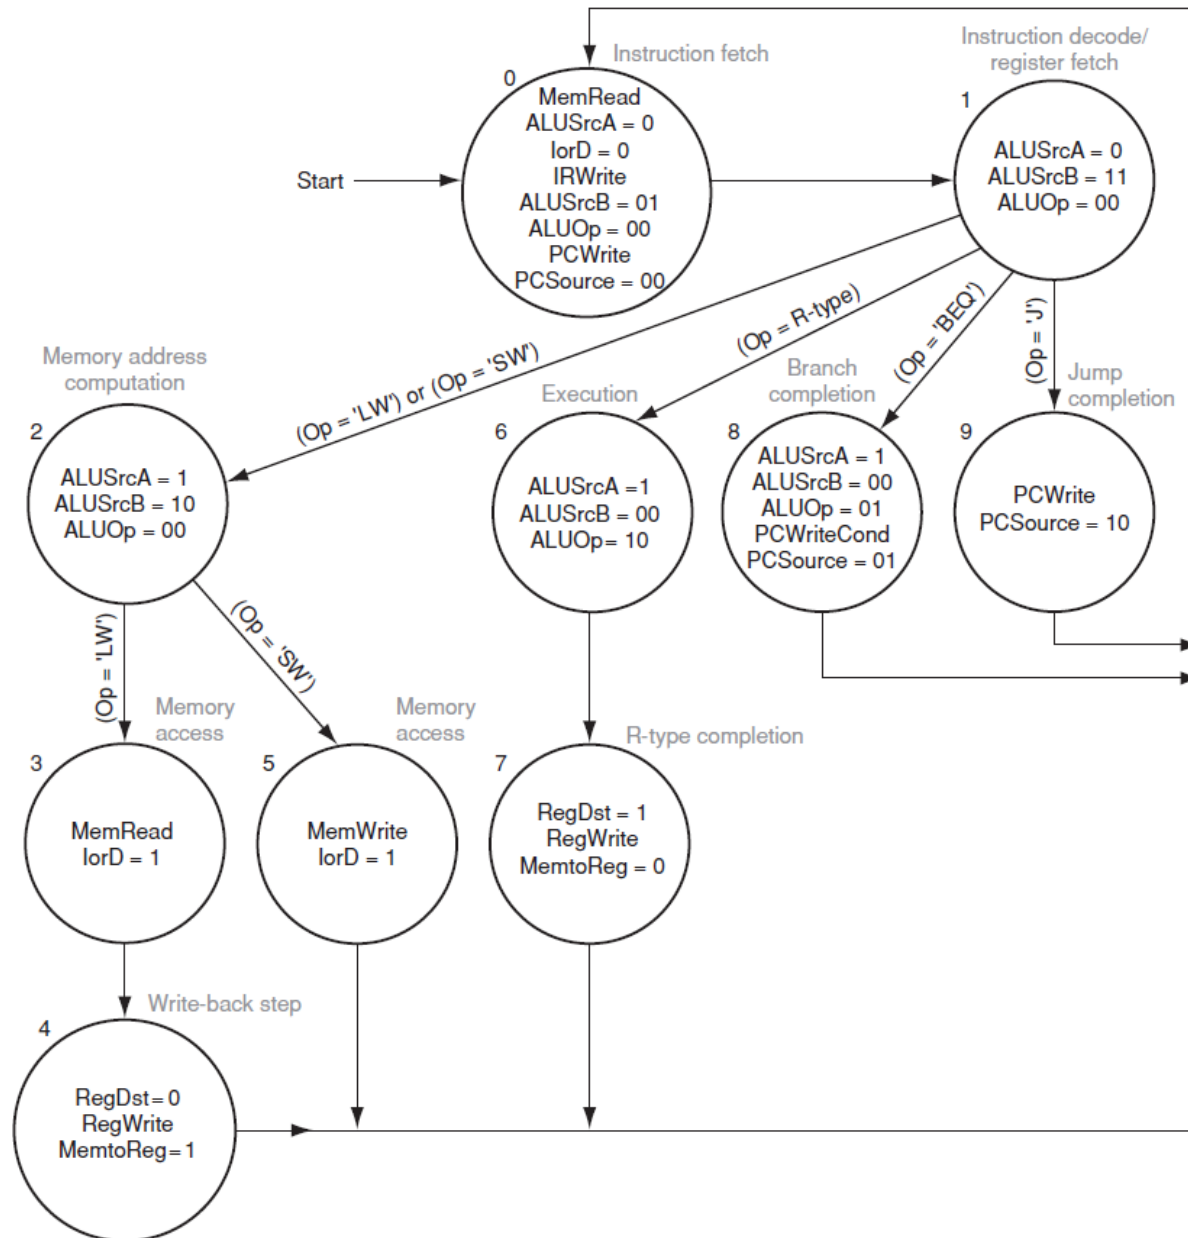
Cycle 5 for load: WriteBack



GPR[IR[20-16]] = MDR



Complete FSM



Our Lectur Today




Topics Covered Today

- **Basics of Pipelined Design**
- **MIPS Pipelined Datapath**
- **Pipelined Performance**
- **Hazards in Pipelined Design**
 - **Structural hazards**
 - **Data hazards**
 - **Control hazards**




Instruction Latencies and Throughput

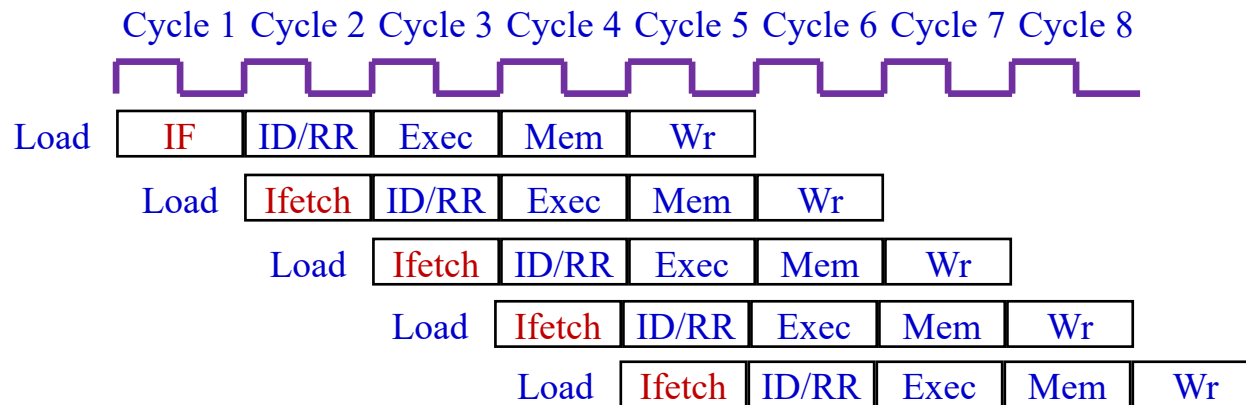
•Single-Cycle CPU



•Multiple Cycle CPU



•Pipelined CPU




Introduction of Pipeline

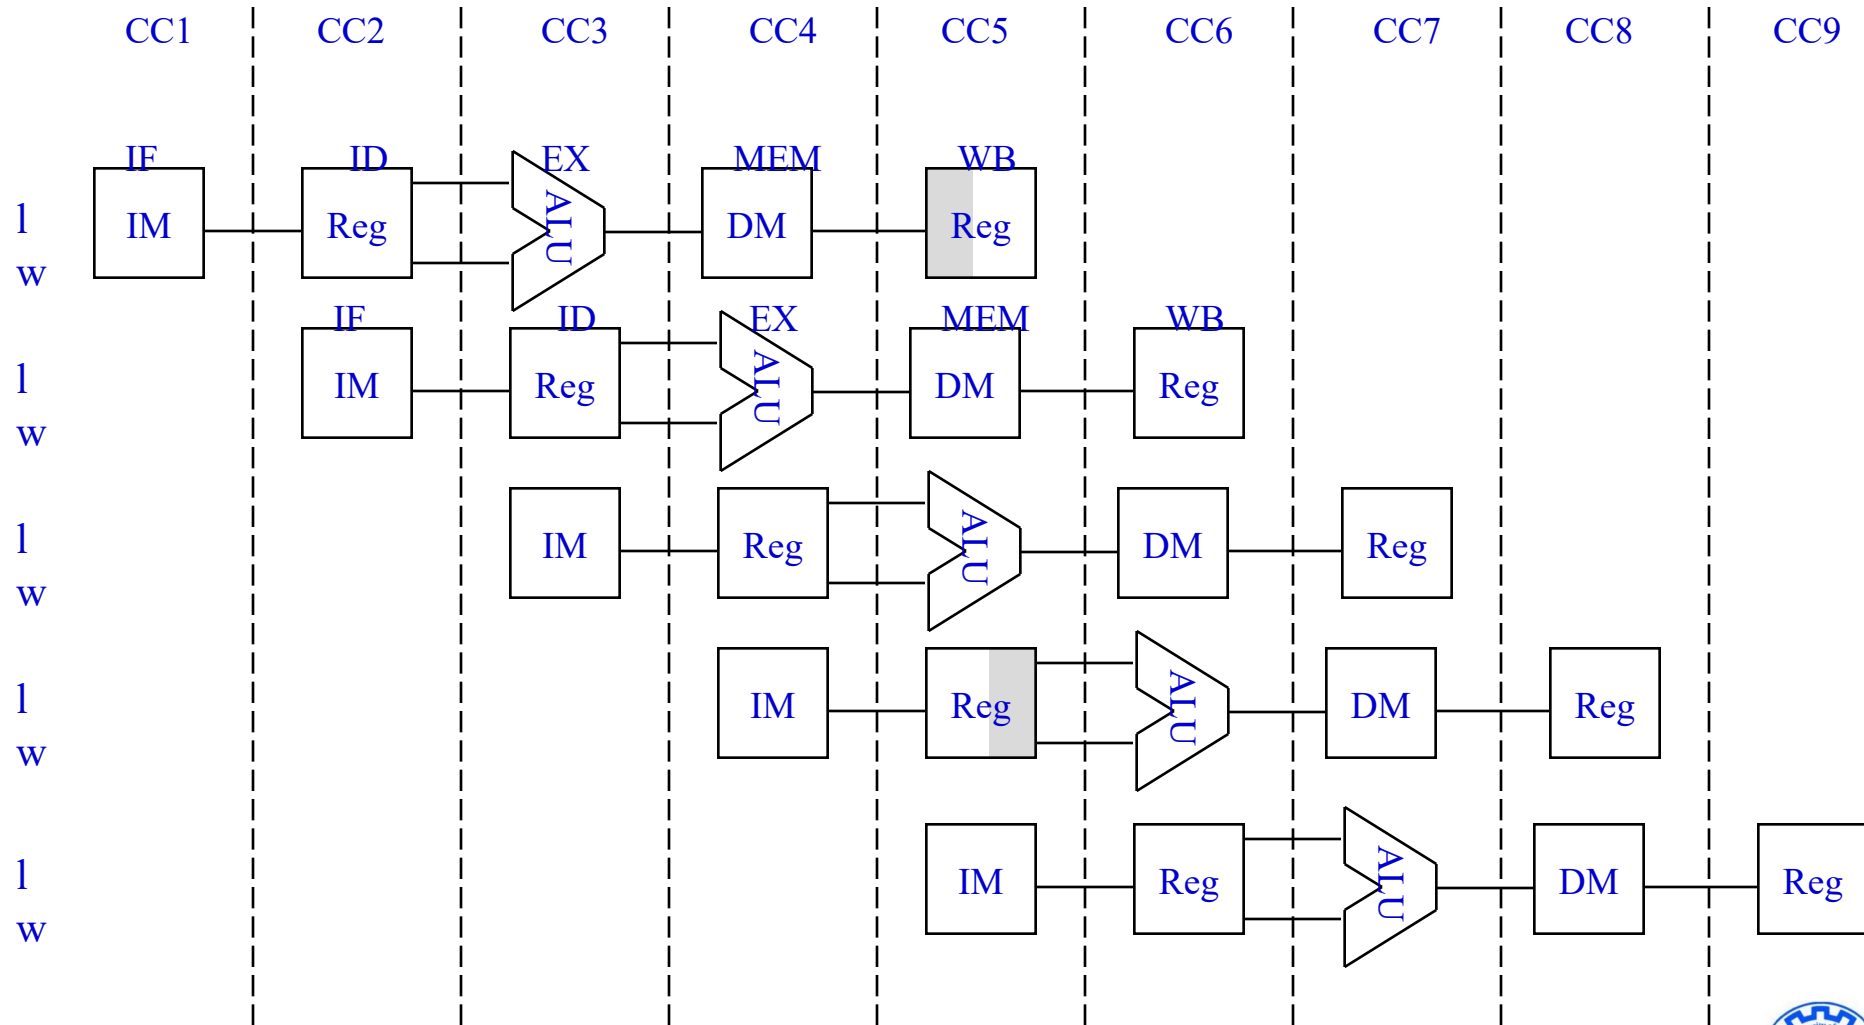
Machine assembly line



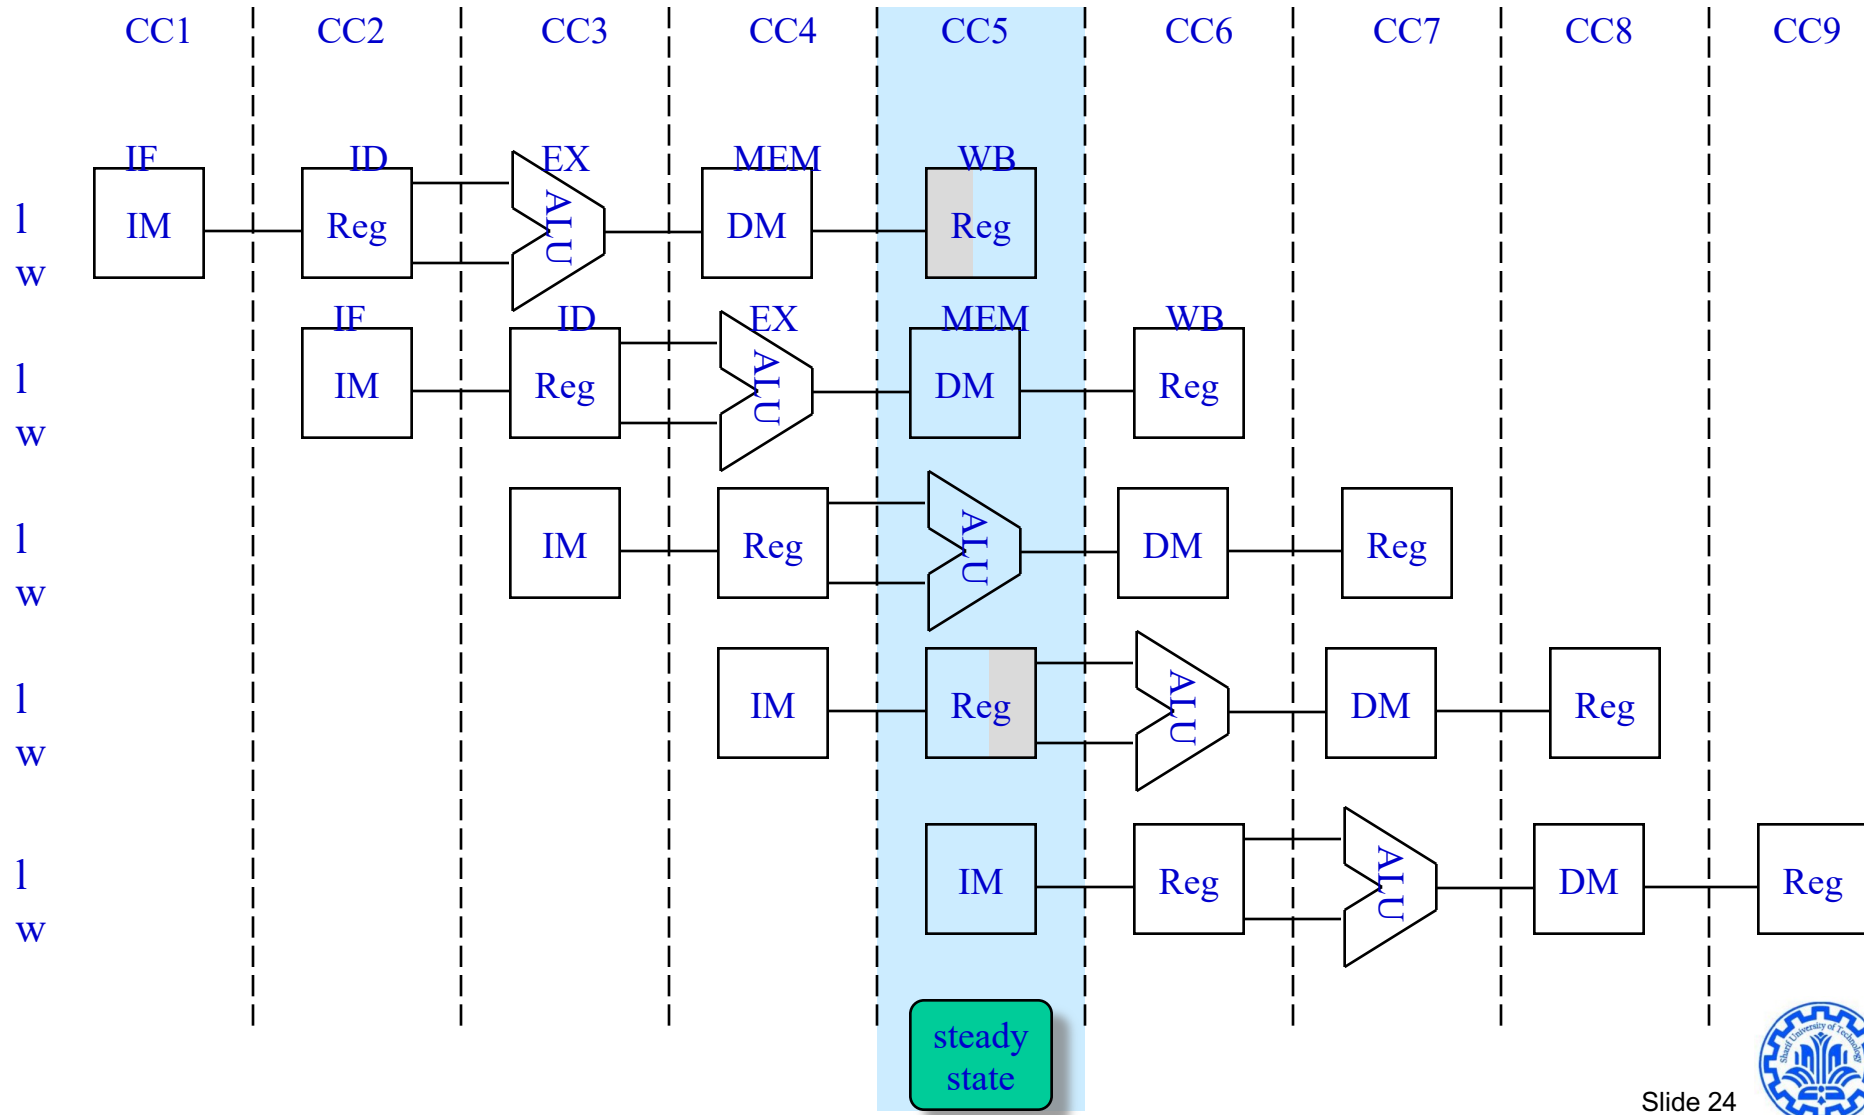
Datapath of pipeline



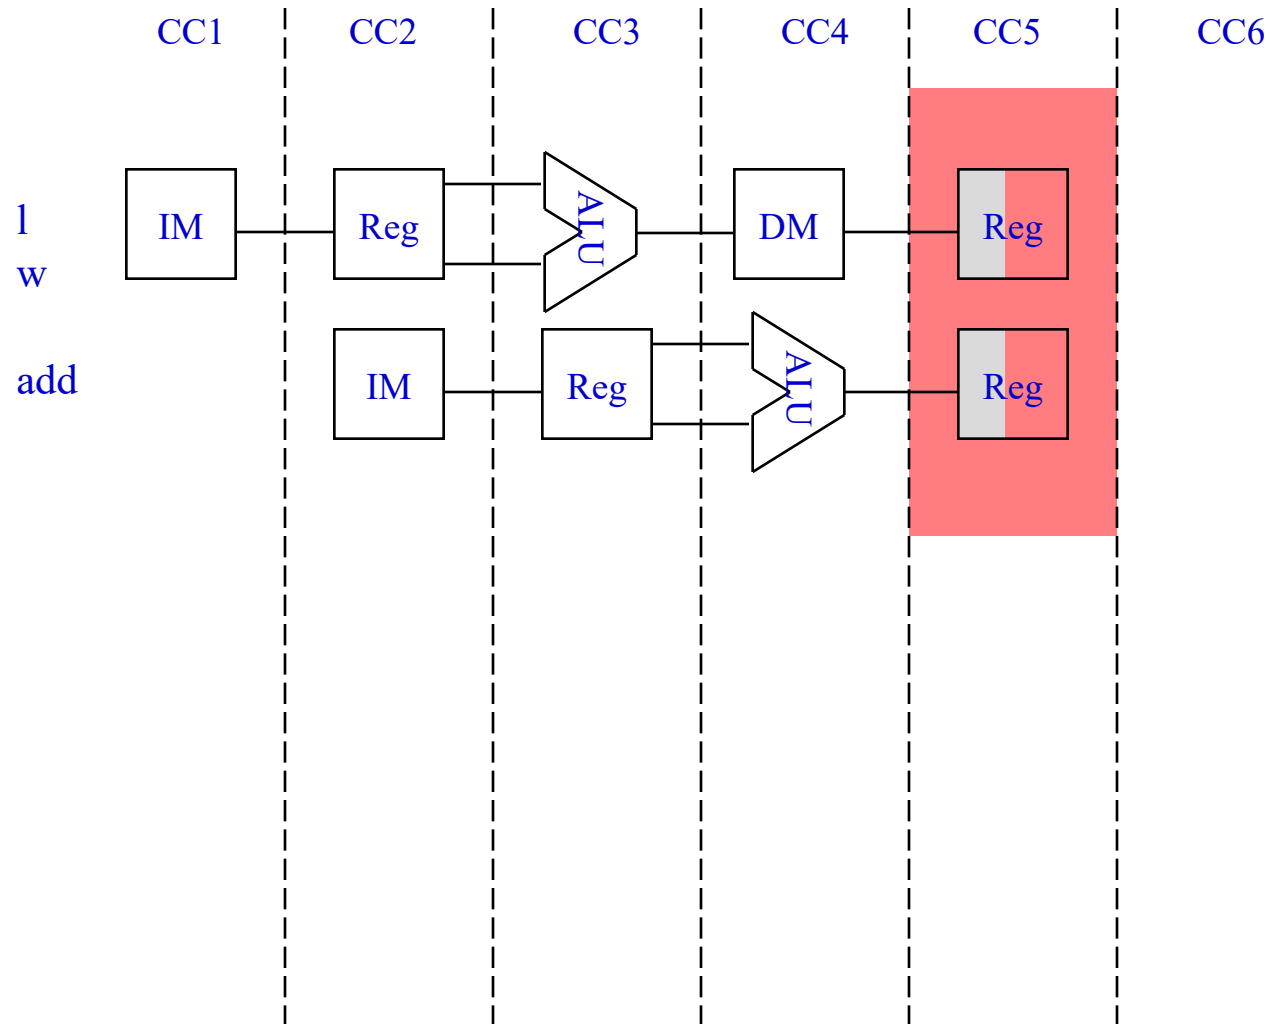
Execution in a Pipelined Datapath



Execution in a Pipelined Datapath




Mixed Instructions in Pipeline




Pipeline Principles

- Principles
 - All instructions that share a pipeline must have same *stages* in same *order*
 - ➔ *add* does nothing during **Mem** stage
 - ➔ *sw* does nothing during **WB** stage




Pipeline Principles (cont.)

- Principles
 - All intermediate values **must be latched** each cycle
 - **No functional block reuse** for an instruction
 - E.g. need 2 adders and ALU (like in single-cycle)

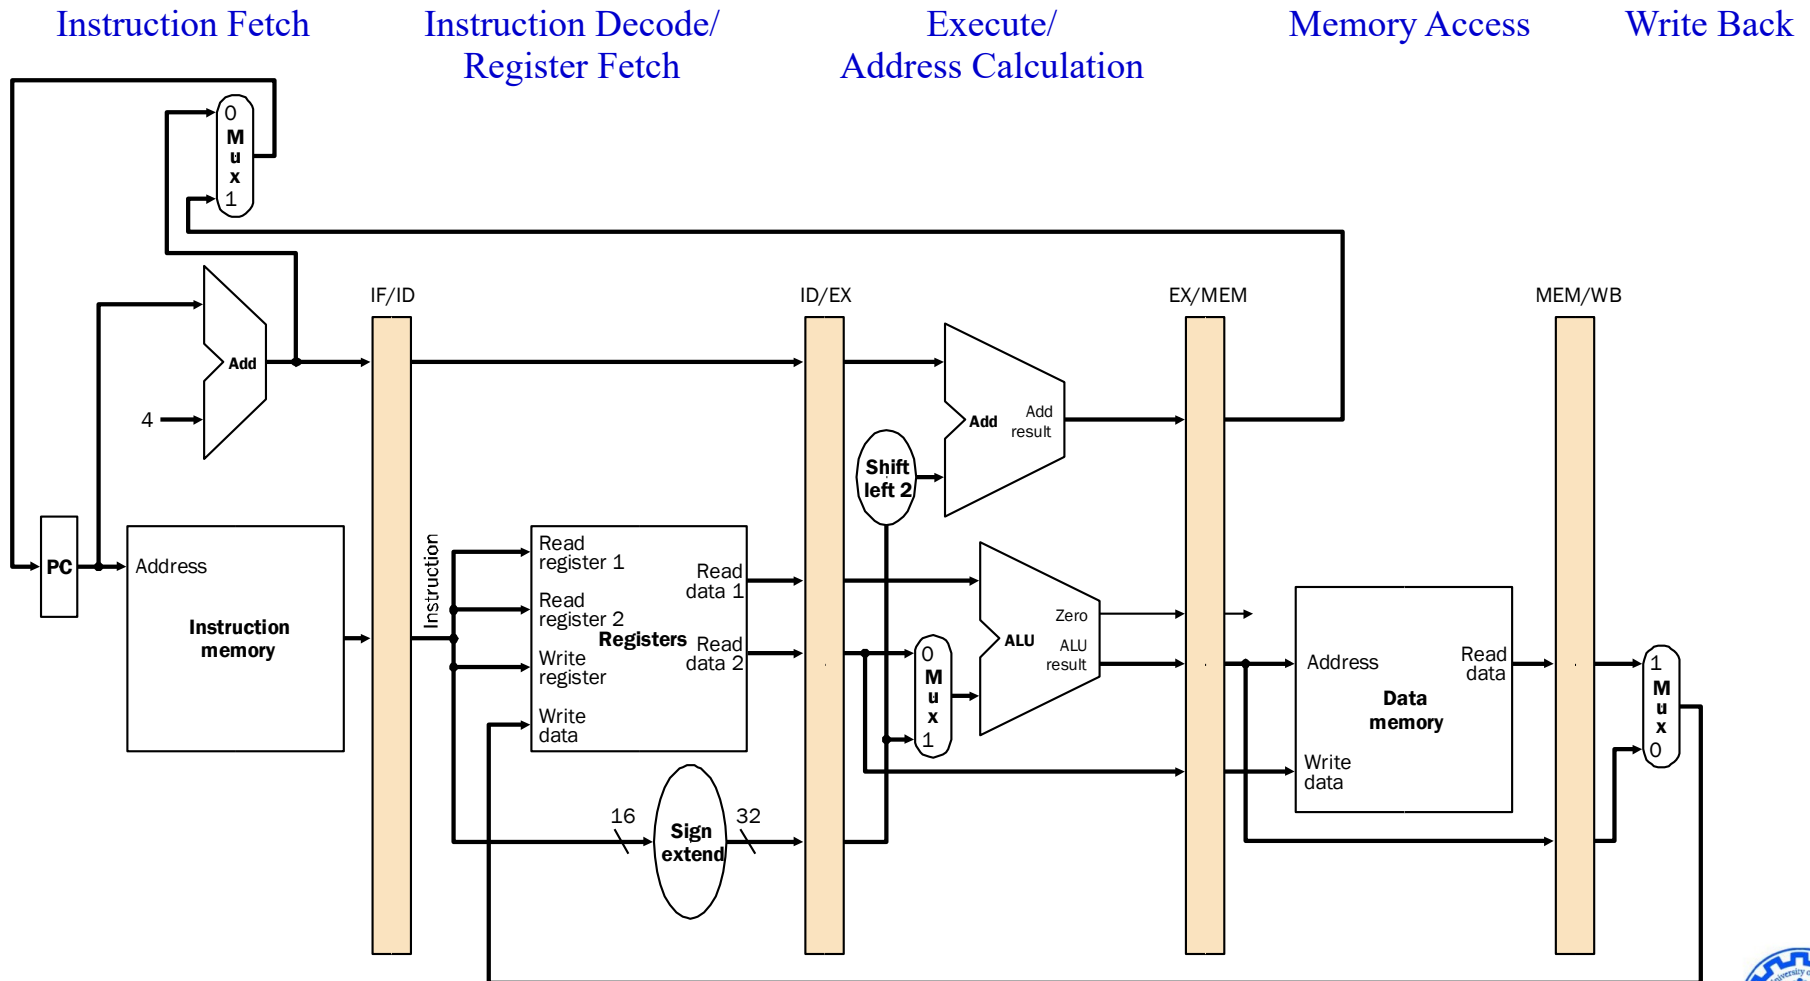


Pipeline Principles

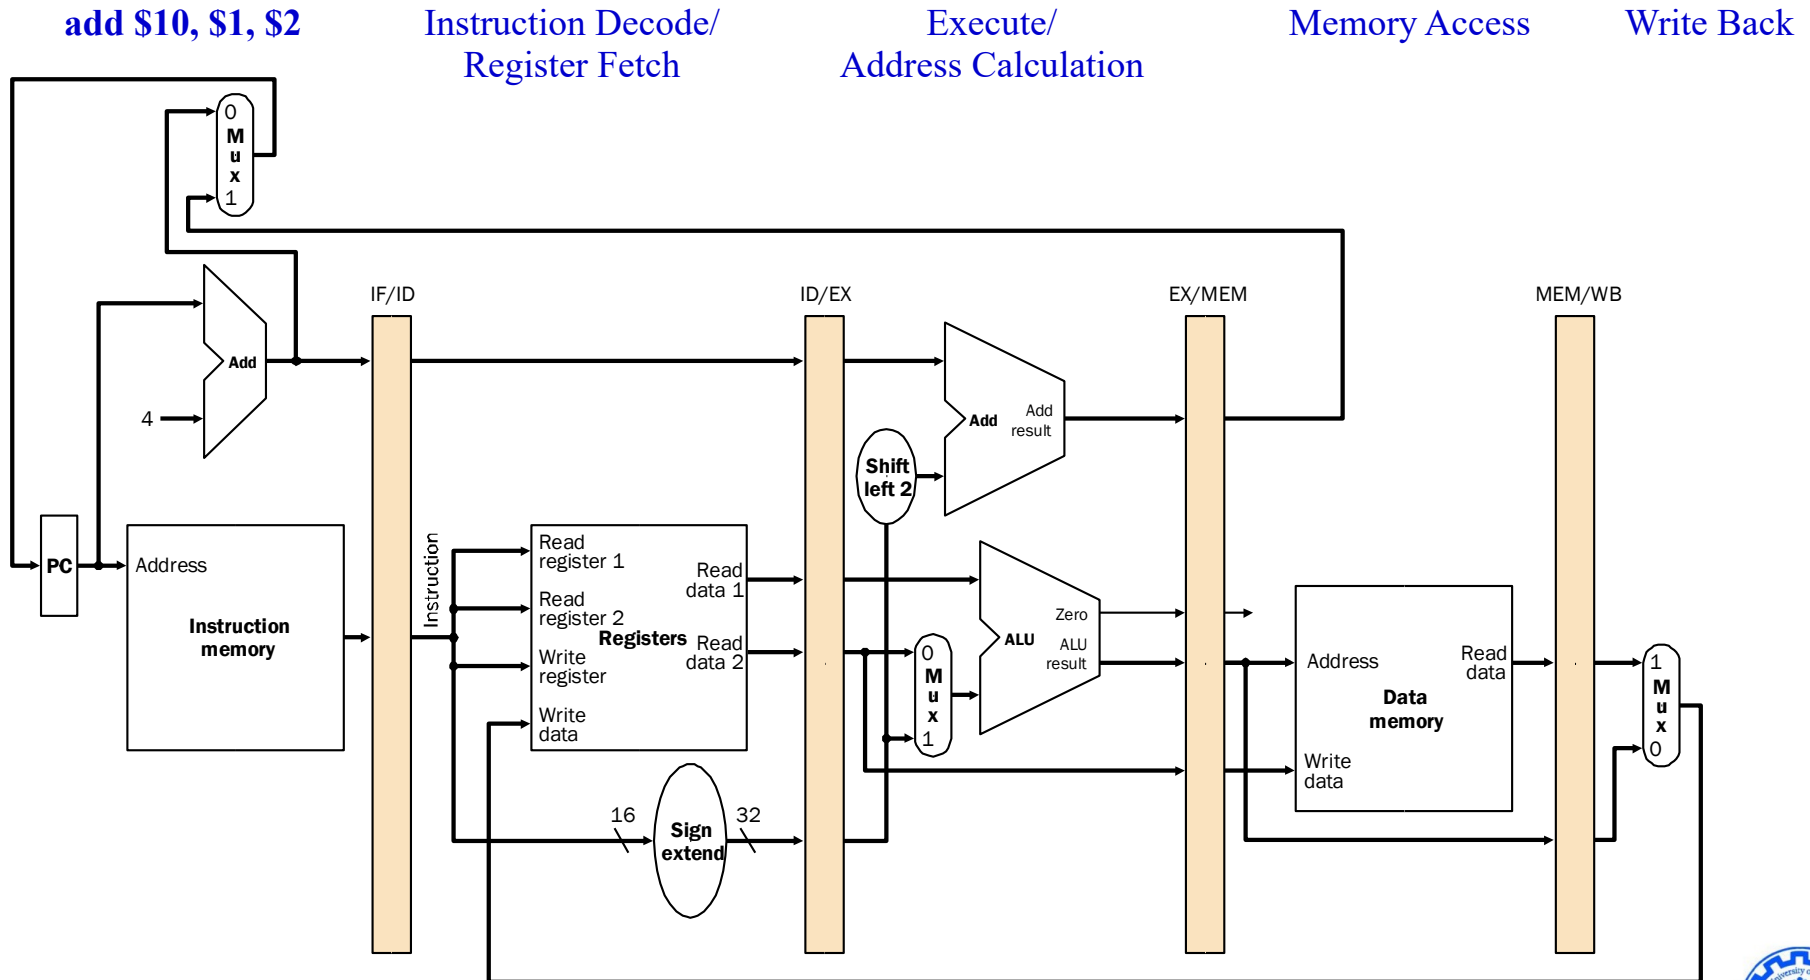
- Question:
 - What if we want to bypass a stage for an instruction?



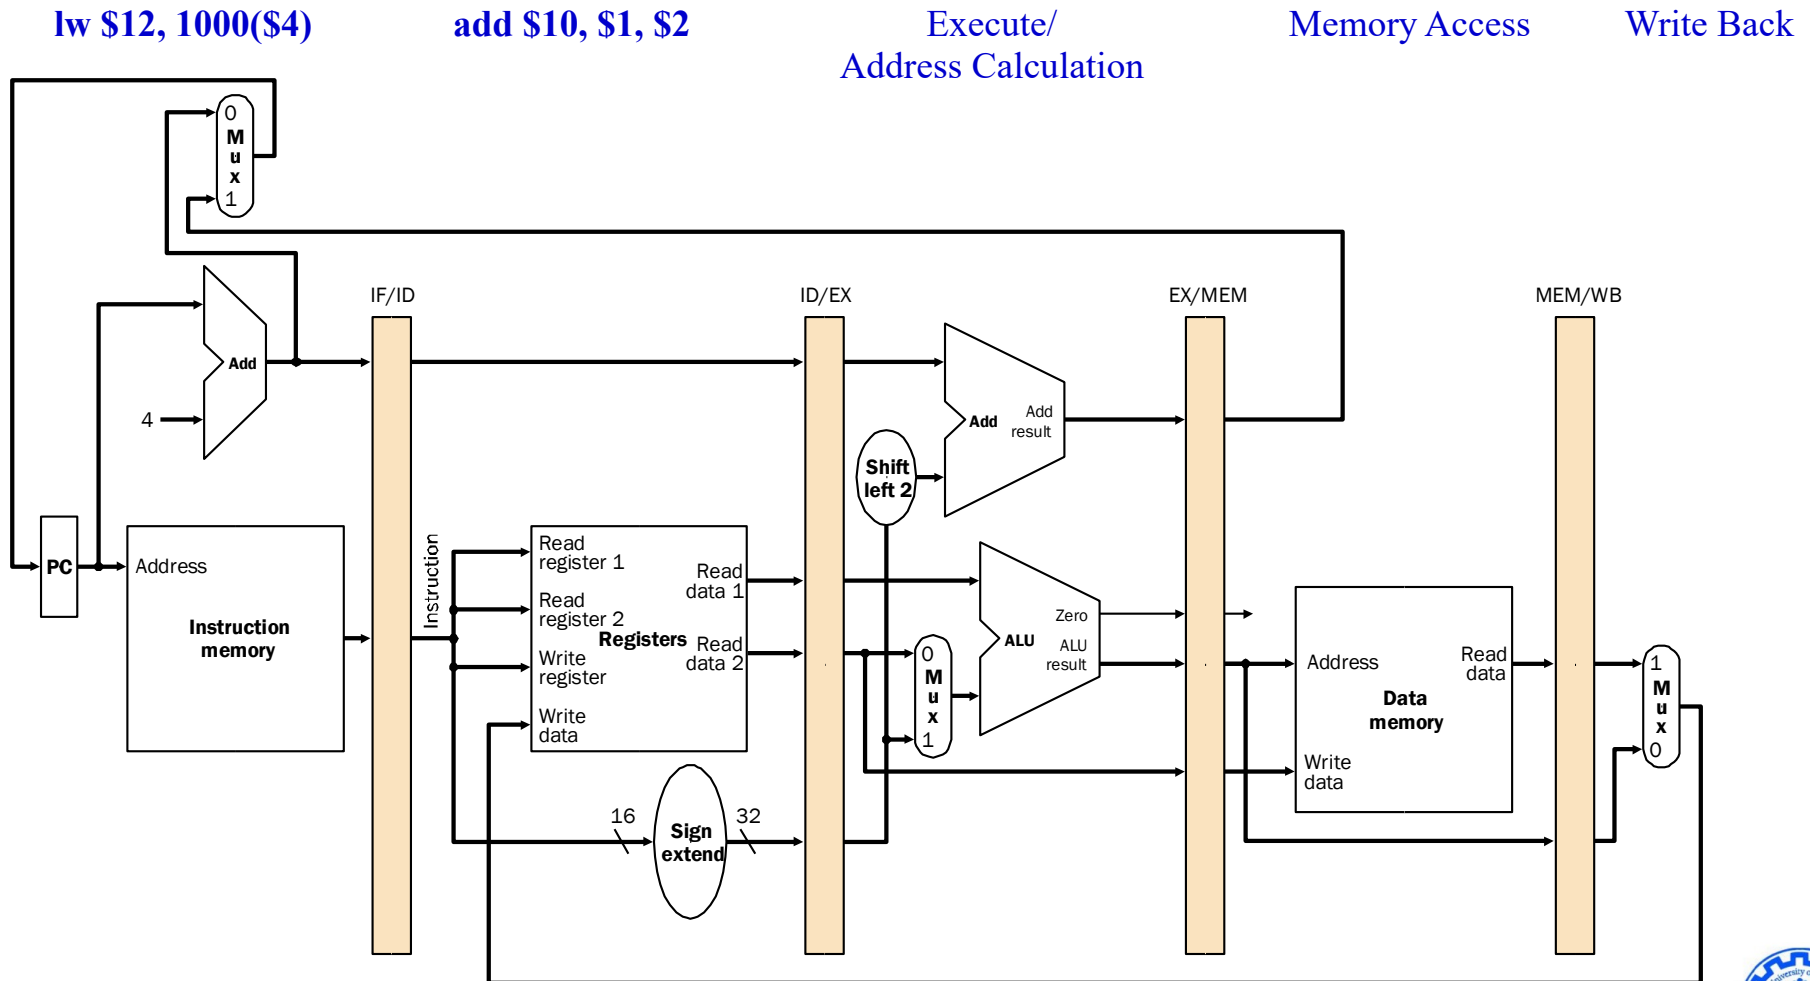
Pipeline in Execution



Pipeline in Execution



Pipeline in Execution



Pipeline in Execution

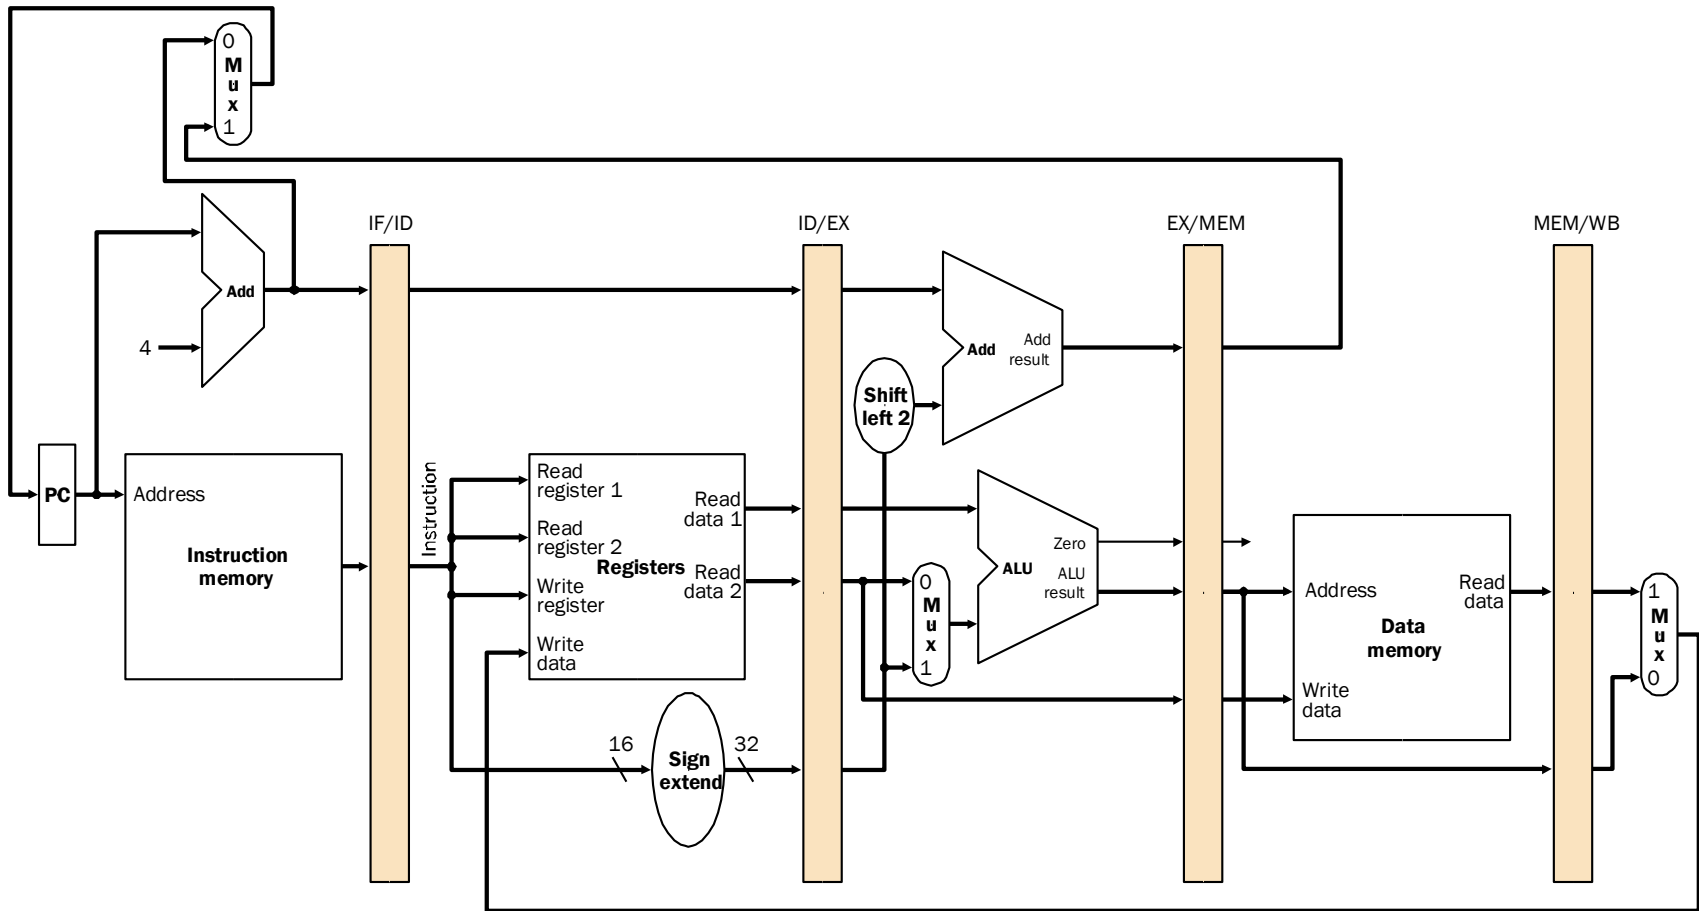
sub \$15, \$4, \$1

lw \$12, 1000(\$4)

add \$10, \$1, \$2

Memory Access

Write Back



Pipeline in Execution

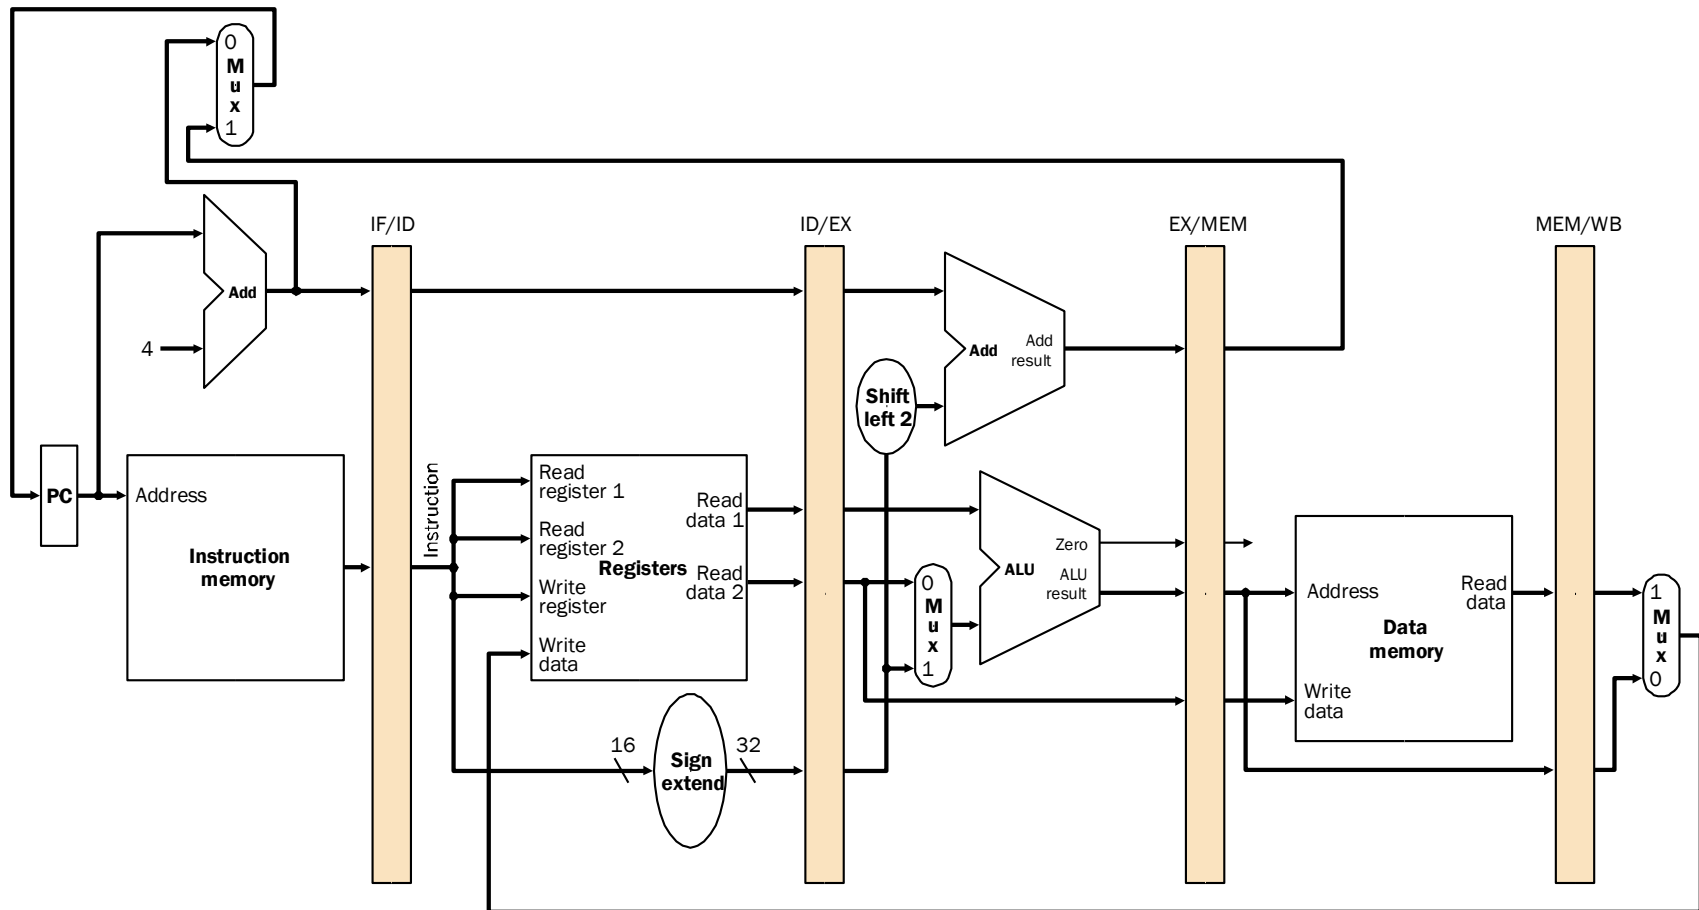
Instruction Fetch

sub \$15, \$4, \$1

lw \$12, 1000(\$4)

add \$10, \$1, \$2

Write Back



Pipelining Performance

- $ET = IC * CPI * CT$
 - Achieve High *throughput*
 - Without reducing instruction *latency*
 - Example:
 - A CPU that takes 5ns to execute an instruction pipelined into 5 equal stages
 - Latch between each stage has a delay of 0.25 ns
1. Min. clock cycle time of this arch?
 2. Max. speedup that can be achieved by this arch (compared to single cycle arch)?



Issues With Pipelining

- **Pipelining Creates Potential Hazards**
 - What happens if two instructions need a same structure?
 - Structure hazard
 - What happens when an instruction needs result of another instruction?
 - Data hazard
 - What happens on a branch?
 - Control hazard

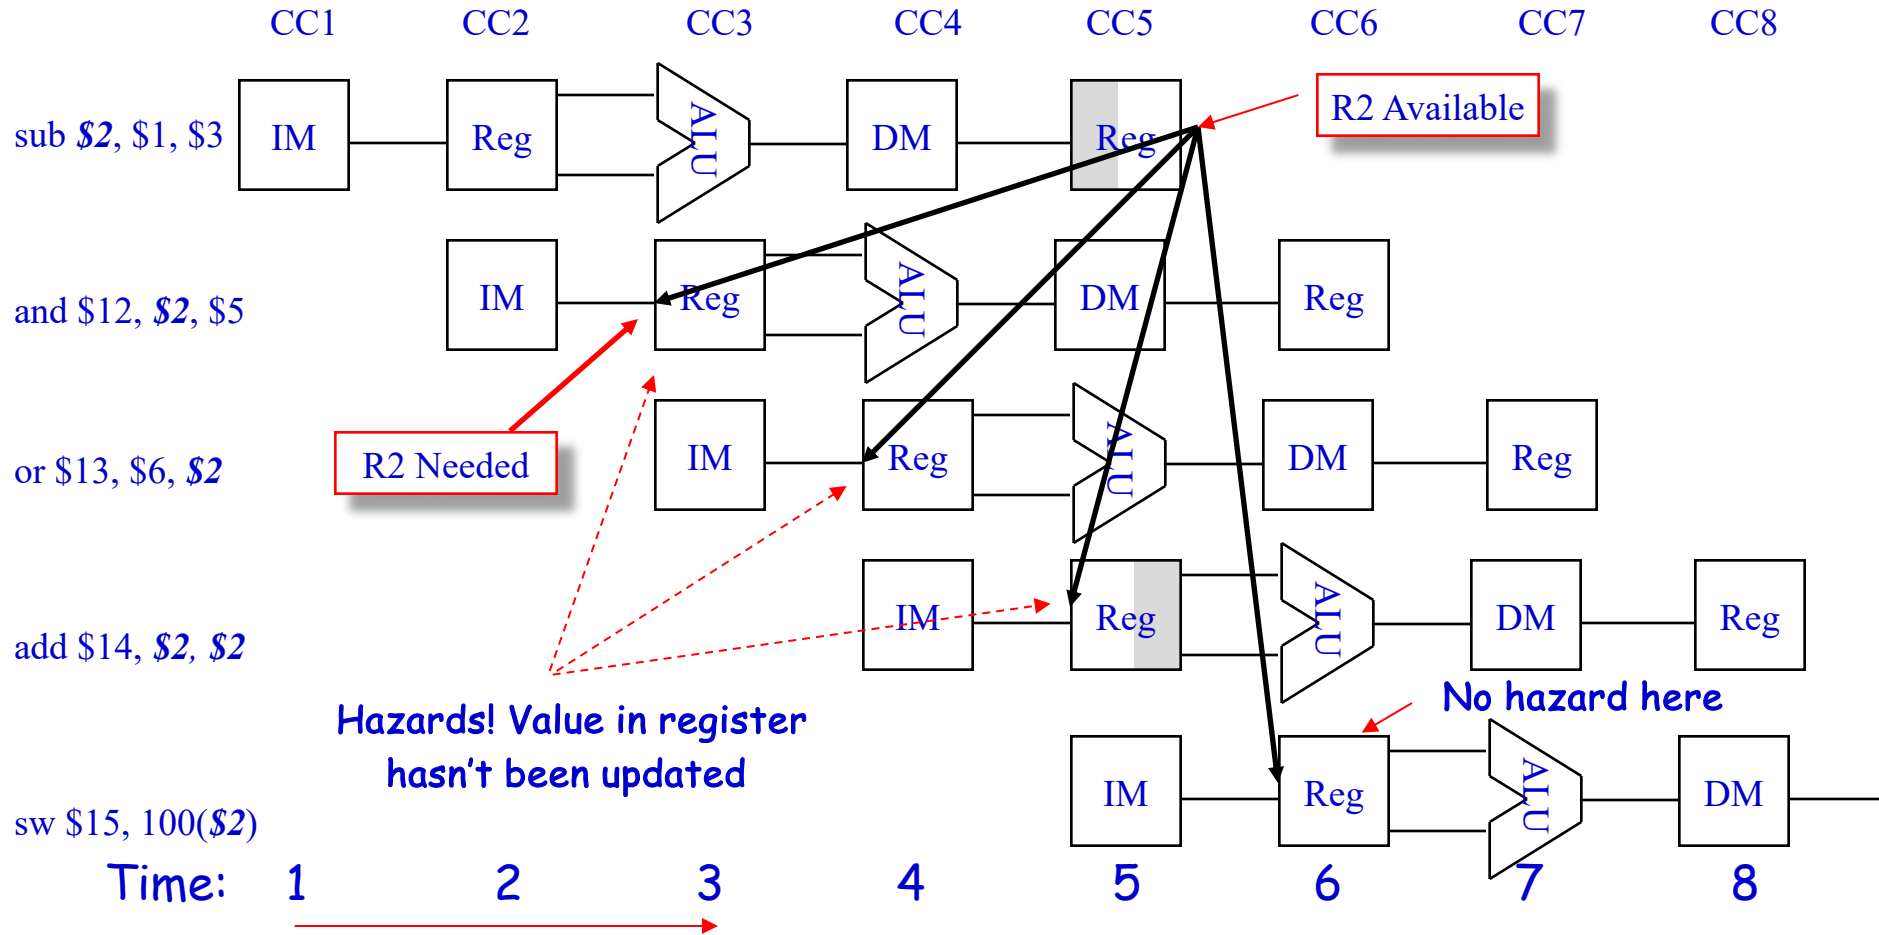


Structural Hazards

- How?
 - Two instructions require **use** of a given **hardware** resource at **same time**
- Access to **Memory**
 - Separate instruction and data caches
- Access to **Register File**
 - Multiple port register file



Data Hazards



Data Hazard: result needed before it is available

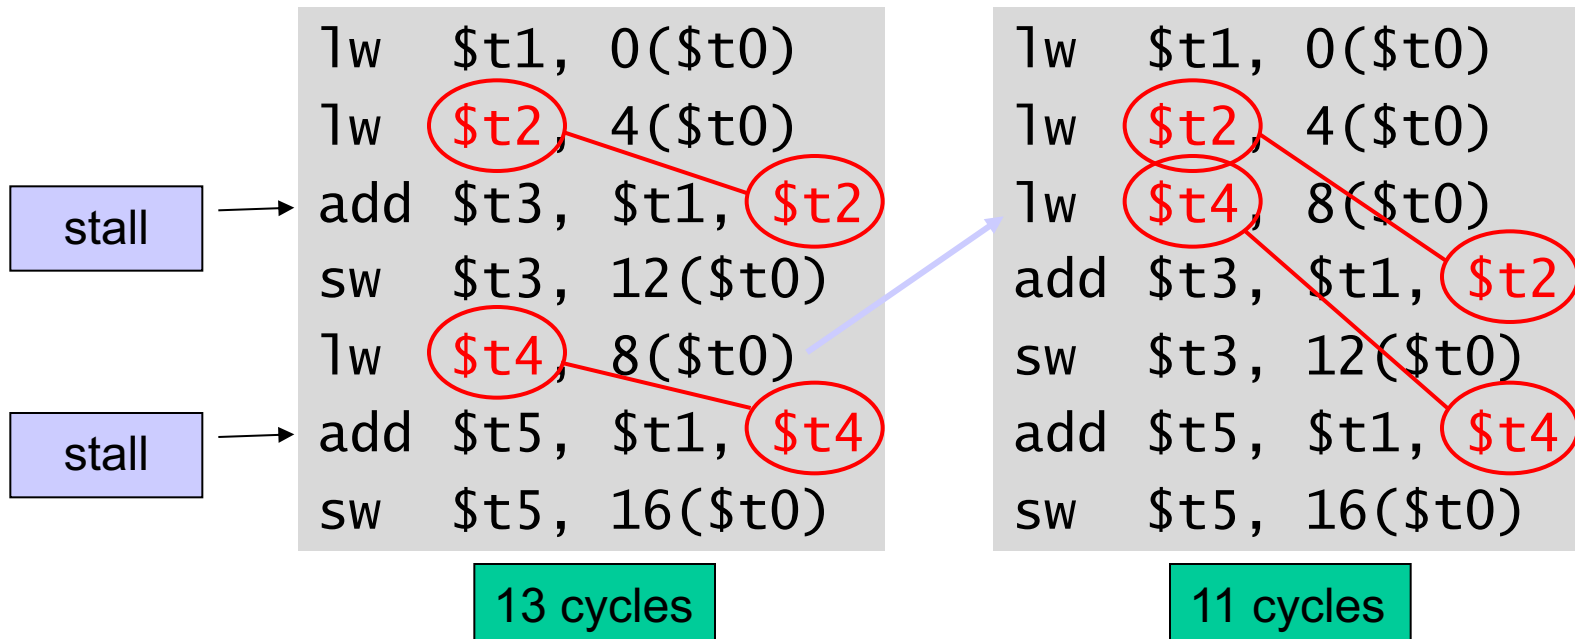
Handling Data Hazards

- **SW** Solution
 - Insert independent instructions
 - No-ops instructions
 - Code reordering
- **HW** Solutions
 - Insert bubbles (i.e. stall pipeline)
 - Data forwarding
 - Latch-based GPR



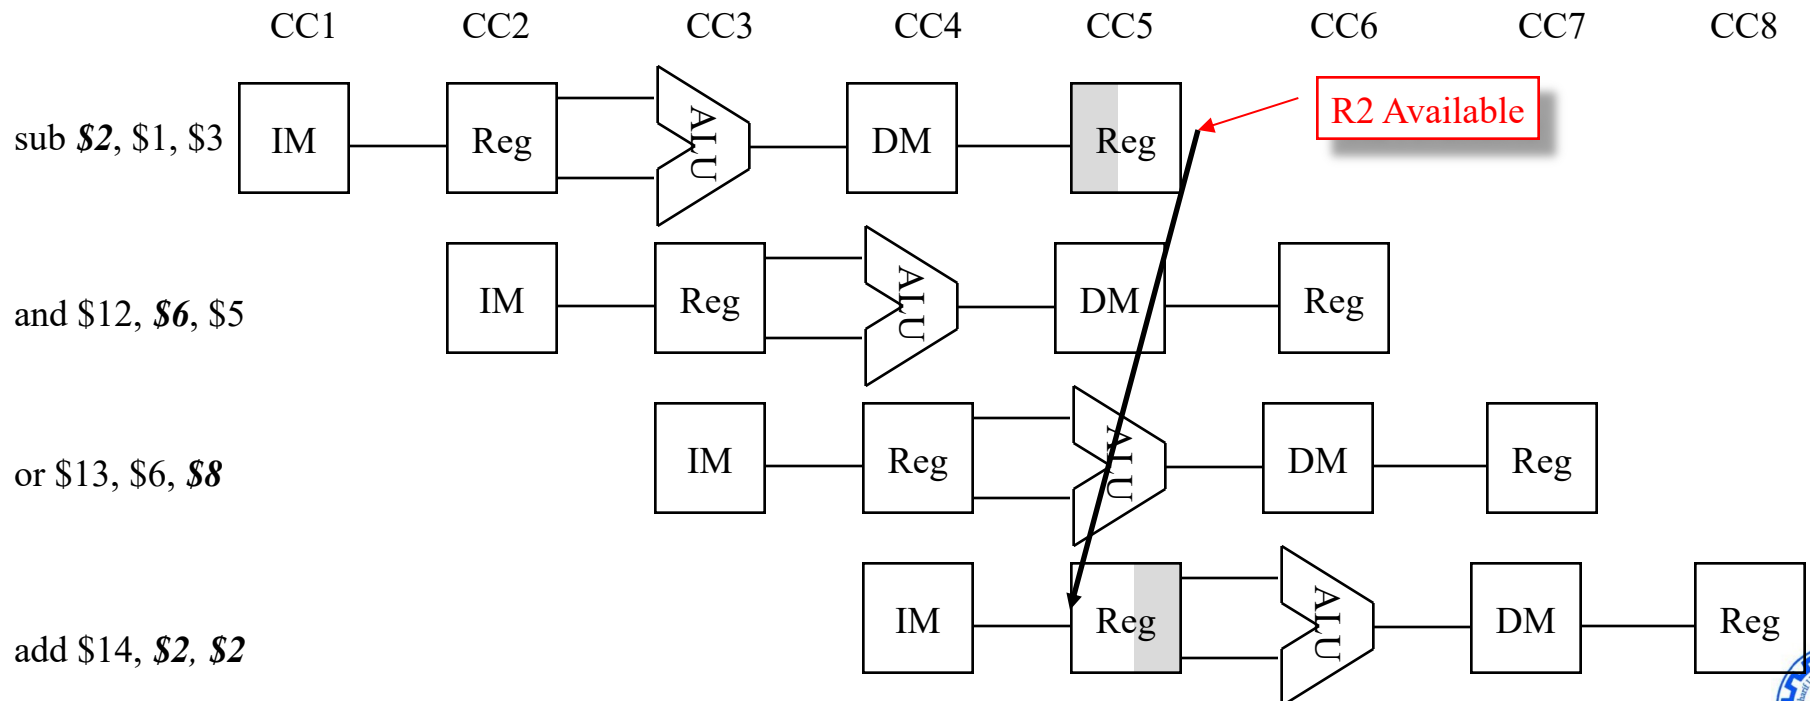
Dealing With Data Hazards

- Reorder code to avoid use of load result in next instruction
- C code for $A = B + E; C = B + F;$

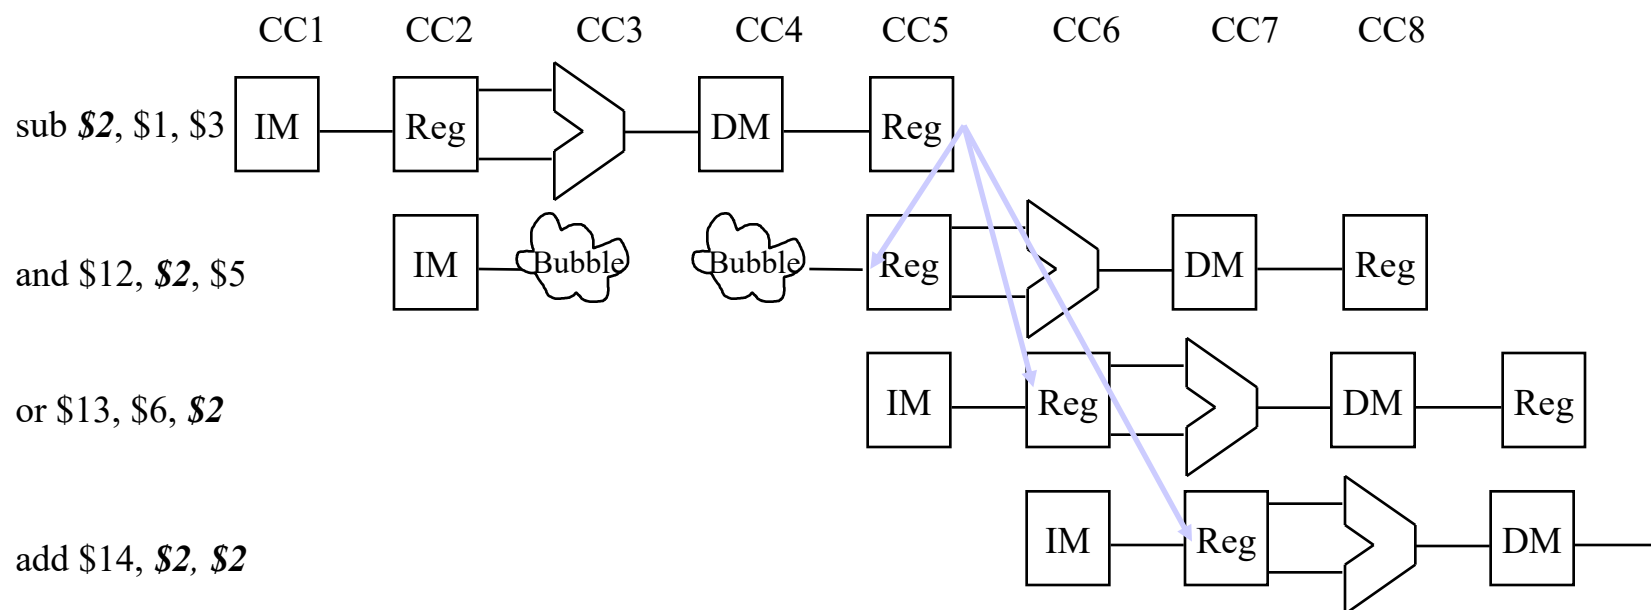


Dealing With Data Hazards

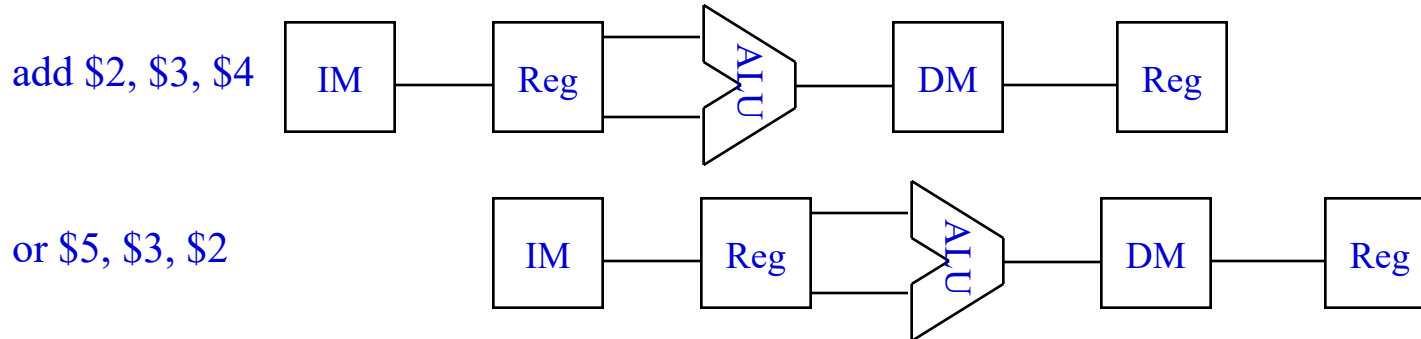
- Using Transparent Register File
 - Use latches rather than flip-flops in RF
 - First half-cycle : register loaded
 - Second half-cycle: new value is read into pipeline state



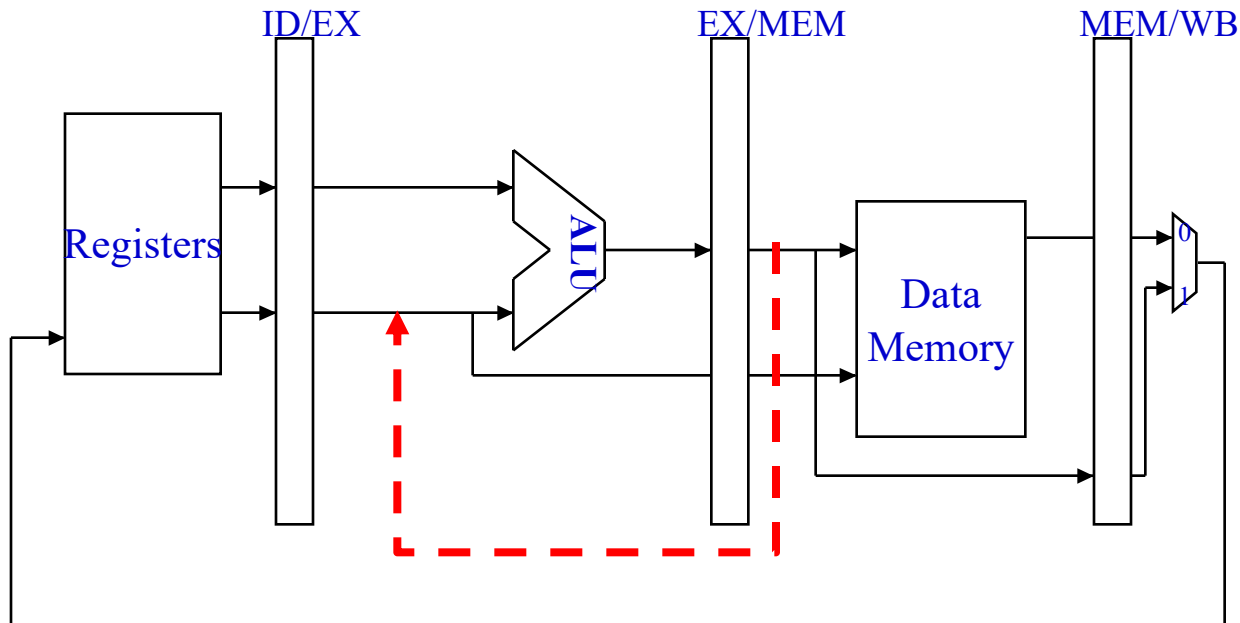
Handling Data Hazards in Hardware: Stall pipeline



Reducing Data Hazards Through Forwarding



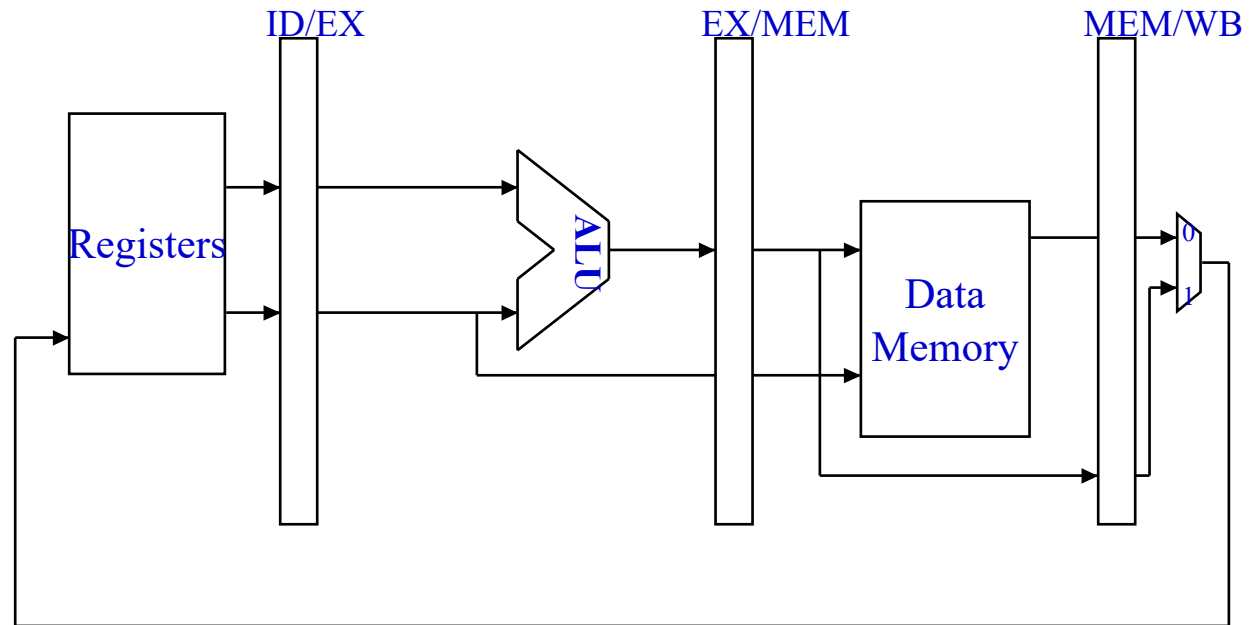
Avoid stalling by forwarding ALU output from "add" to ALU input for "or"



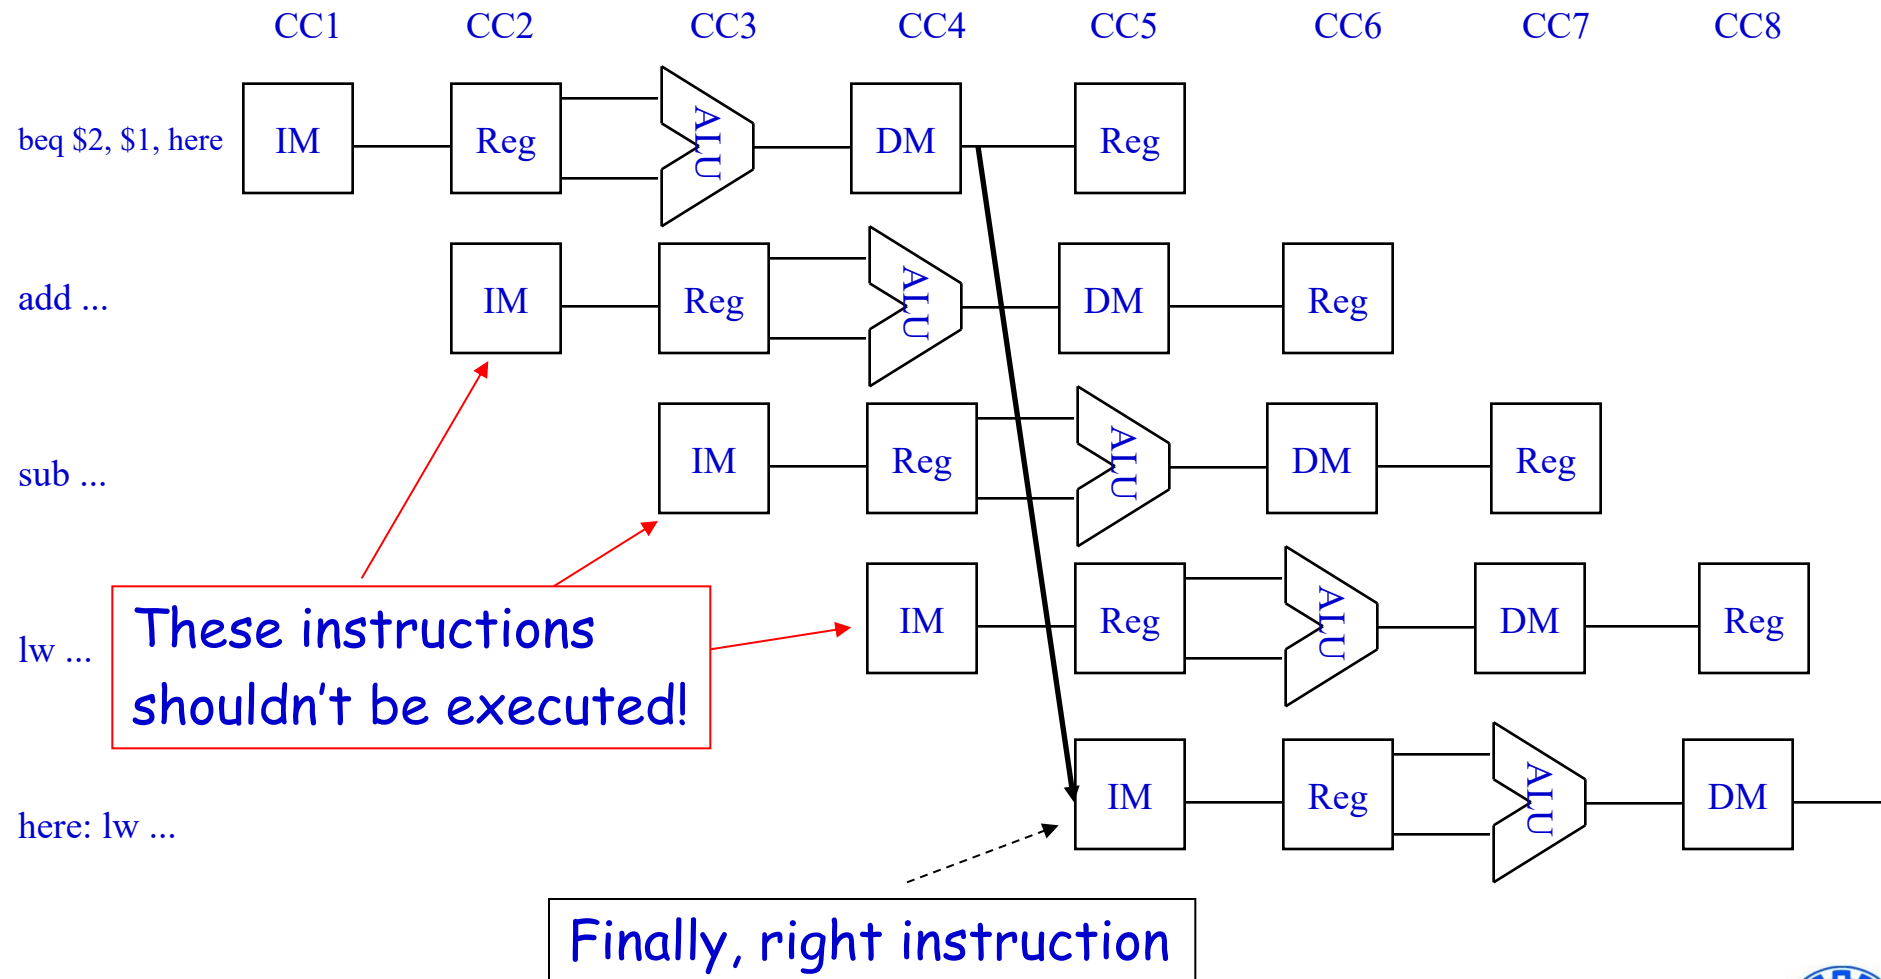
Practice

- Consider the following code:
 - What types of data forwarding are required?

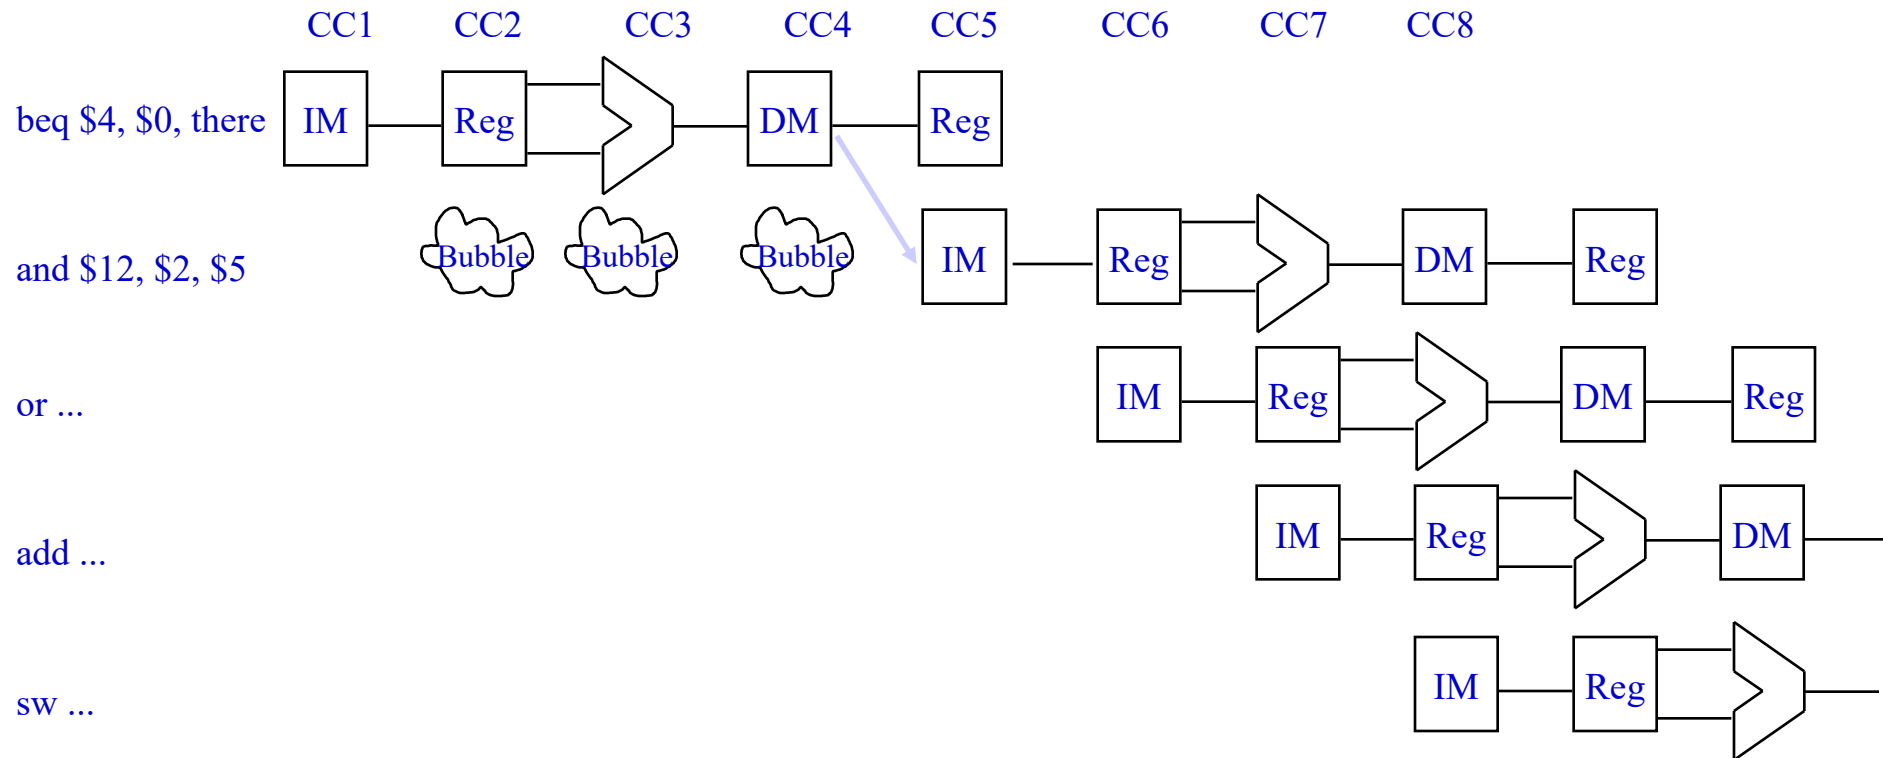
```
sub $2, $1, $3  
and $12, $2, $5  
or $13, $6, $2  
add $14, $2, $2  
sw $15, 100($2)
```



Control or Branch Hazard

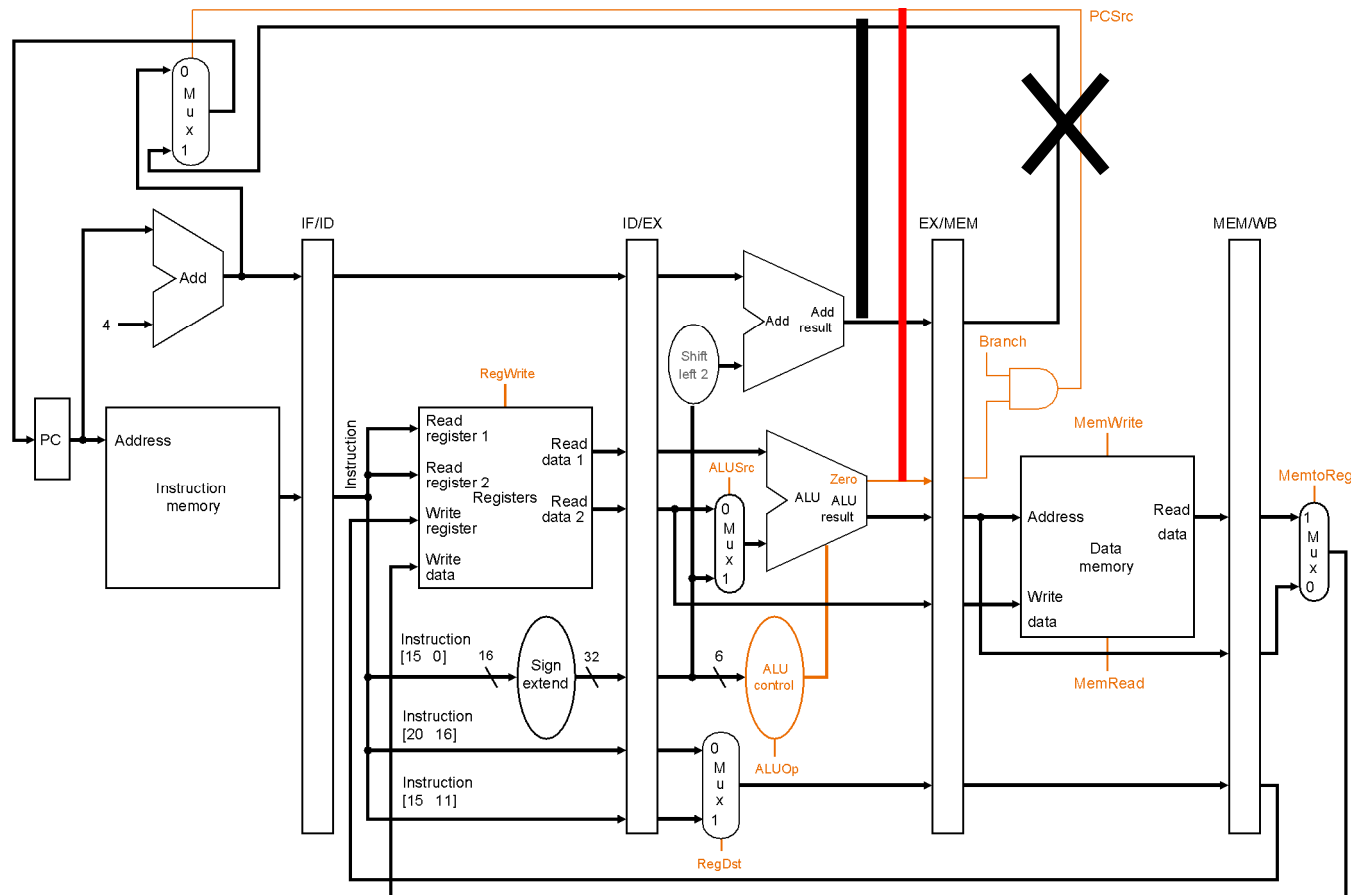


Stalling for Branch Hazards

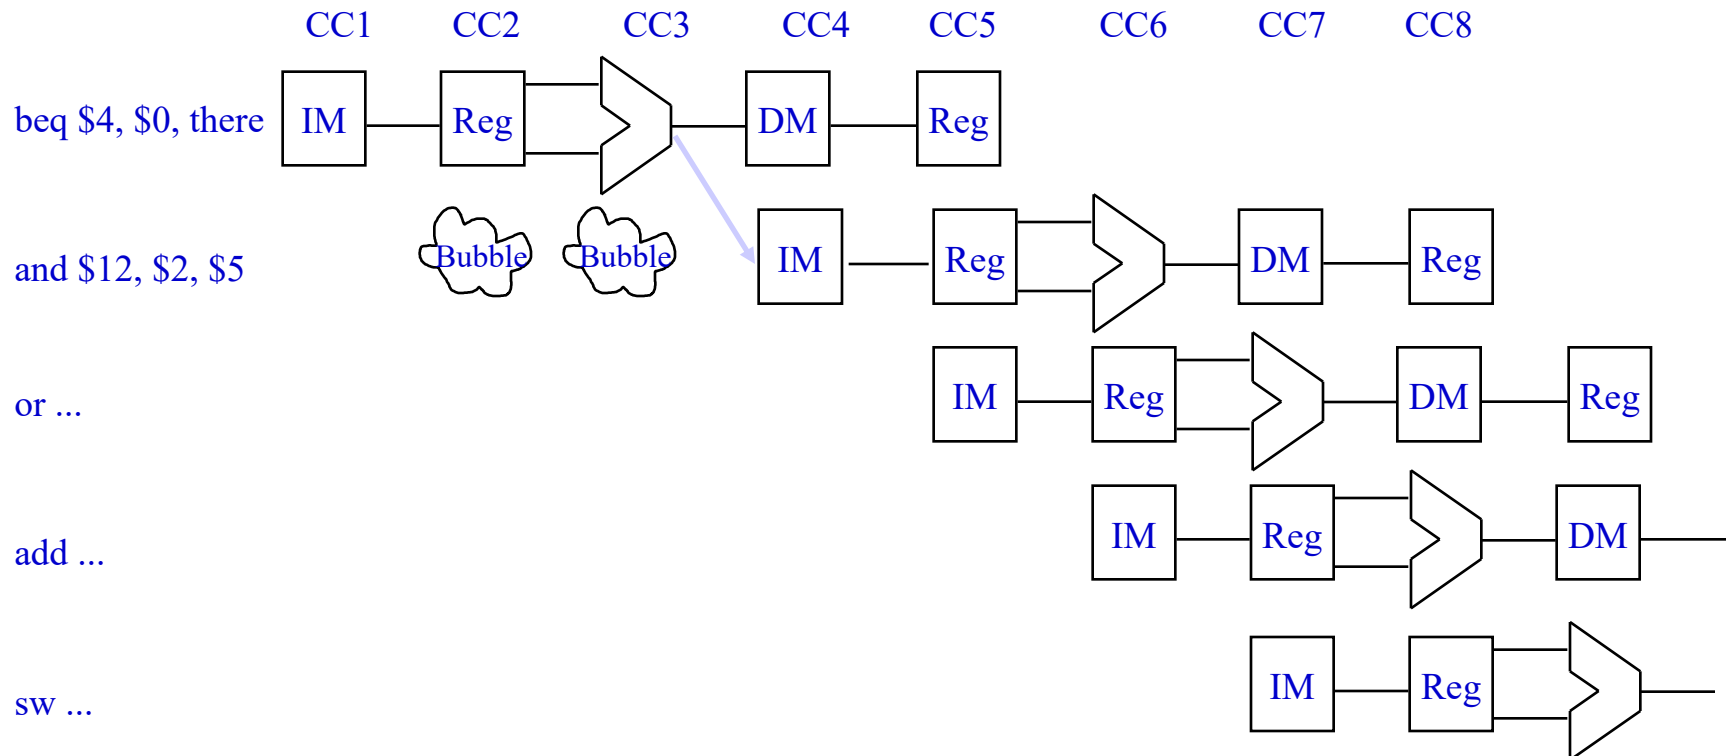


Reducing Branch Delay

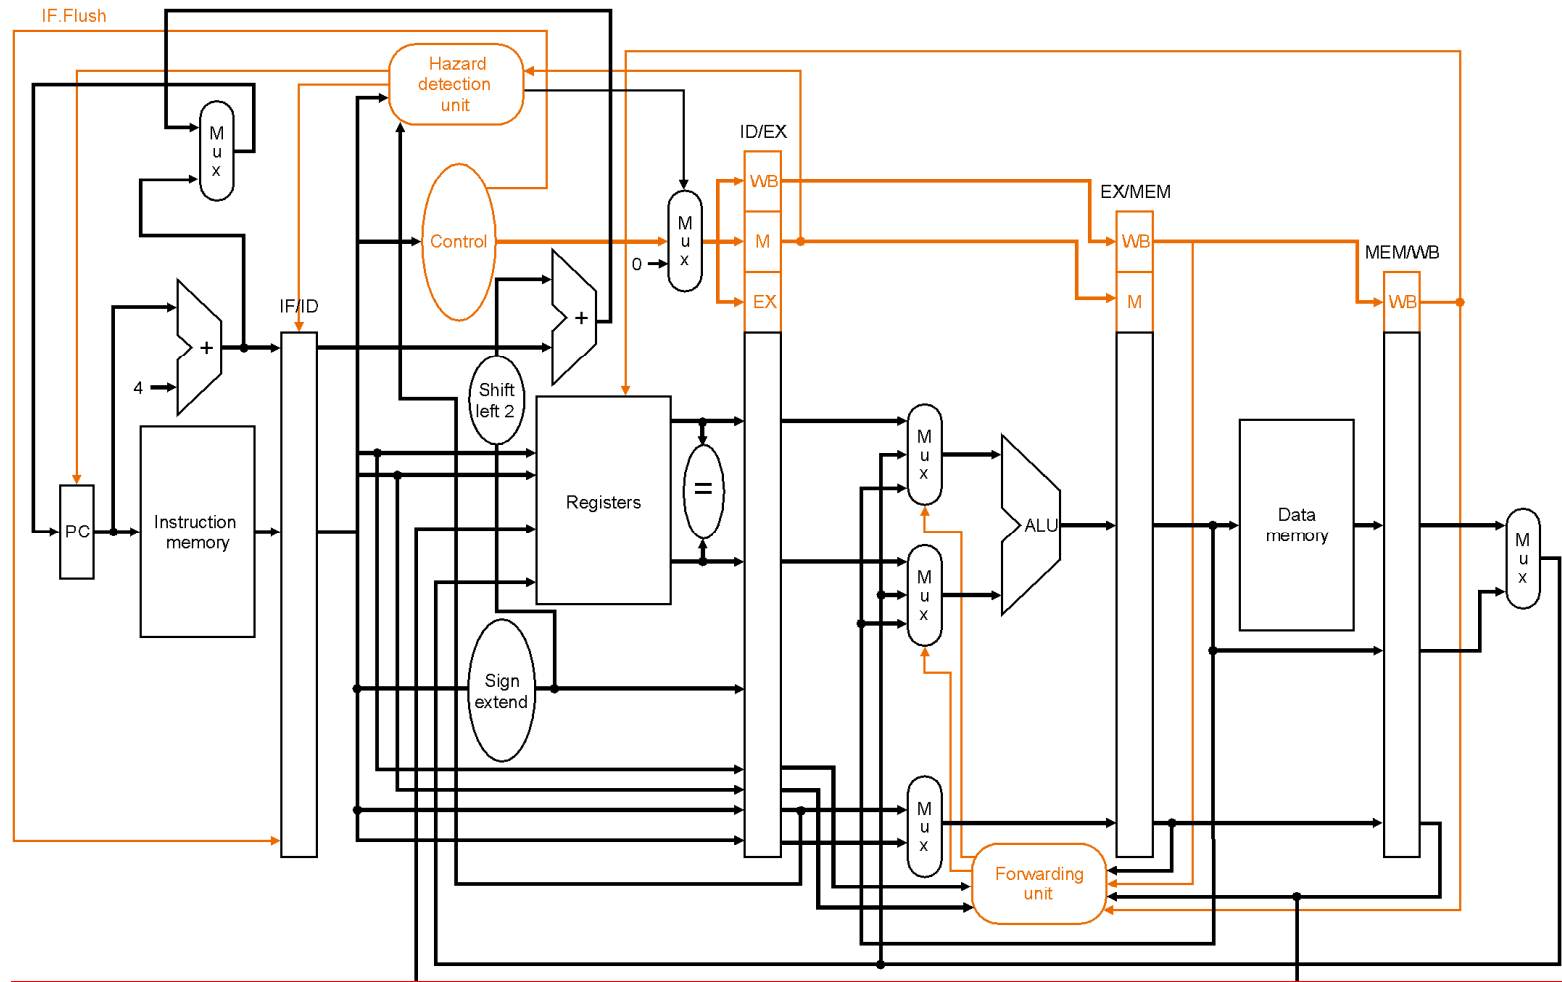
It's easy to reduce stall to 2-cycles



Stalling for Branch Hazards



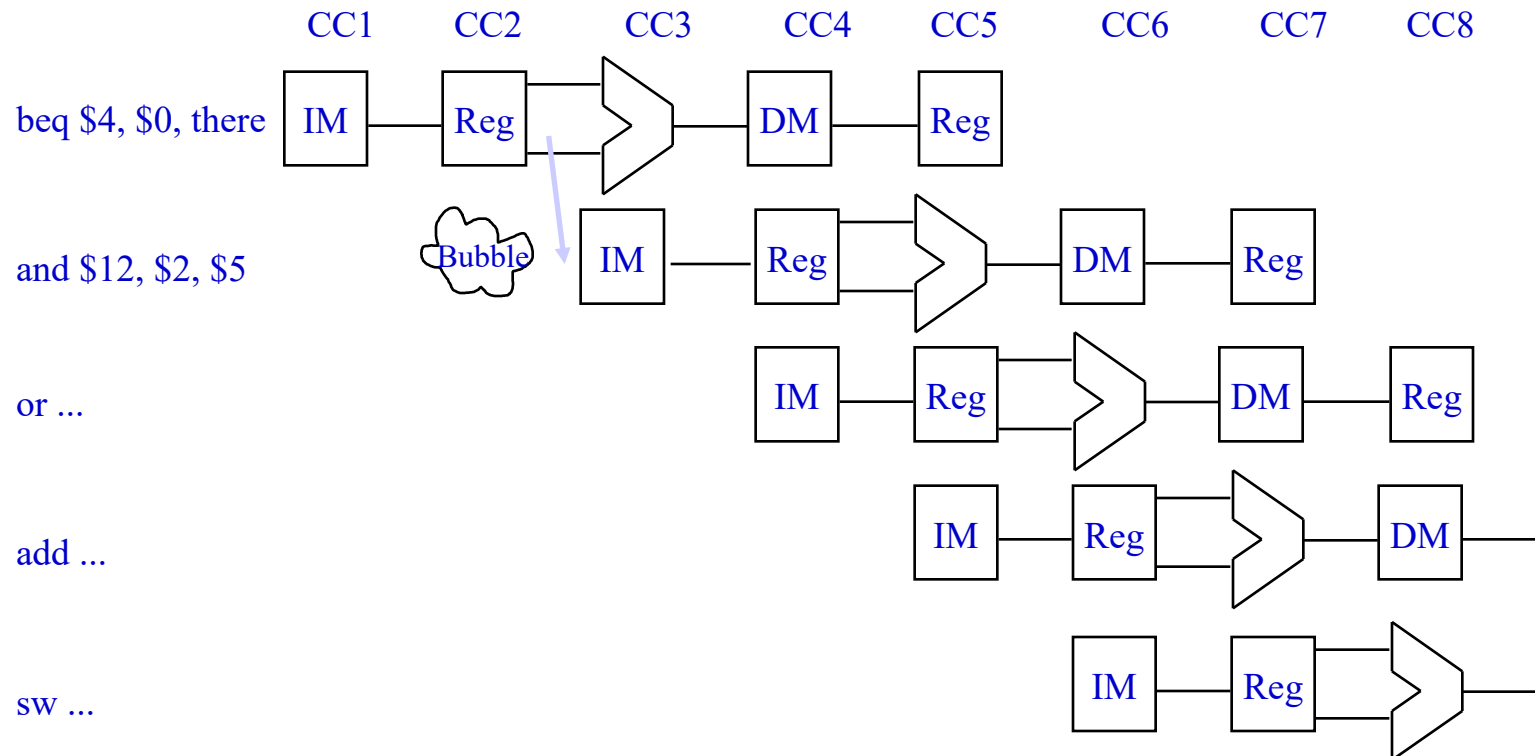
One-Cycle Branch Misprediction Penalty



- Target computation & equality check in ID phase



Stalling for Branch Hazards



Practice

- Which program has more IC?
- Which one has less bubbles?
- Which one runs faster?
- How many clock cycles?

```
    move $s0,$zero
    li $s1, 100
L1: beq $s0, $s1, L2
    add $s2,$s2,$s0
    addi $s0,$s0,1
    j L1
L2:
```

```
    move $s0,$zero
    li $s1, 100
L1: add $s2,$s2,$s0
    addi $s0,$s0,1
    bne $s0, $s1, L1
L2:
```

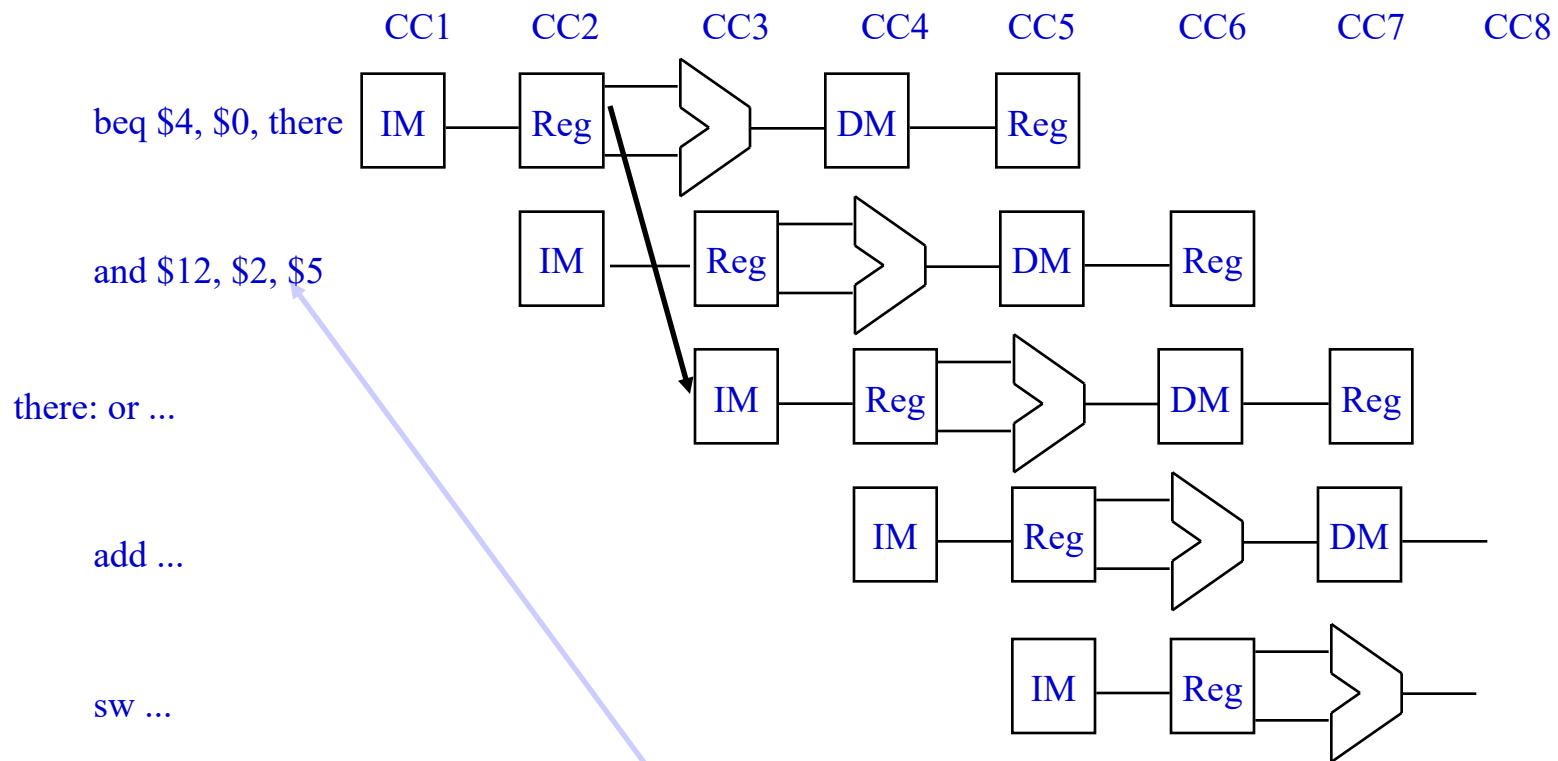


Eliminating Branch Stall

- SPARC and MIPS
 - Use a single *branch delay slot* to eliminate single-cycle stalls after branches
 - Instruction after a conditional branch is *always executed* in those machines
 - Regardless of whether branch is taken or not!



Branch Delay Slot



Branch delay slot instruction (next instruction after a branch) is executed even if the branch is taken.



Filling Branch Delay Slot

add \$5, \$3, \$7

sub \$6, \$1, \$4

and \$7, \$8, \$2

beq \$6, \$7, there

nop /* branch delay slot */

add \$9, \$1, \$2

sub \$2, \$9, \$5

...

there:

mult \$2, \$10, \$11



More-Realistic Branch Prediction

- **Static** Branch Prediction
 - Based on typical branch behavior
 - Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken
- **Dynamic** Branch Prediction
 - Hardware measures actual branch behavior
 - e.g., record recent history of each branch
 - Assume future behavior will continue trend
 - When wrong → stall while re-fetching & update history



Correlated Branches

- Two Types of branch
 - No correlation
 - Simple correlation
 - Complex correlation

Code A

```
for (i=0; i++; i<n)  
    sum+=i;  
  
for (j=0; j++; j<m)  
    sum+=2*j;
```

Code B

```
for (i=0; i++; i<n)  
    for (j=0; j++; j<i)  
        sum+=i+j;
```


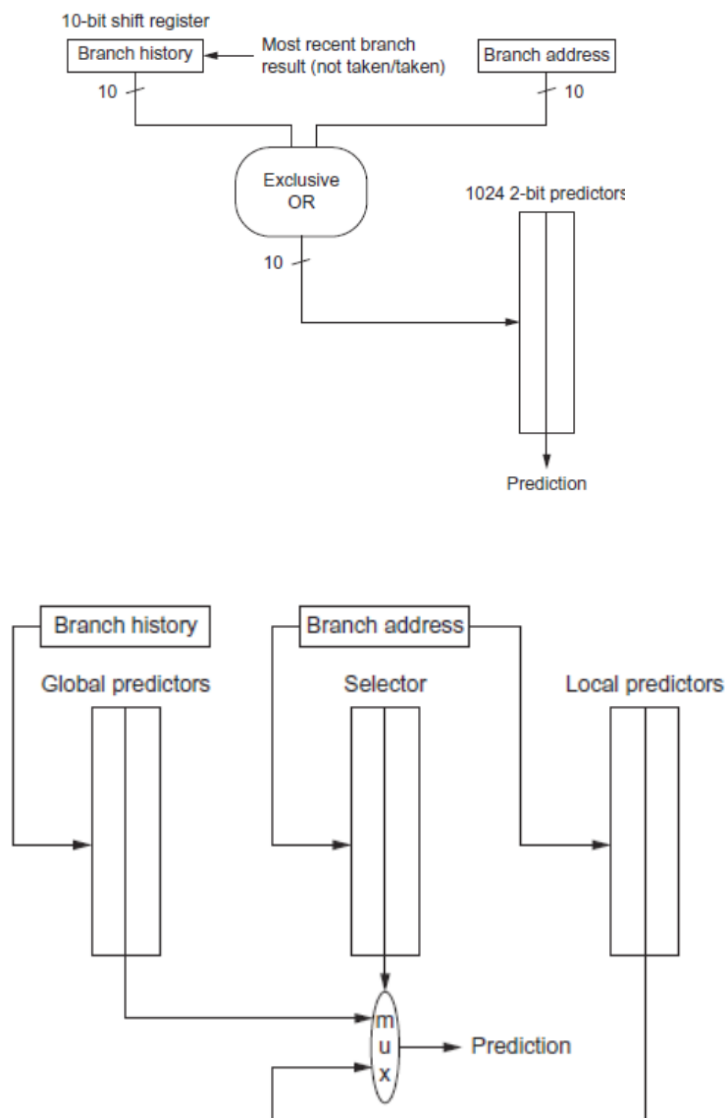
Code C

```
If (aa==2)  
    aa=0;  
If (bb==2)  
    bb=0;  
if (aa!=bb)  
    cc = 1;
```



Branch Prediction

- Different Configurations → Different mispredictions



Compiler Techniques

- Compiler Techniques to Mitigate Branch Misspenalty
 - Loop unrolling
 - Macro (instead of procedure call)

```
for (i=0; i++; i<n)  
  for (j=0; j++; j<2)  
    sum+=i+j;
```





Thanks for Your Attention!

