

Computer Architecture: Cache Design: Part One

Hossein Asadi (asadi@sharif.edu)

Department of Computer Engineering

Sharif University of Technology

Spring 2024



Copyright Notice

- Some Parts (text & figures) of this Lecture adopted from following:
 - D.A. Patterson and J.L. Hennessy, “[Computer Organization and Design: the Hardware/Software Interface](#)” (MIPS), 6th Edition, 2020.
 - J.L. Hennessy and D.A. Patterson, “[Computer Architecture: A Quantitative Approach](#)”, 6th Edition, Nov. 2017.
 - “Intro to Computer Architecture” handouts, by Prof. Hoe, CMU, Spring 2009.
 - “Computer Architecture & Engineering” handouts, by Prof. Kubiawicz, UC Berkeley, Spring 2004.
 - “Intro to Computer Architecture” handouts, by Prof. Hoe, UWisc, Spring 2021.
 - “Computer Arch I” handouts, by Prof. Garzarán, UIUC, Spring 2009.



Topics Covered in This Lecture: I

- **Locality in Memory**
- **Spatial & Temporal Locality**
- **Cache Structure**
- **Cache Configurations**
- **Write-Through vs. Write-Back**
- **Direct-Mapped**
- **Set Associative vs. Full Associative**



Topics Covered in This Lecture: II

- **Replacement Policy**
- **Write Policy: Write-Through vs. Write-Back**
- **Sample Cache Configurations**




Ideal Memory

- **Processors**
 - Would run one instruction per cycle if:
 - Every memory access takes one cycle
 - Every request to memory is successful
- **Our Ideal Memory?**
 - Very large
 - Can be accessed in one clock cycle
- **Reality**
 - Any GB-size memory running at Ghz?



Why Memory a Big Deal?



The Law of Storage

- **Bigger is Slower**
 - FFs, 512 Bytes, sub-nanosec
 - SRAM, KByte~MByte, ~nanosec
 - DRAM, Gigabyte, ~50 nanosec
 - Hard Disk, Terabyte, ~10 millisec
- **Faster is More Expensive (\$ and chip area)**
 - SRAM < 10\$ per Megabyte
 - DRAM, < 1\$ per Megabyte
 - Hard Disk < 1\$ per Gigabyte
- *Note* these sample values scale with time



Question is:

- How to Make Memory?
 - Bigger,
 - Faster, &
 - Cheaper?



Principle of Locality

- **Locality**

- One's recent past is a very good predictor of his/her near future

- **Temporal Locality**

- If you just did something, it is very likely that you will do same thing again **soon**

- **Spatial Locality**

- If you just did something **there**, it is very likely you will do something similar/related **around** again



Locality in Memory

- Locality in Memory
 - A “typical” program has a lot of locality in memory references
 - Programs are sequential and composed of “loops”
- Temporal
 - A program tends to reference **same memory location many times** and all within a small window of time
- Spatial
 - A program tends to **reference a cluster of memory locations** at a time



Locality in Memory (cont.)

- Example 1:
 - This sequence of addresses has both types of locality

1, 2, 3, 1, 2, 3, 8, 8, 47, 9, 10, 8, 8 ...

spatial temporal non-local



Locality in Memory (cont.)


- Example 2:
 - Data
 - Reference array elements in succession (spatial)
 - Reference to sum & i (temporal)
 - Instructions
 - Reference instructions in sequence (spatial)
 - Cycle through loop repeatedly (temporal)

```
sum = 0;  
for (i = 0; i < n; i++)  
    sum += a[i];  
*v = sum;
```



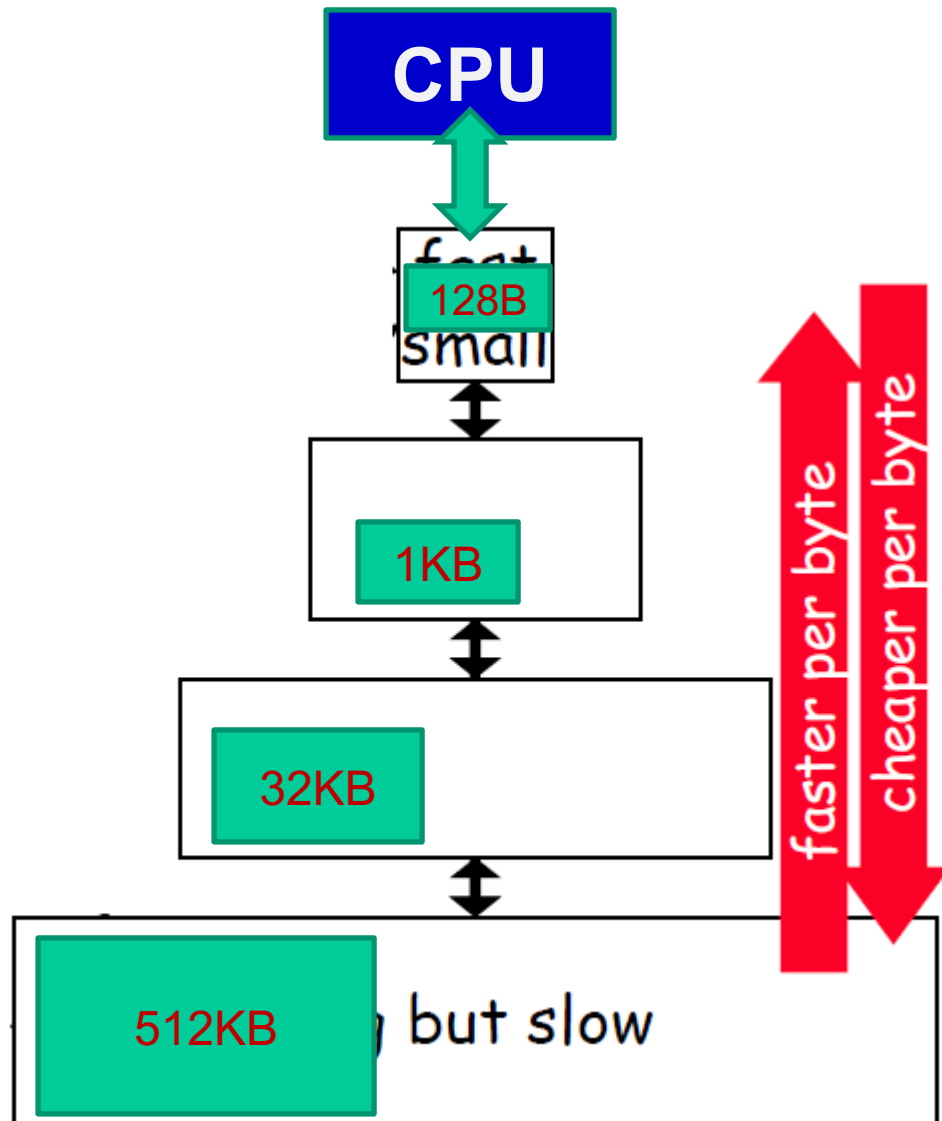
Probability of Reference

Probability
of reference

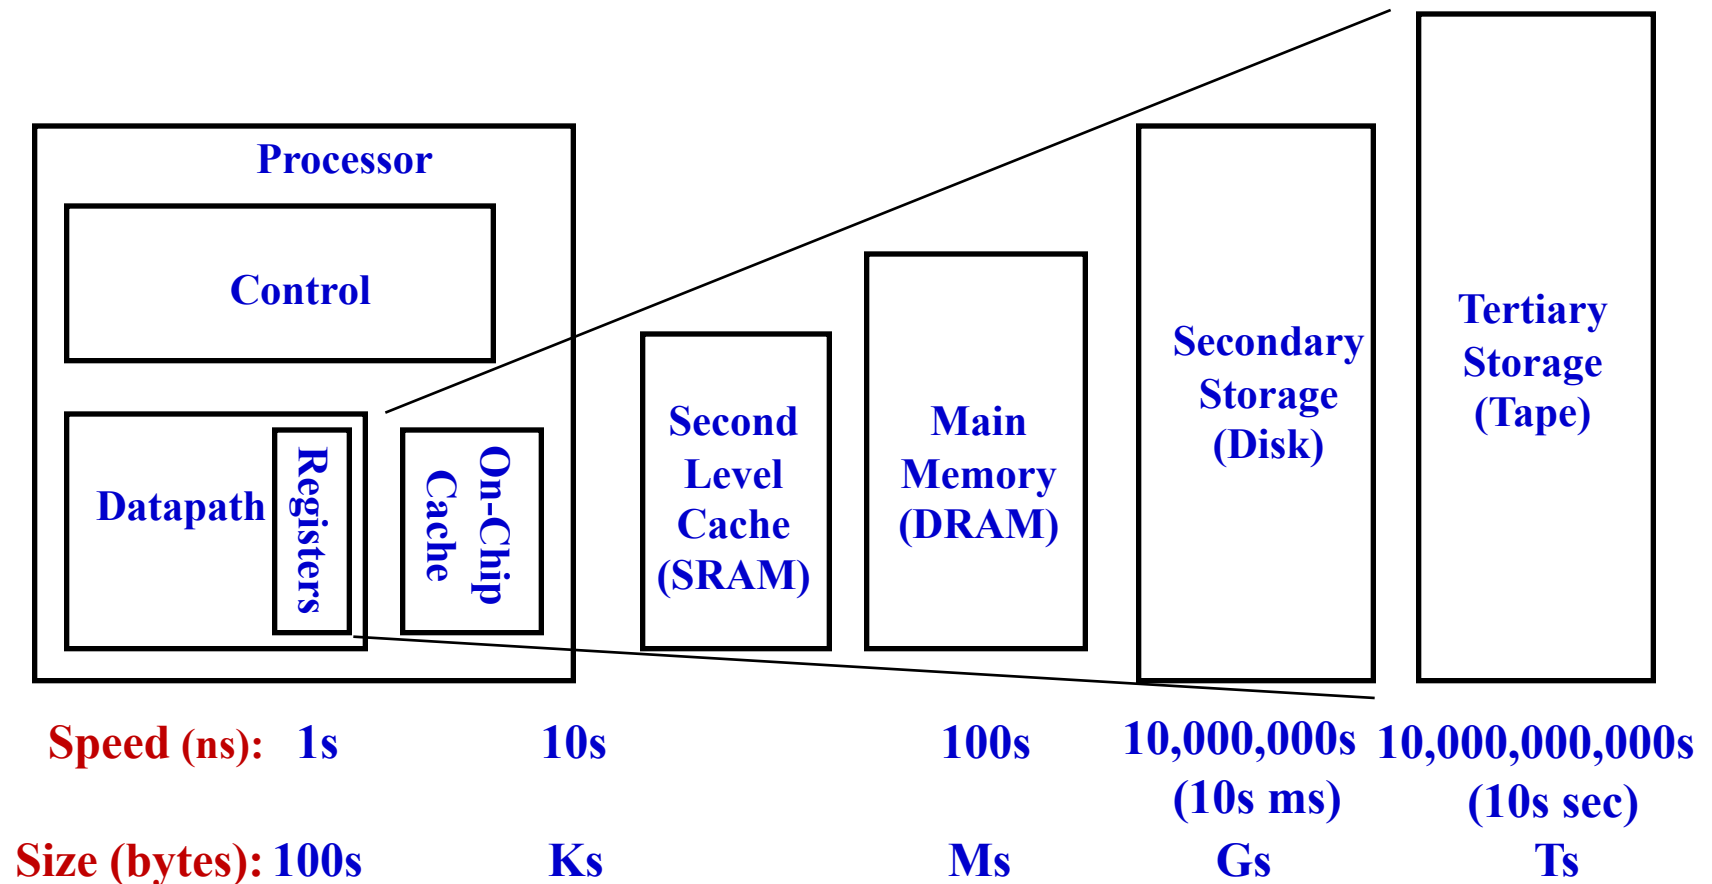


Memory Hierarchy

- Faster & Small
- Bigger & Slower



Memory Hierarchy (cont.)

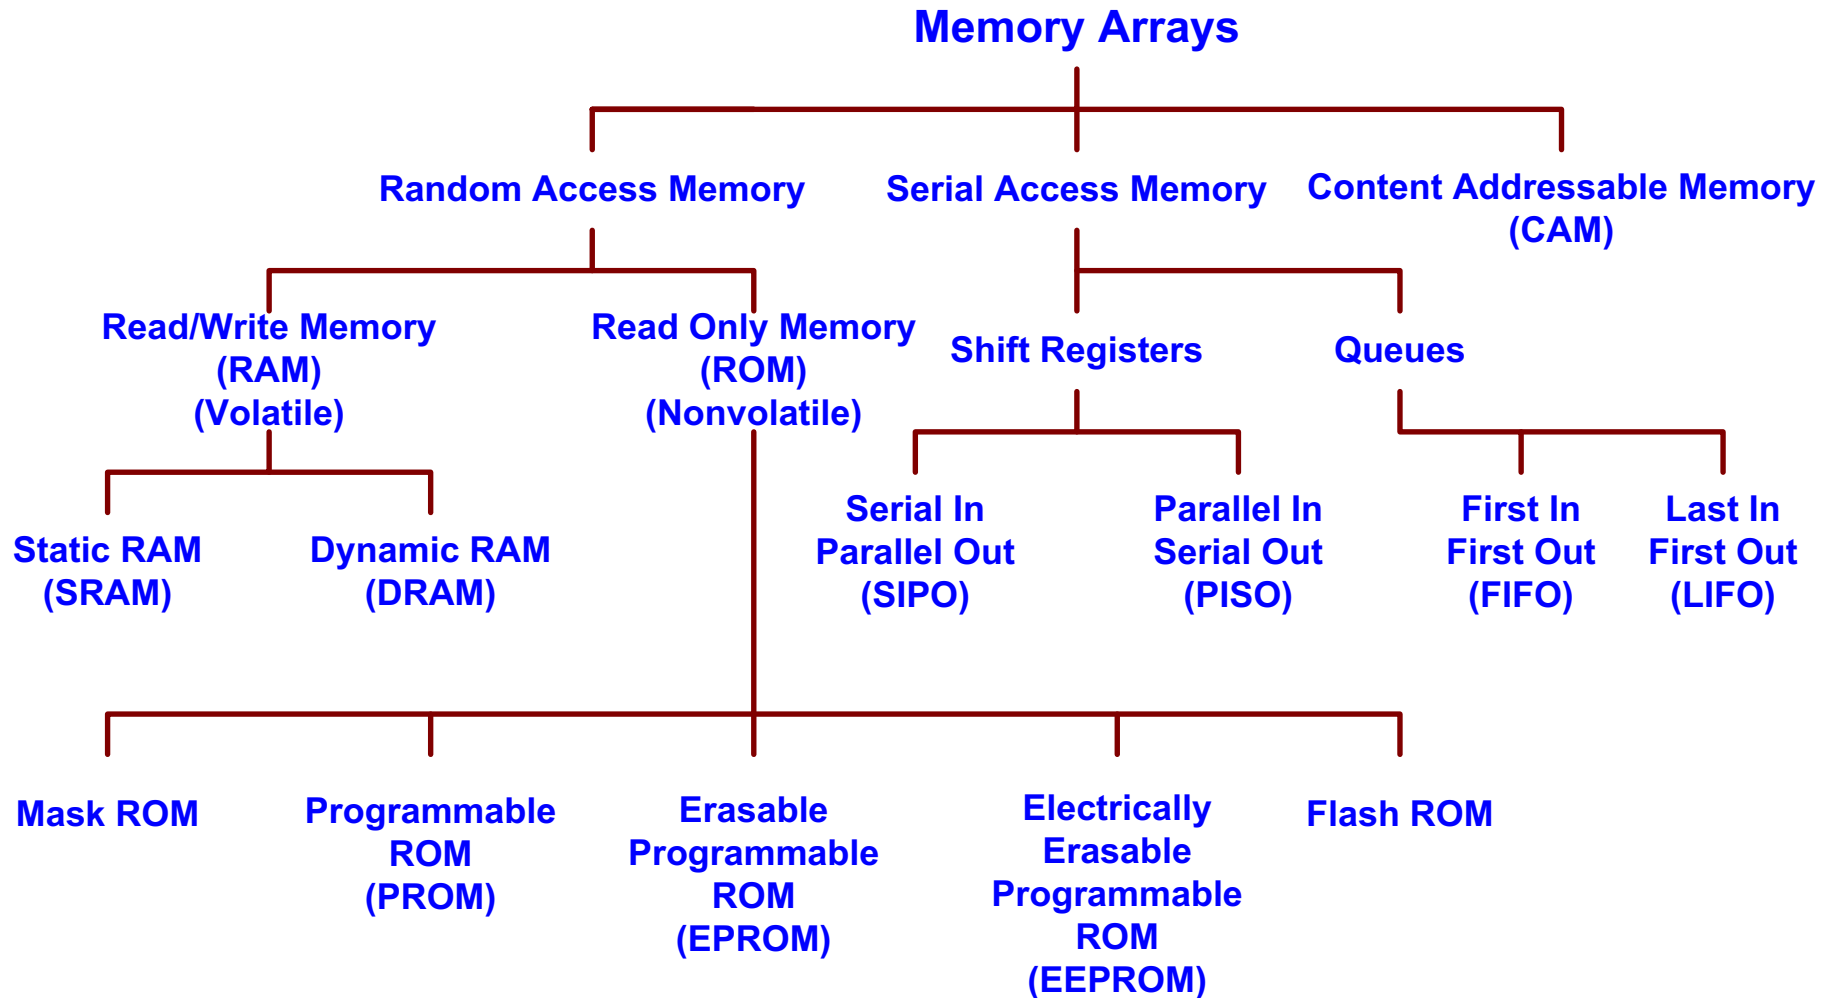


Technology Used in Main Memory & Cache Memory

- **SRAM** (Static Random Access Memory)
 - No refresh (6 transistors/bit vs. 1 transistor)
 - # of transistors per bit: DRAM < SRAM
 - Cycle time: DRAM > SRAM
- **DRAM** (Dynamic Random Access Memory)
 - Dynamic since needs to be refreshed periodically
 - Addresses divided into 2 halves
 - Memory as a 2D matrix
 - RAS or Row Address Strobe
 - CAS or Column Address Strobe

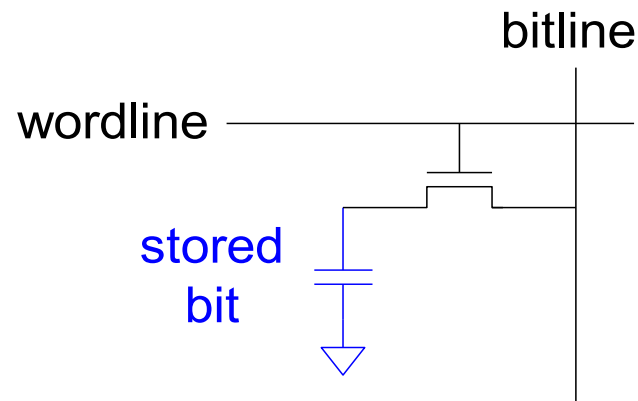



Memory Arrays



SRAM Cell vs. DRAM Cell

6-Transistor SRAM Cell



Slow Memory in Pipeline Datapath

- Freeze pipeline in *Mem* stage:

IF0	ID0	EX0	Mem0	Wr0	Noop	...	Noop	Noop
	IF1	ID1	EX1	Mem1	stall...	stall	Mem1	
	Wr1							
		IF2	ID2	EX2	stall...	stall	Ex2	
	Mem2	Wr2						
			IF3	ID3	stall...	stall	ID3	
	Ex3	Mem3	Wr3					
				IF4	stall...	stall	IF4	
	ID4	Ex4	Mem4	Wr4				
								IF5
	ID5	Ex5	Mem5					

- Stall detected by end of Mem1 stage



CPU Performance

- **CPU Time** = (CPU execution clock cycles + memory stall clock cycles) x clock cycle time
- Memory **Stall** Clock Cycles =
(reads x read miss rate x read miss penalty +
writes x write miss rate x write miss penalty)
- Memory **Stall** Clock Cycles =
Memory accesses x Miss rate x Miss penalty



CPU Performance (cont.)

- Different Measure:
 - Average Memory Access Time (AMAT)
- Expressed in Terms of:
 - Hit time
 - Miss rate
 - Miss penalty



CPU Performance (cont.)

- **Hit Time**
 - Time required to **access a level of memory hierarchy** including time required to determine whether **access** is **hit or miss**
- **Hit Rate** (Hit Ratio)
 - Fraction of memory accesses found in a cache
 - **Miss Rate** = $1 - \text{Hit Rate}$



CPU Performance (cont.)

- **Miss Penalty:**
 - Time required to fetch a block into a level of memory hierarchy from lower level:
 - Time to **access block** +
 - Time to **transmit** it to higher level +
 - Time to **insert** it in appropriate **block**
- **Block**
 - **Minimum unit** of information transferred between two levels of memory hierarchy
 - Also called **line**
- **AMAT** = Hit Time +
(Miss Rate x Miss Penalty)



CPU Performance (cont.)

- Example 1
 - A memory system consists of a cache and a main memory
 - Cache hit = 1 cycle
 - Cache miss = 100 cycles
 - What is average memory access time if hit rate in cache is 97%?



CPU Performance (cont.)

- Example 2
 - A memory system has a cache, a main memory, and a **virtual** memory
 - Hit rate = 98%
 - Hit rate in main memory = 99%
 - 2 cycles to access cache
 - 150 cycles to fetch a line from main mem.
 - 100,000 cycles to access virtual memory
 - What is average memory access time?



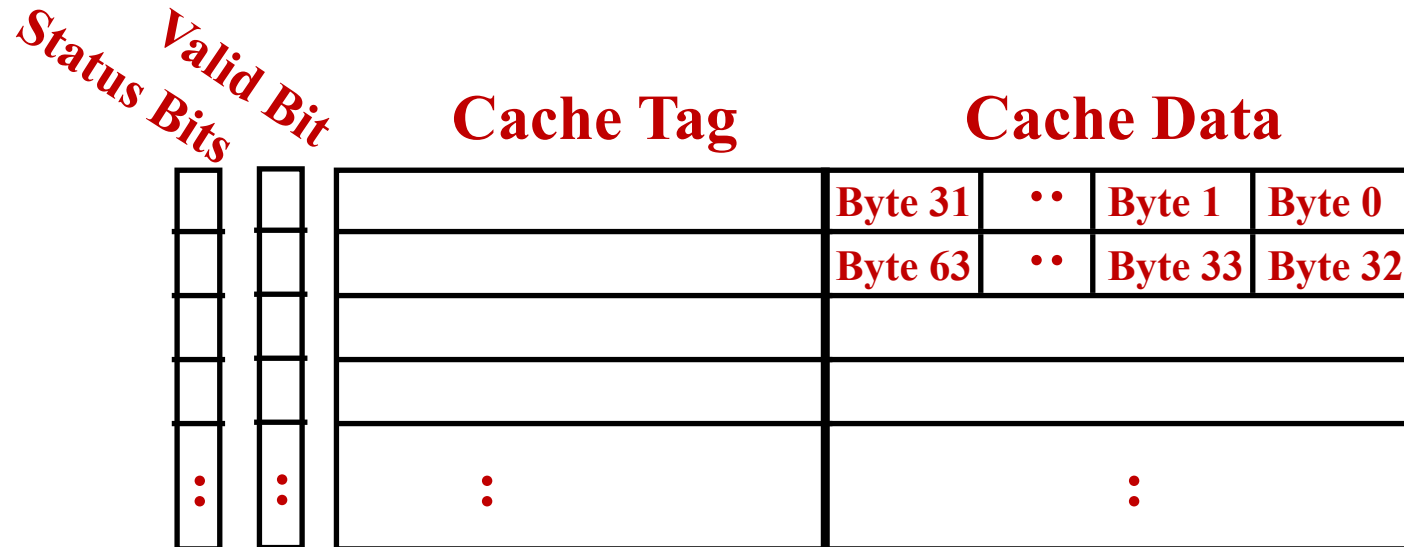
Improving Cache Performance

- AMAT =
Hit Time + (Miss Rate x Miss Penalty)
- Options to Reduce AMAT
 - Reduce time to hit in cache
 - Use smaller cache size
 - Reduce miss rate
 - Increase cache size!
 - Reduce miss penalty
 - Use multi-level cache hierarchy



Cache Structure

- Data Bits
- Tag Bits
- Invalid Bit
- Status Bits (LRU bit, Dirty Bit, ...)



Cache Configuration

- Q1:
 - Where can a block be placed in upper level?
 - Block placement
- Q2:
 - How is a block found if it is in upper level?
 - Block identification




Cache Configuration (cont.)

- Q3:
 - Which block should be replaced on a miss?
 - **Block replacement**
- Q4:
 - What happens on a write?
 - How to propagate changes?
 - **Write strategy**




Q1: Where can a block be placed in upper level?


Fully associative:
block 12 can go
anywhere



Direct mapped:
block 12 can go
only into block 4
($12 \bmod 8$)




Set associative:
block 12 can go
anywhere in set0
($12 \bmod 4$)

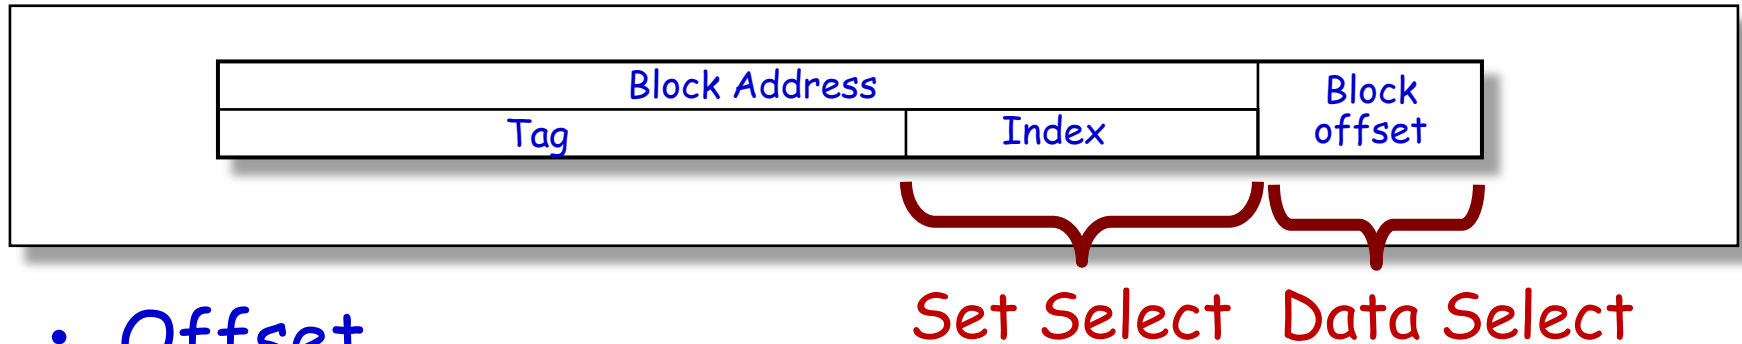


Set 0 Set 1 Set 2 Set 3
0 1 2 3

Block-frame address



Q2: How is a block found if it is in upper level?




- Offset
 - Identifies a byte/word within a block
- Index
 - Identifies corresponding set
- Tag
 - Identifies whether associated block corresponds to a requested word or not



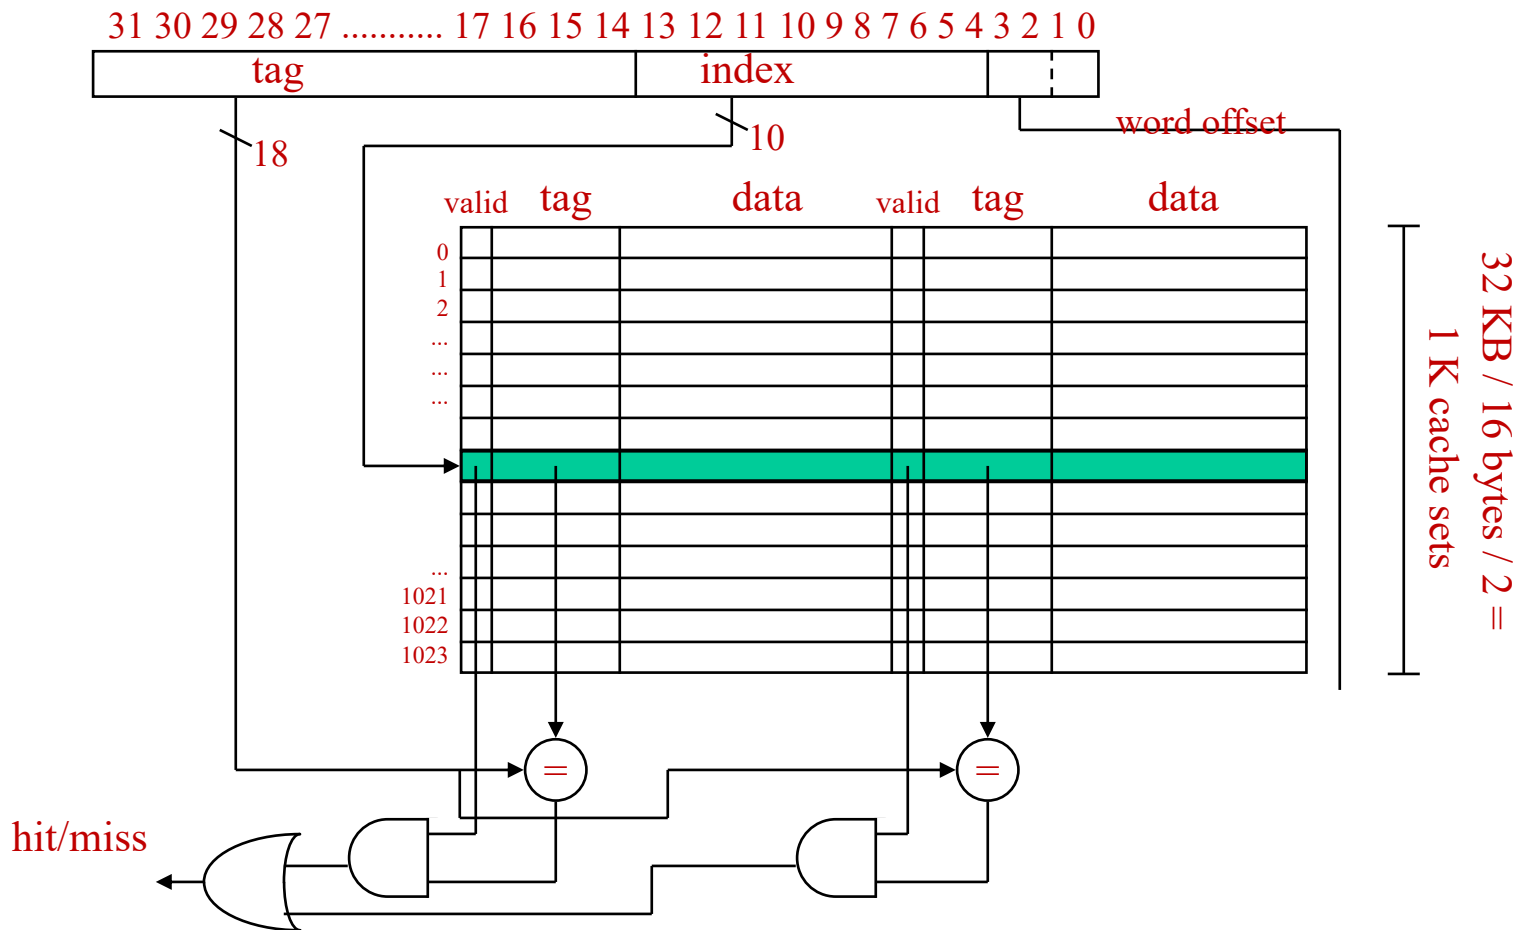
Direct-Mapped

64 KB cache, 32-byte cache block



Set-Associative

32 KB cache, 2-way, 16-byte blocks



Cache Parameters

- Cache size =
 $\# \text{ of sets} * \text{block size} * \text{associativity}$
- Example 1
 - 128 blocks, 32-byte blocks, direct mapped, size = ?
- Example 2
 - 128 KB cache, 64-byte blocks, 512 sets, associativity = ?



Direct-Mapped vs. Fully-Associative

- Direct-Mapped
 - Requires **less area**
 - Only one comparator
 - Fewer tag bits required
 - **Fast hit time**
 - Less bits to compare
 - Cache **block is available BEFORE Hit/Miss**
 - Return data to CPU in parallel with hit process
 - Possible to assume a hit and continue recover later if miss
 - Conflict misses **reduce hit rate**



Direct-Mapped vs. Fully-Associative (cont.)

- Fully-Associative
 - More area compared to direct-mapped
 - Need one comparator for each line in cache
 - Longer hit time
 - Too many comparison required
 - Less miss rate vs. direct-mapped
 - No conflict misses (conflict Miss = 0)
 - No cache index
 - Compare tag with all tags of all cache entries in parallel



Q3: Which block should be replaced on a miss?

- Easy for Direct Mapped; why?
- Set Associative or Fully Associative:
 - Random
 - LRU (Least Recently Used): in status bits
 - FIFO

Associativity:

Data Cache Misses Per 1000 Instructions

Two-way				Four-way			Eight-way		
Size	LRU	Random	FIFO	LRU	Random	FIFO	LRU	Random	FIFO
16 KiB	114.1	117.3	115.5	111.7	115.1	113.3	109.0	111.8	110.4
64 KiB	103.4	104.3	103.9	102.4	102.3	103.1	99.7	100.5	100.3
256 KiB	92.2	92.1	92.5	92.1	92.1	92.5	92.1	92.1	92.5

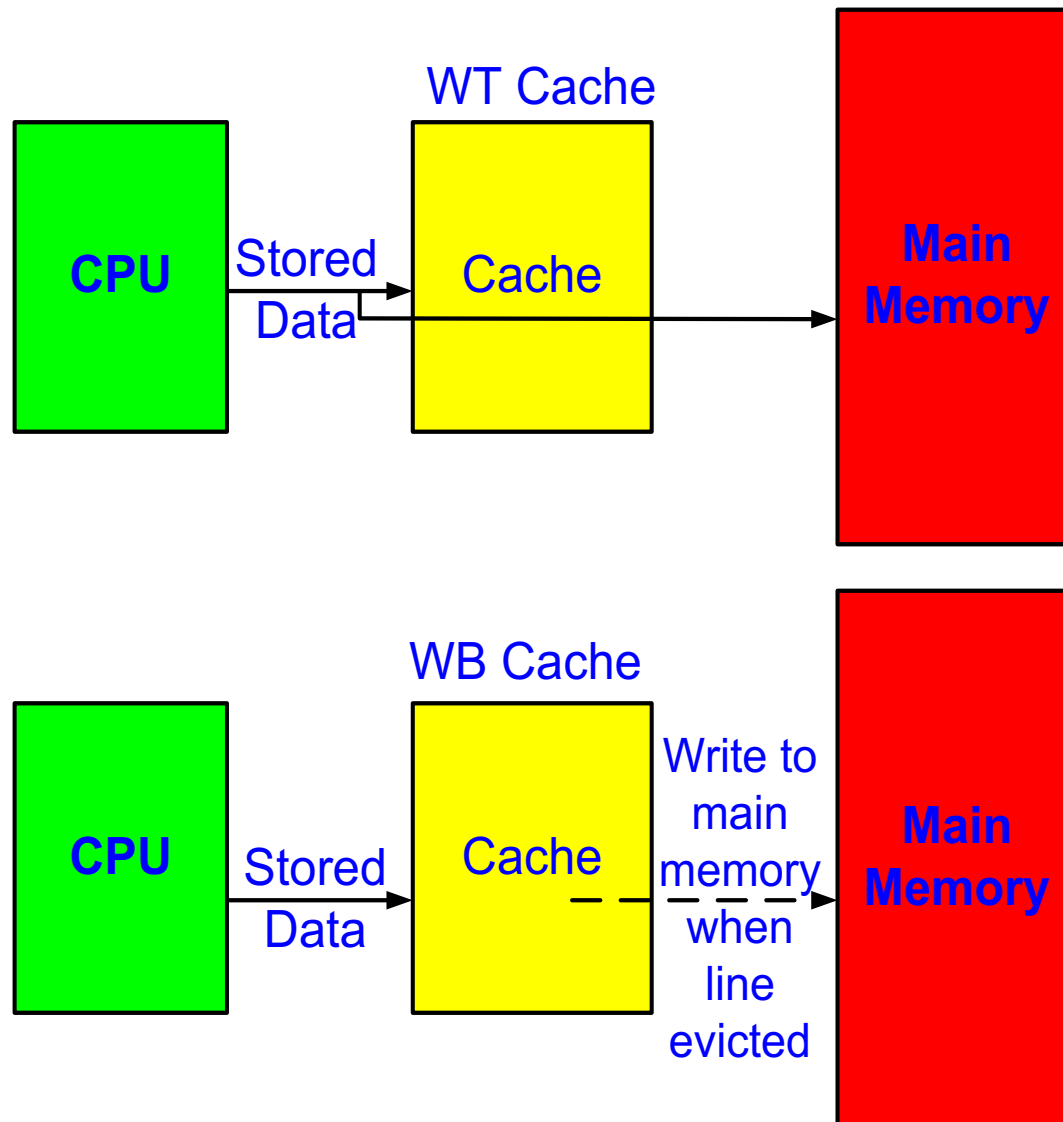


Q4: What happens on a write?

- **Write Through (Allocate/Non-Allocate)**
 - Information written to both block in cache and to block in lower-level memory
- **Write Back**
 - Information written only to block in cache
 - Modified cache block written to main memory only when it is replaced
 - **Inconsistent**
 - Need cache **coherency** policy for multi-core chips
 - Is block clean or dirty? (status bits)



Write-Through (WT) vs. Write-Back (WB)



WT Cache

- **Pros**

- Simpler to implement
- Don't need dirty bit
- No interface issues with I/O devices
 - Cache memory consistent with memory

- **Cons**

- Less performance vs. WB cache
- Processor held up on writes unless writes buffered

- **Write Buffer**

- Stores data while waiting to be written to memory



WB Cache

- **Pros**

- Tends to have better performance
 - Repeated writes not sent to DRAM
 - Processor not held up on writes
 - Combines multiple writes into one line WB
- Virtual memory systems use write-back
 - Because of huge penalty for going out to disk



WB Cache (cont.)

- Cons

- More complex
 - Read miss may require writeback of dirty data
- Need to implement cache coherency
- Typically requires two cycles on writes
 - Can't overwrite data and do tag comparison at same time as block may be **dirty**
 - Unless using **store buffer**



Write Policy in WT Caches

- **Allocate-on-Write (Write Allocate)**
 - Fetch line into cache
 - Then perform write in cache
 - Also called, fetch-on-miss, fetch-on-write
- **No-Allocate-on-Write (No-Write Allocate)**
 - Pass write through to main memory
 - Don't bring line into cache
 - Also called Write-Around or Read-Only



Write Policy in WT Caches

- **Allocate-on-Write Pros**
 - Better performance if data referenced again before it is evicted
- **No-Allocate-on-Write Pros**
 - Simpler write hardware
 - May be better for small caches if written data won't be read again soon



L1 Cache Configuration

- **Split Cache**

- Two independent caches
 - Instruction cache (IL1)
 - Data cache (DL1)

- **Unified**

- One unified L1 cache
- Usually better hit ratio (same size); why?

- **Question:**

- Most processors use split caches; why?



Practice

- Consider a 16KB Cache
 - 4-way, 32-bit address, byte-addressable memory, 32-byte cache blocks
- Q1:
 - How many tag bits?
 - Total tag bits in cache?
- Q2:
 - Where to find word with address = 0x200356A4?



Direct-Mapped

Address

01101

Cache

V d tag data

0				
0				
0				
0				

Block Offset (1-bit)

Line Index (2-bit)

Tag (2-bit)

Memory

00000	78	23
00010	29	218
00100	120	10
00110	123	44
01000	71	16
01010	150	141
01100	162	28
01110	173	214
10000	18	33
10010	21	98
10100	33	181
10110	28	129
11000	19	119
11010	200	42
11100	210	66
11110	225	74



2-Way Set Associative

Address

01101

Cache

V d tag data

0				
0				
0				
0				

Block Offset (unchanged)

1-bit Set Index

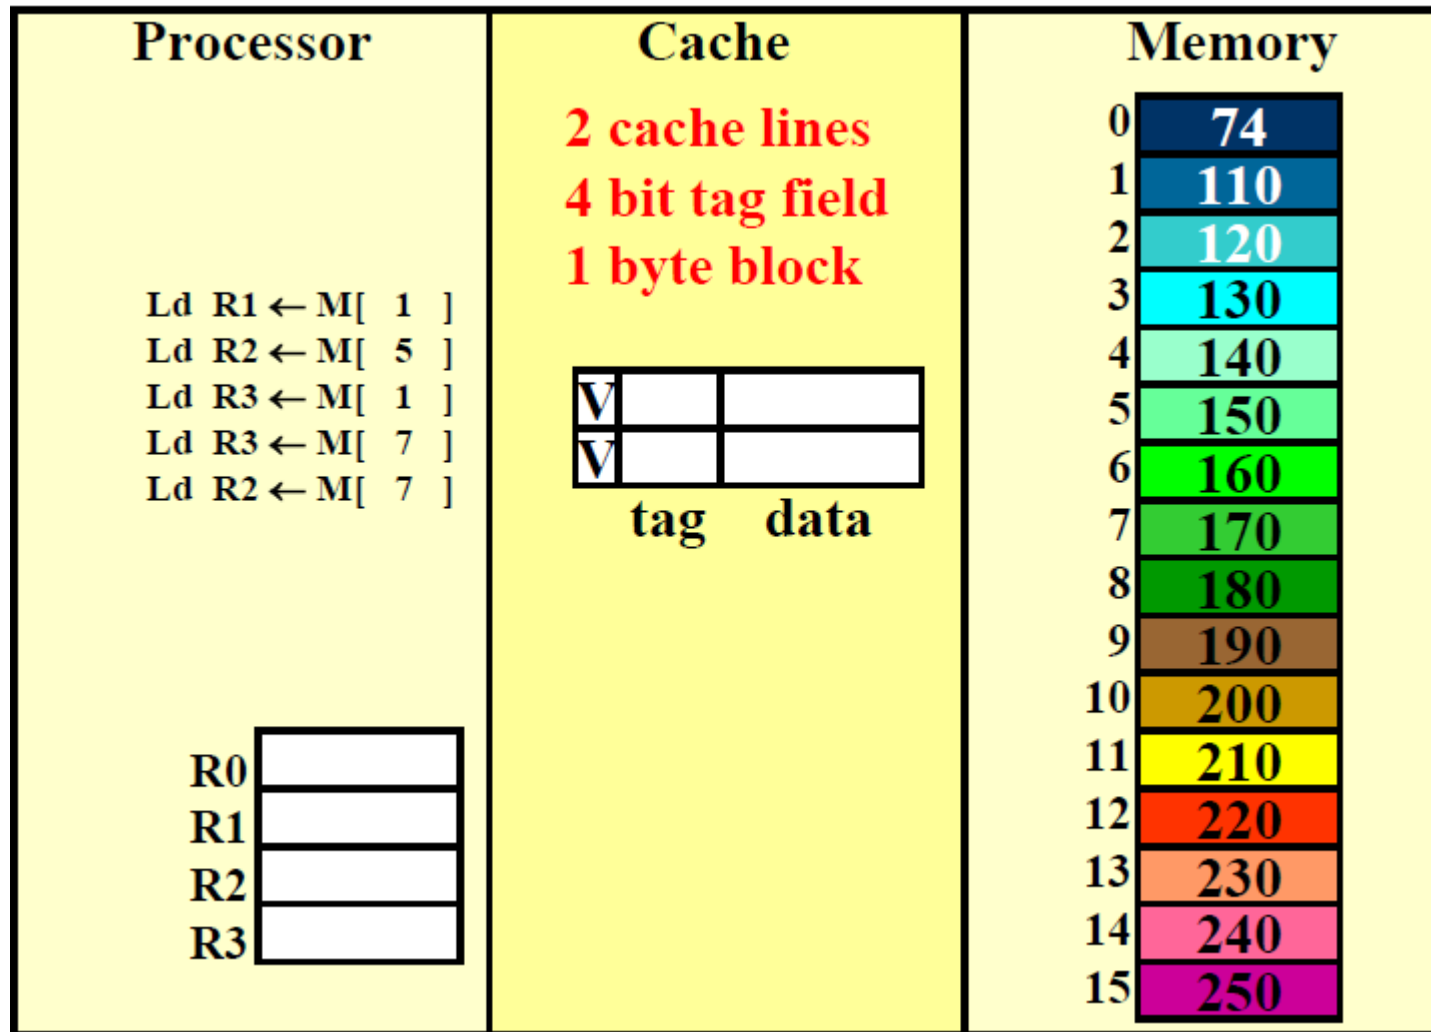
Larger (3-bit) Tag

Memory

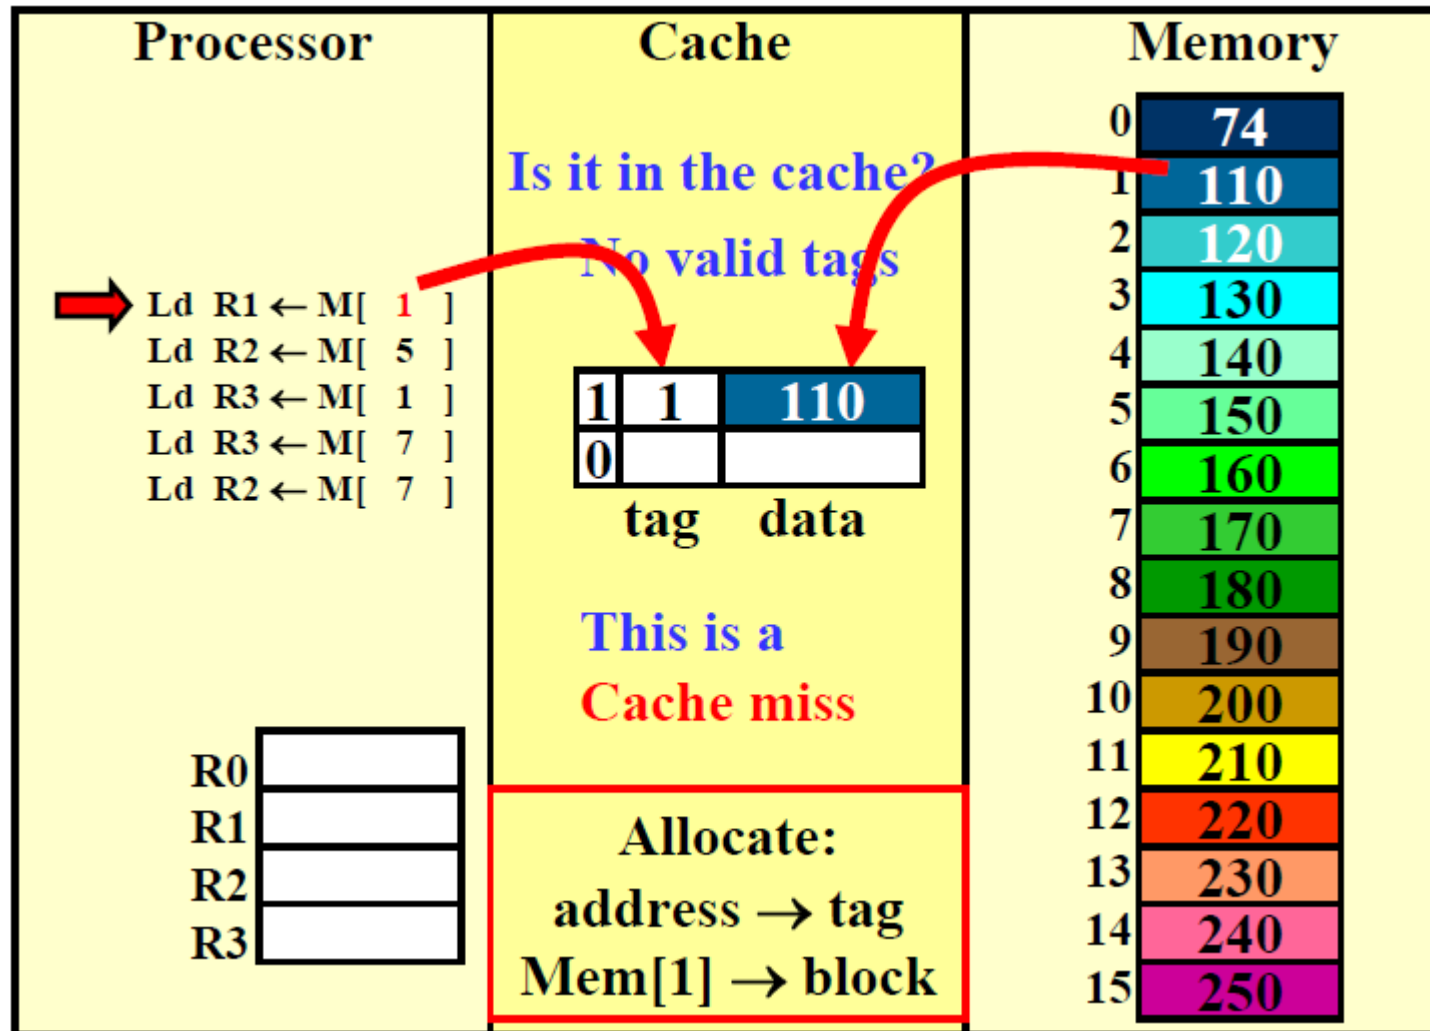
00000	78	23
00010	29	218
00100	120	10
00110	123	44
01000	71	16
01010	150	141
01100	162	28
01110	173	214
10000	18	33
10010	21	98
10100	33	181
10110	28	129
11000	19	119
11010	200	42
11100	210	66
11110	225	74



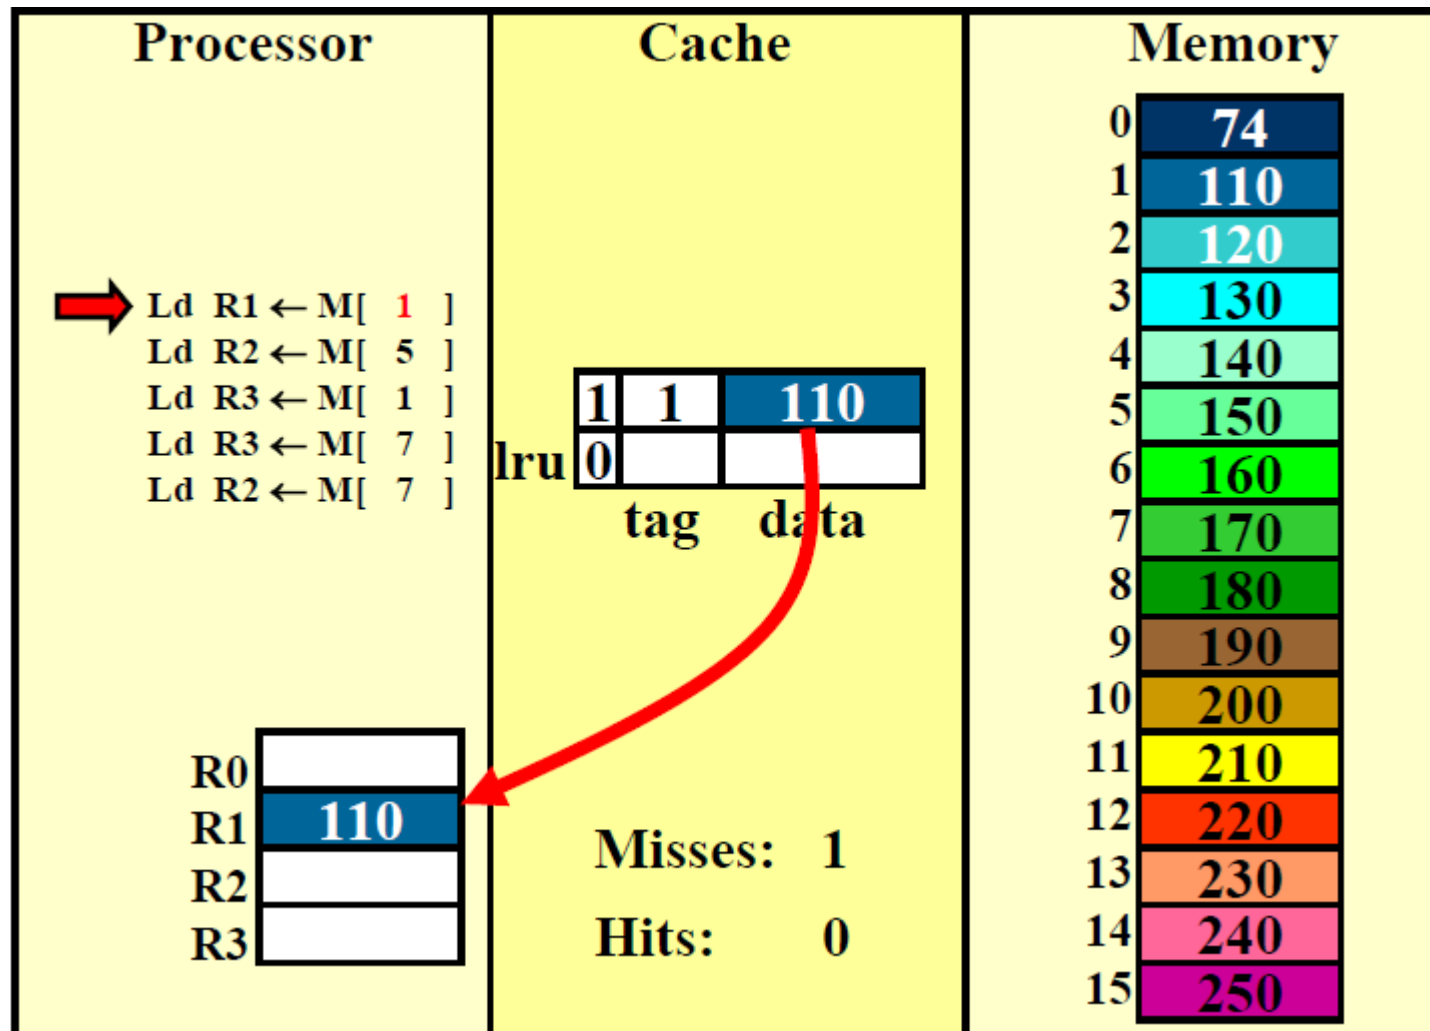
Simple Memory System



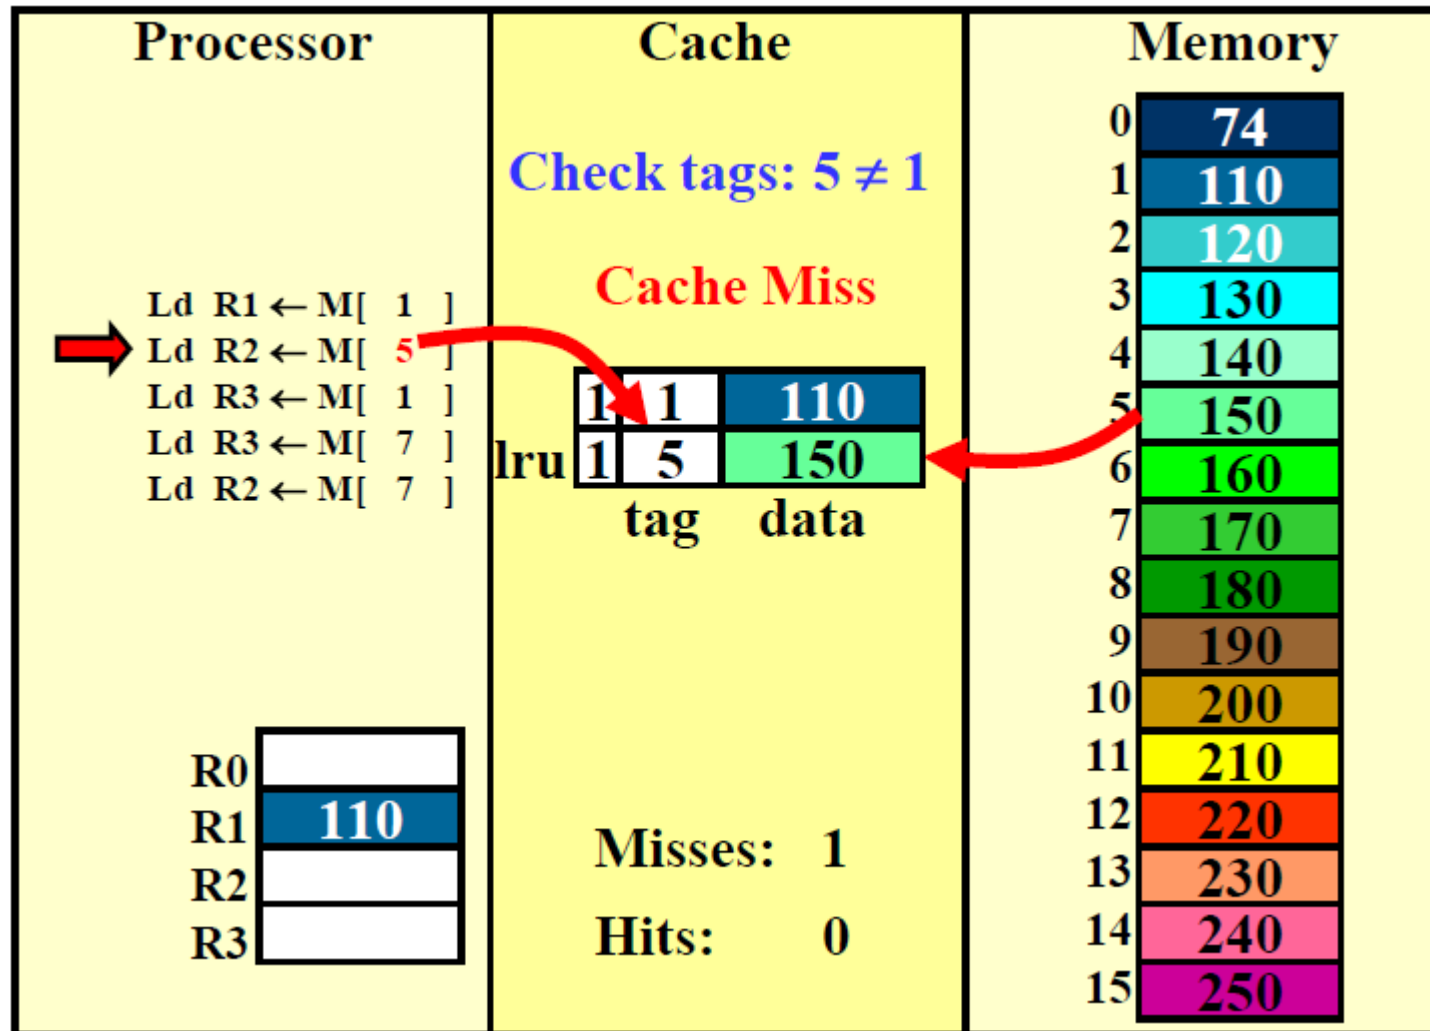
Simple Memory System




Simple Memory System



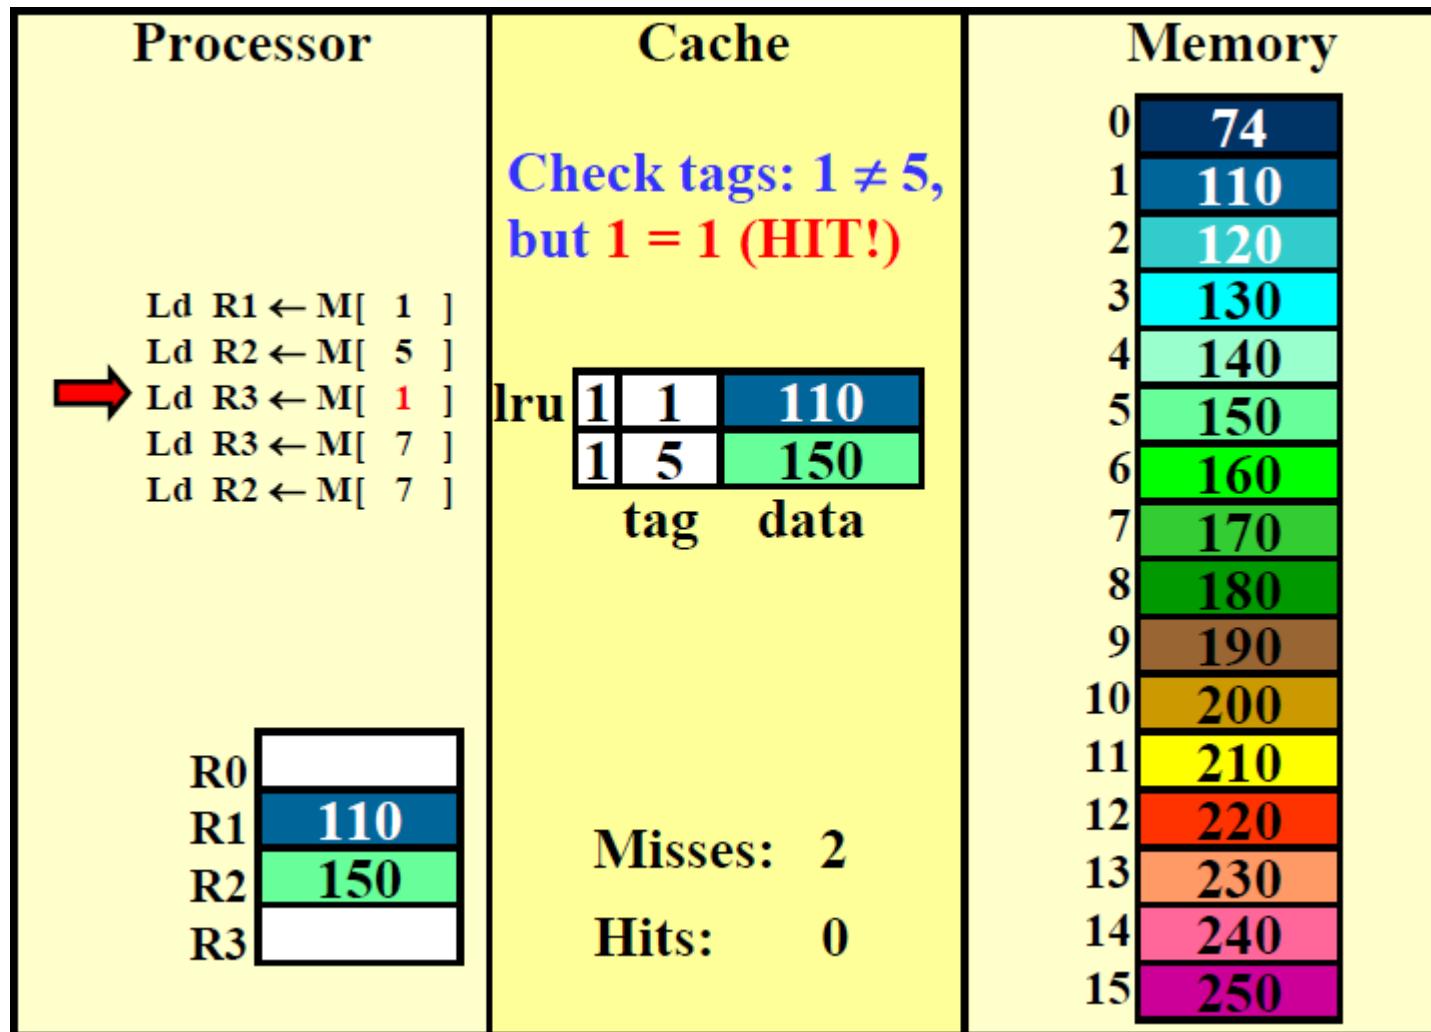
Simple Memory System



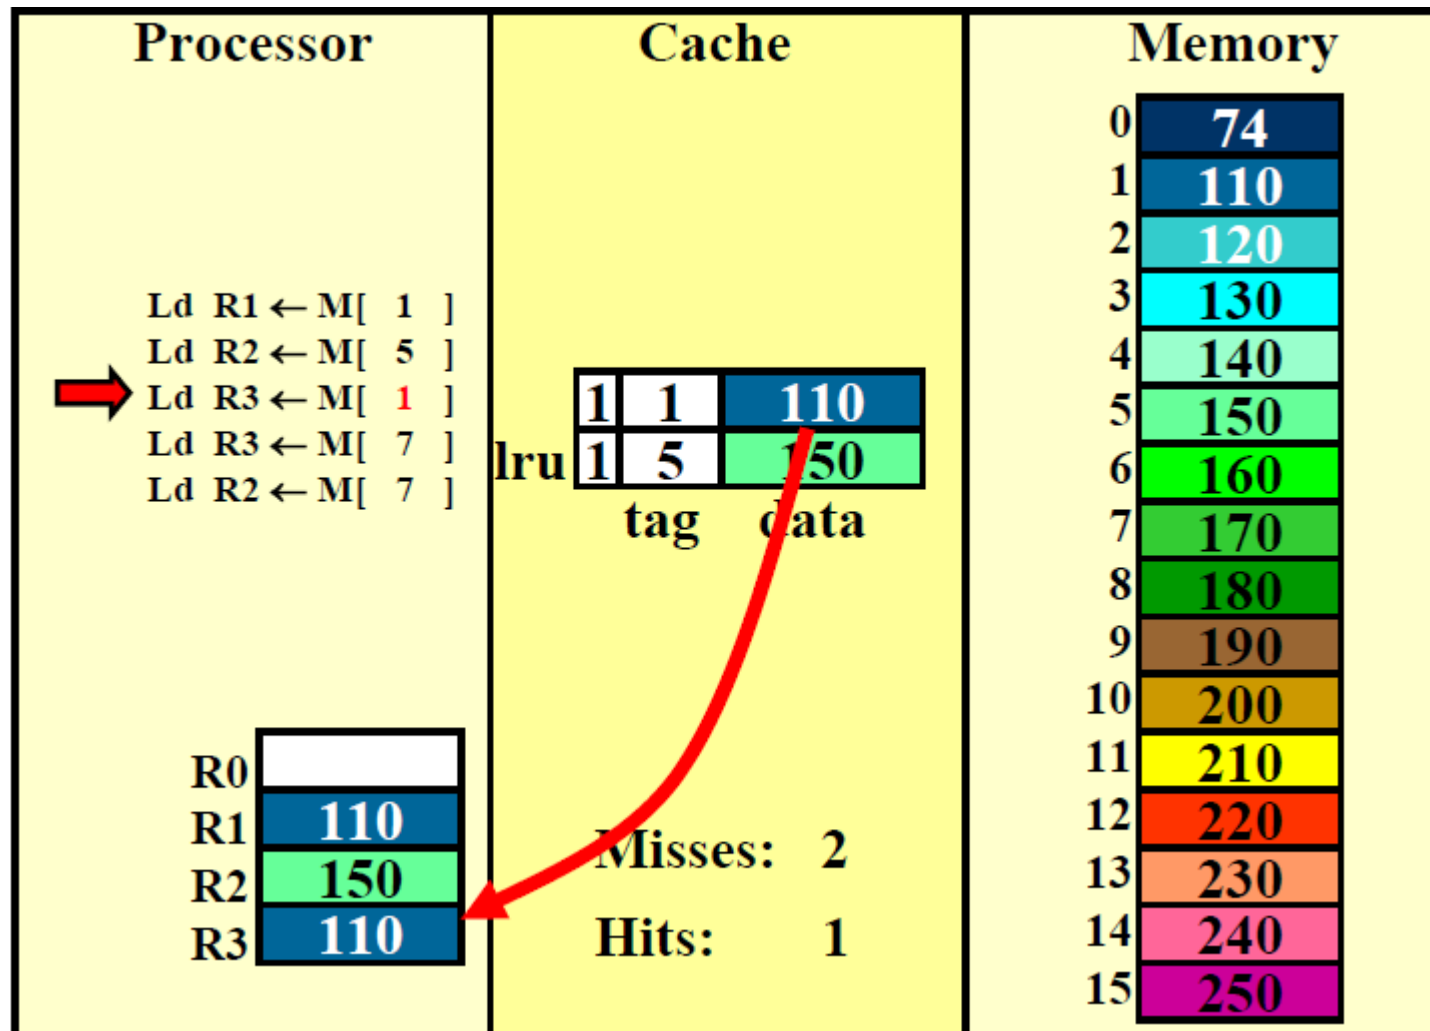
Simple Memory System



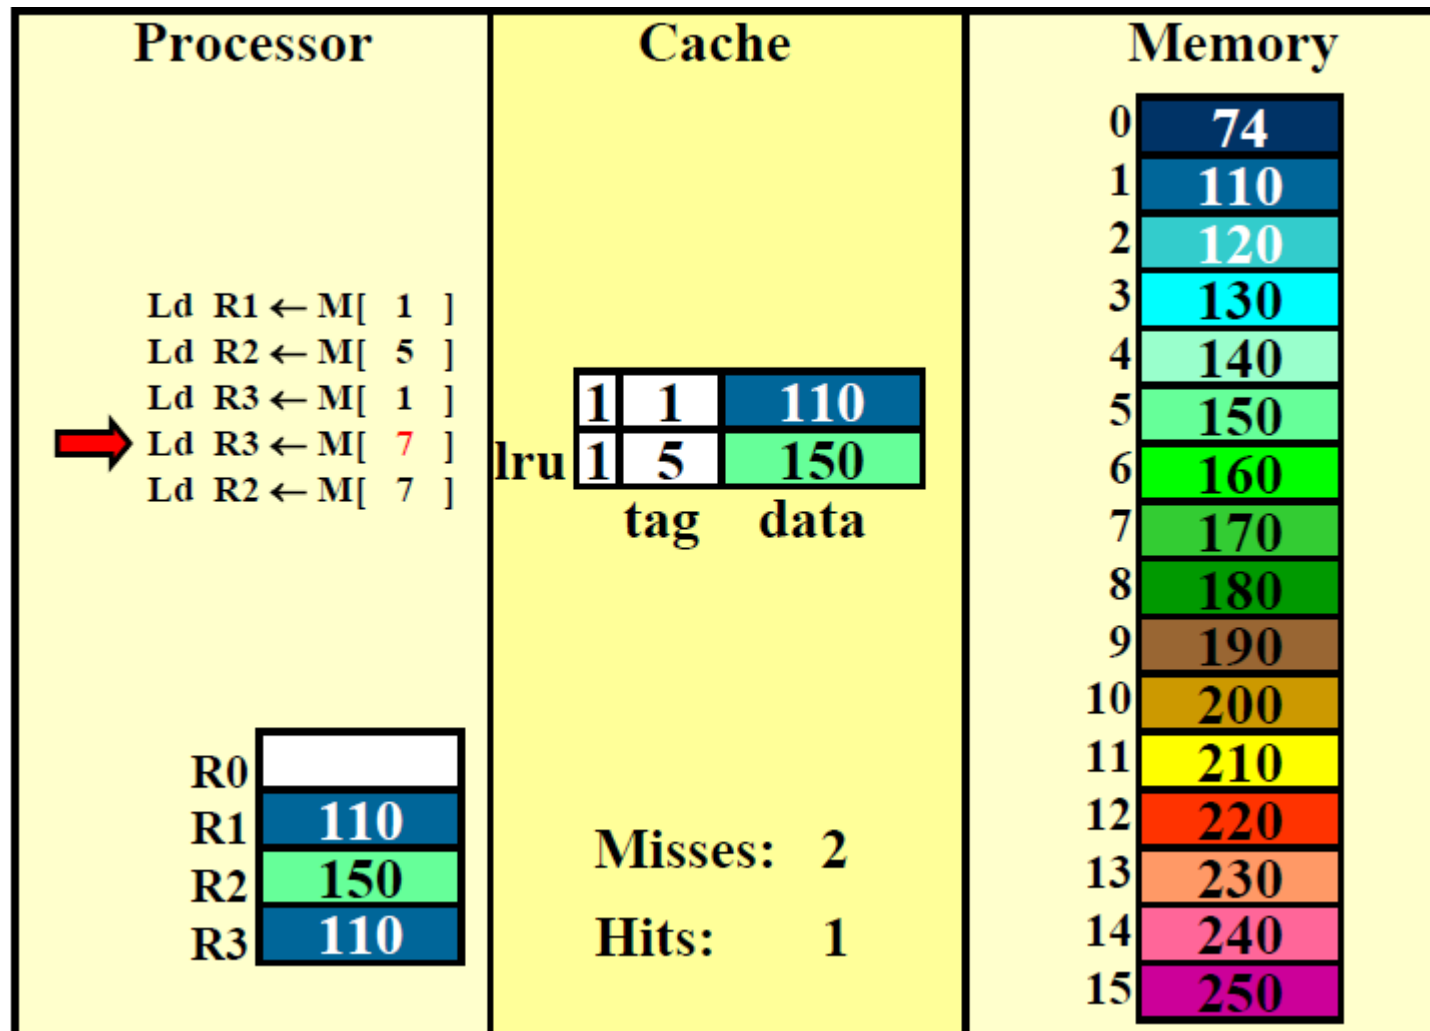
Simple Memory System



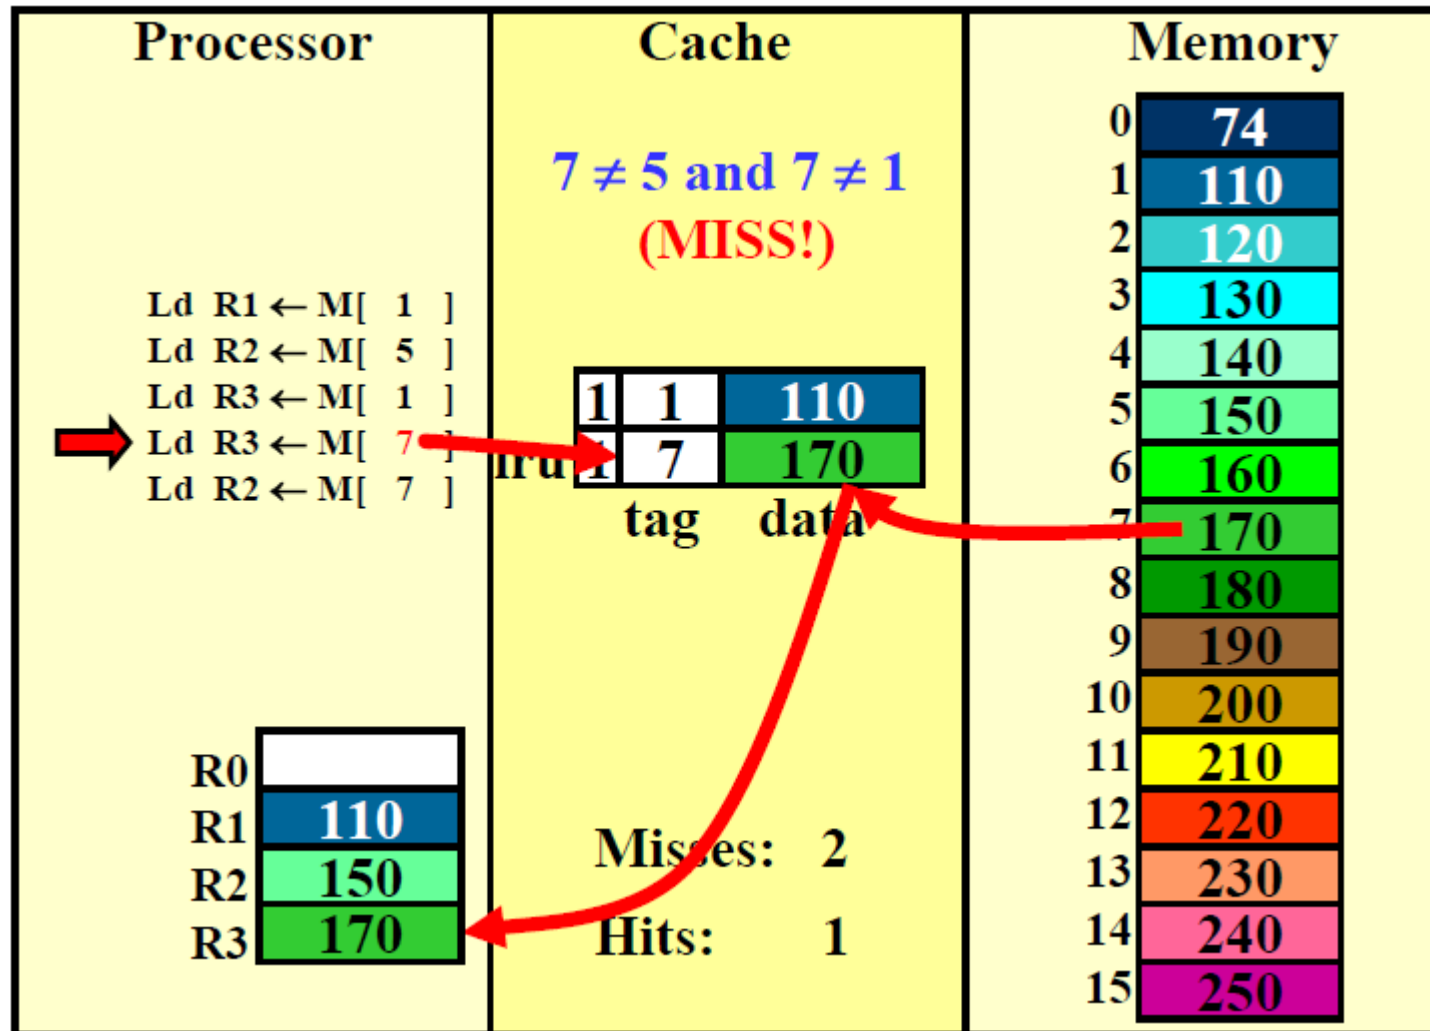
Simple Memory System



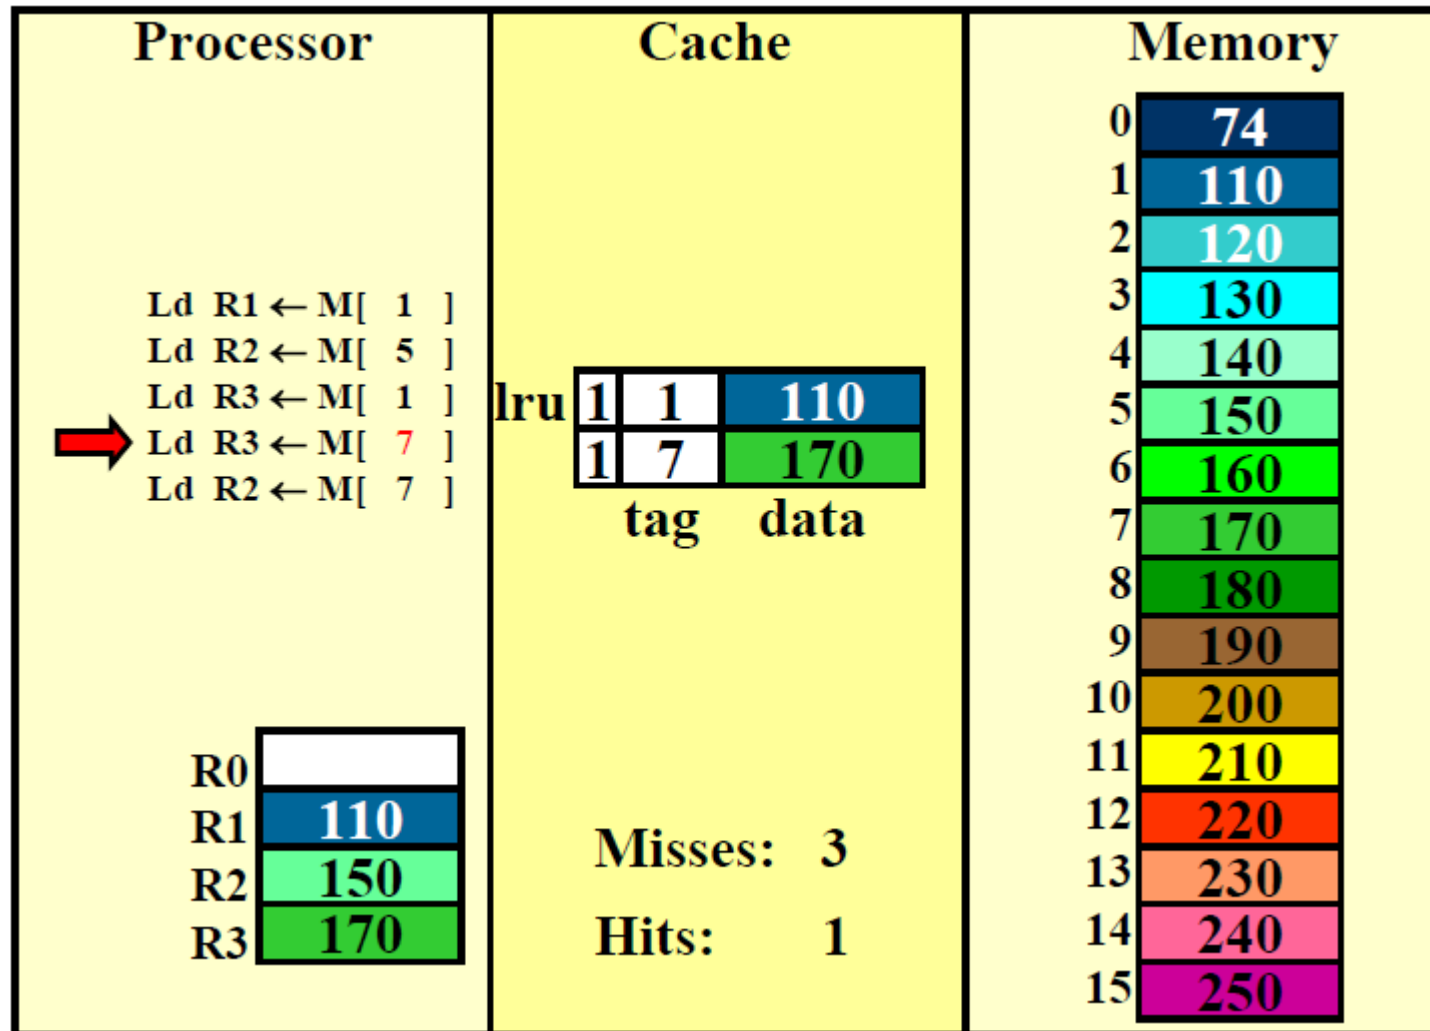
Simple Memory System



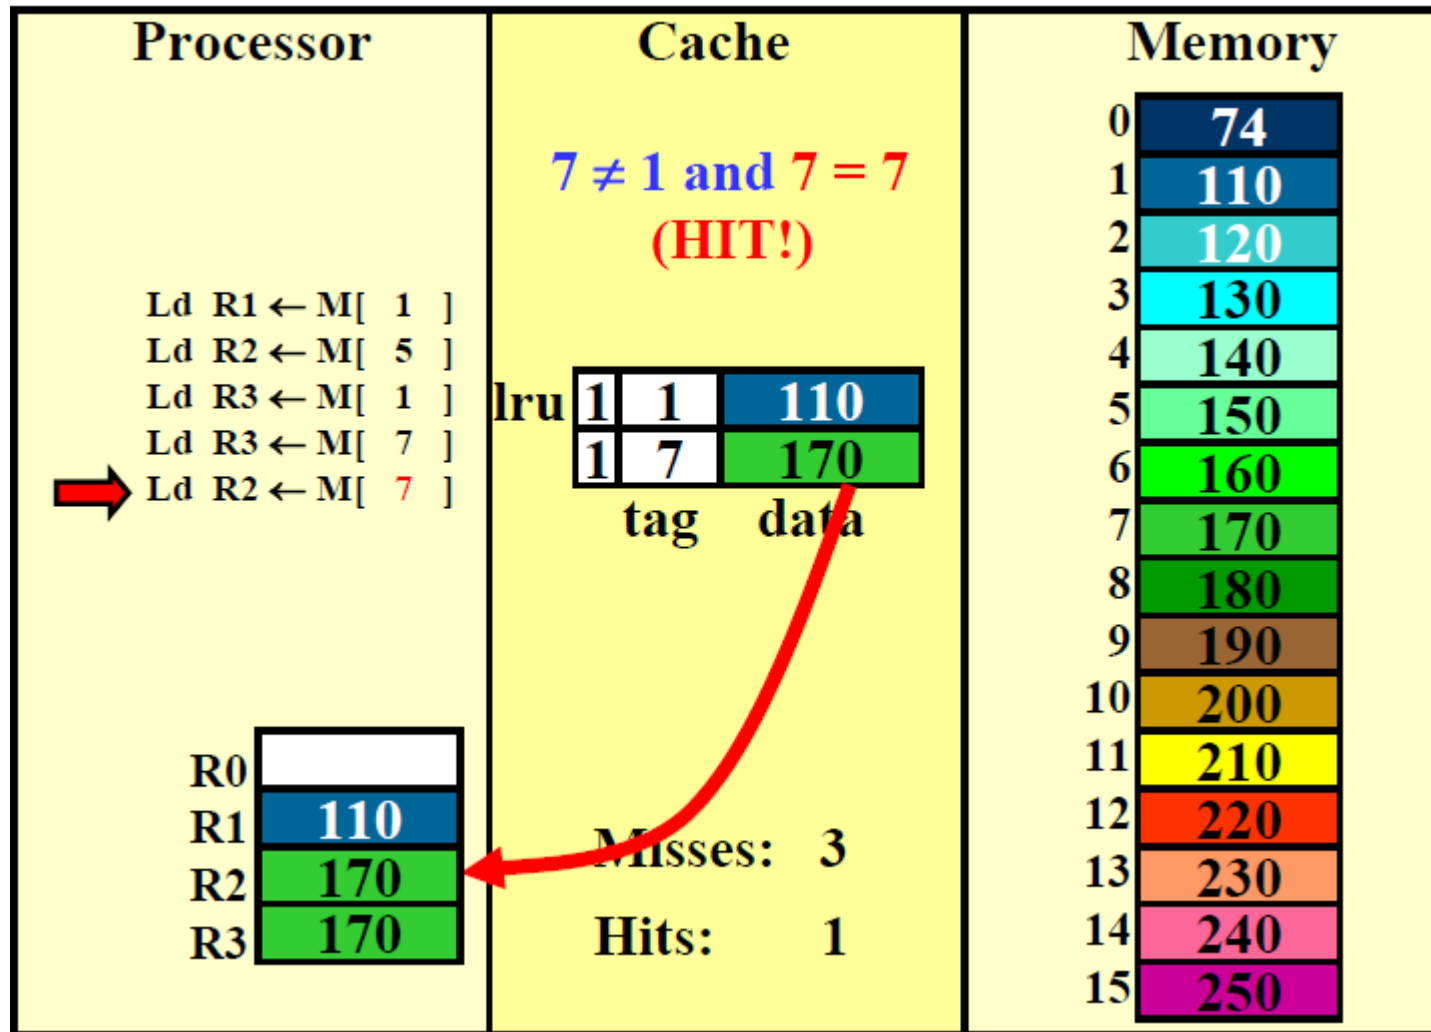
Simple Memory System



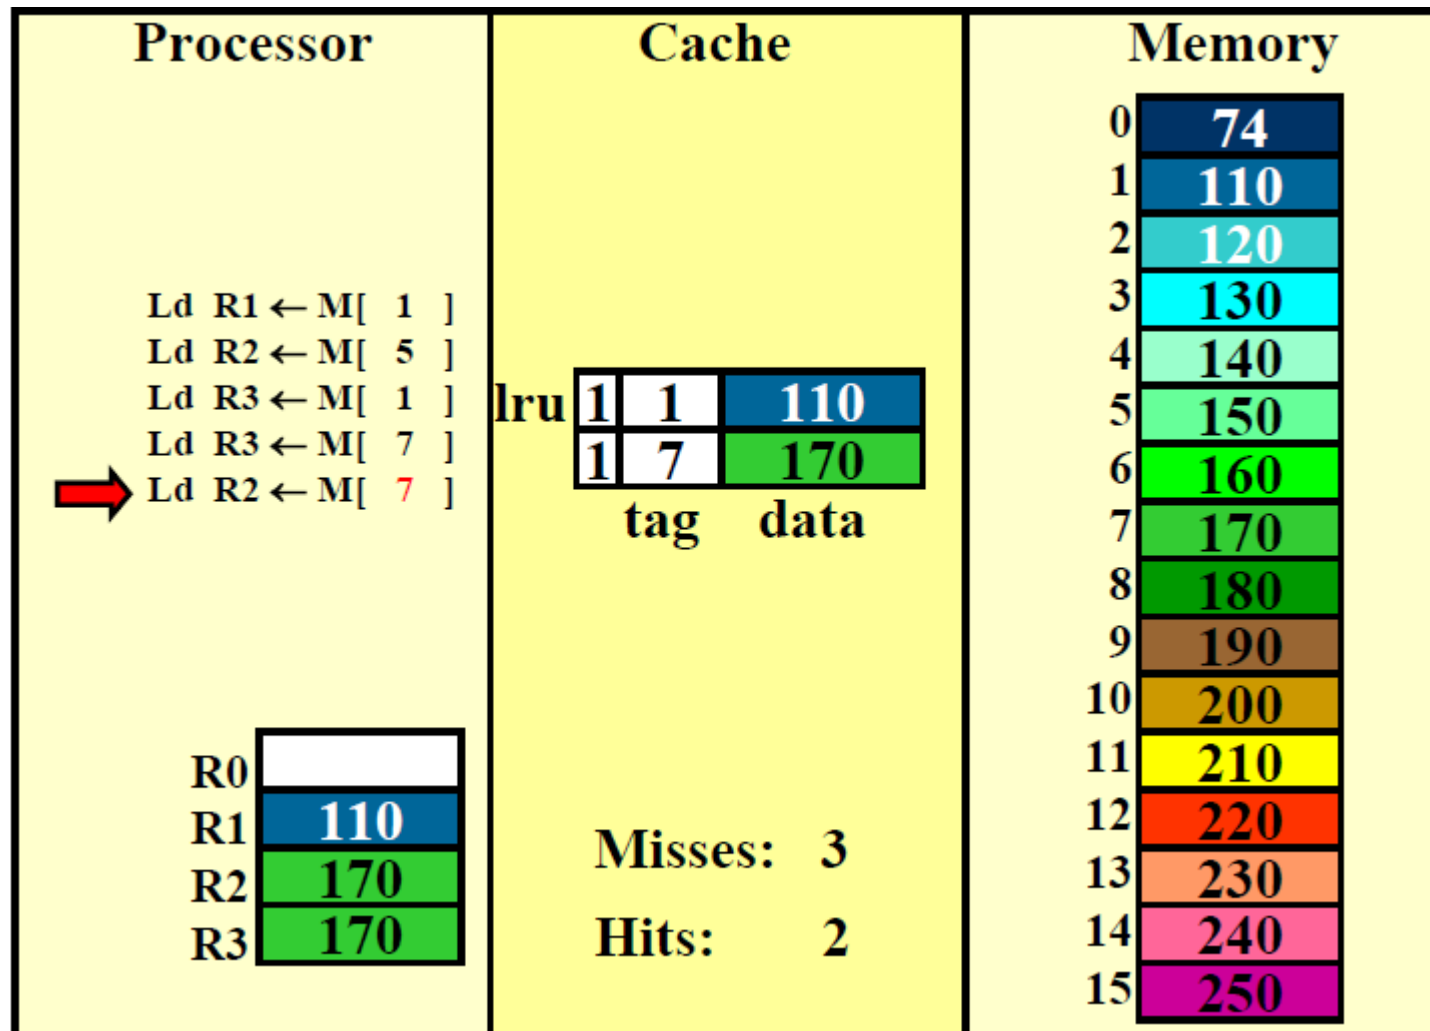
Simple Memory System



Simple Memory System



Simple Memory System





Thanks for Your Attention!

