

Computer Architecture: I/O and Handshaking

Hossein Asadi (asadi@sharif.edu)

Department of Computer Engineering

Sharif University of Technology

Spring 2024



Copyright Notice

- Some Parts (text & figures) of this Lecture adopted from following:
 - D.A. Patterson and J.L. Hennessy, “[Computer Organization and Design: the Hardware/Software Interface](#)” (MIPS), 6th Edition, 2020.
 - J.L. Hennessy and D.A. Patterson, “[Computer Architecture: A Quantitative Approach](#)”, 6th Edition, Nov. 2017.
 - “Intro to Computer Architecture” handouts, by Prof. Hoe, CMU, Spring 2009.
 - “Computer Architecture & Engineering” handouts, by Prof. Kubiawicz, UC Berkeley, Spring 2004.
 - “Intro to Computer Architecture” handouts, by Prof. Hoe, UWisc, Spring 2021.
 - “Computer Arch I” handouts, by Prof. Garzarán, UIUC, Spring 2009.

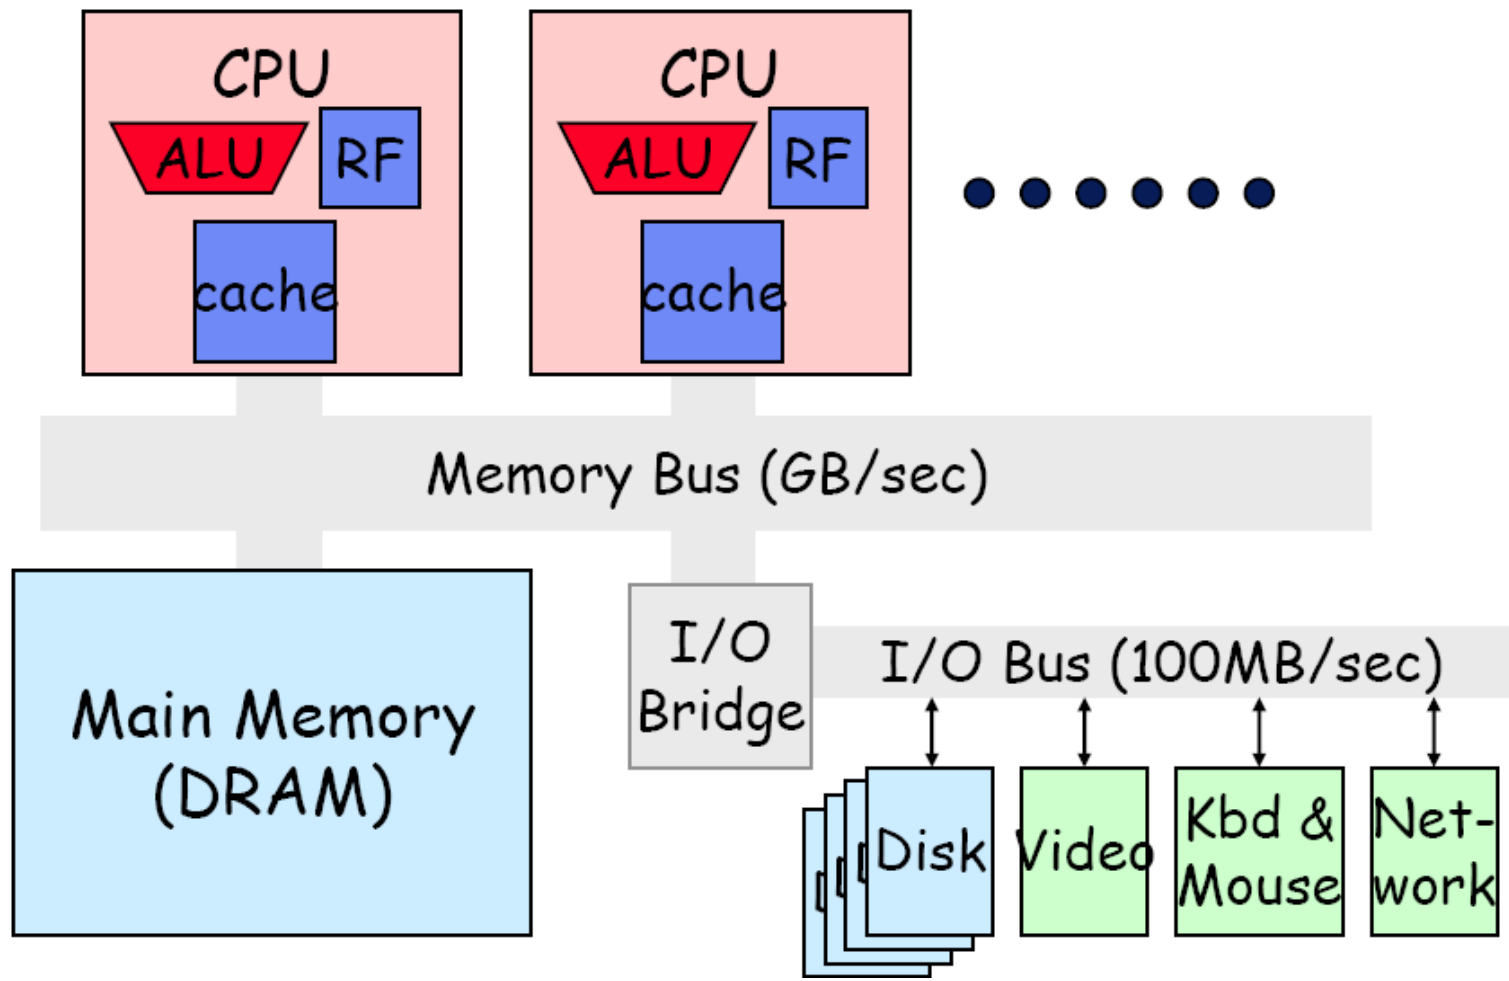


Topics Covered in This Lecture

- **I/O & Mass Storage**
 - Busses
 - I/O Handshaking



Computer Organization



Input/Output

- Fact
 - Input/outputs very slow devices
 - Average response time: few milli-seconds
 - CPUs very fast devices
 - Average running time: nano-seconds
- Question:
 - How Input/Output devices are connected to CPUs?
 - I/O controller (also called I/O bridge)



Input/Output (cont.)

- I/O Controller
 - Input/Output handshaking
 - Interrupt
 - Polling

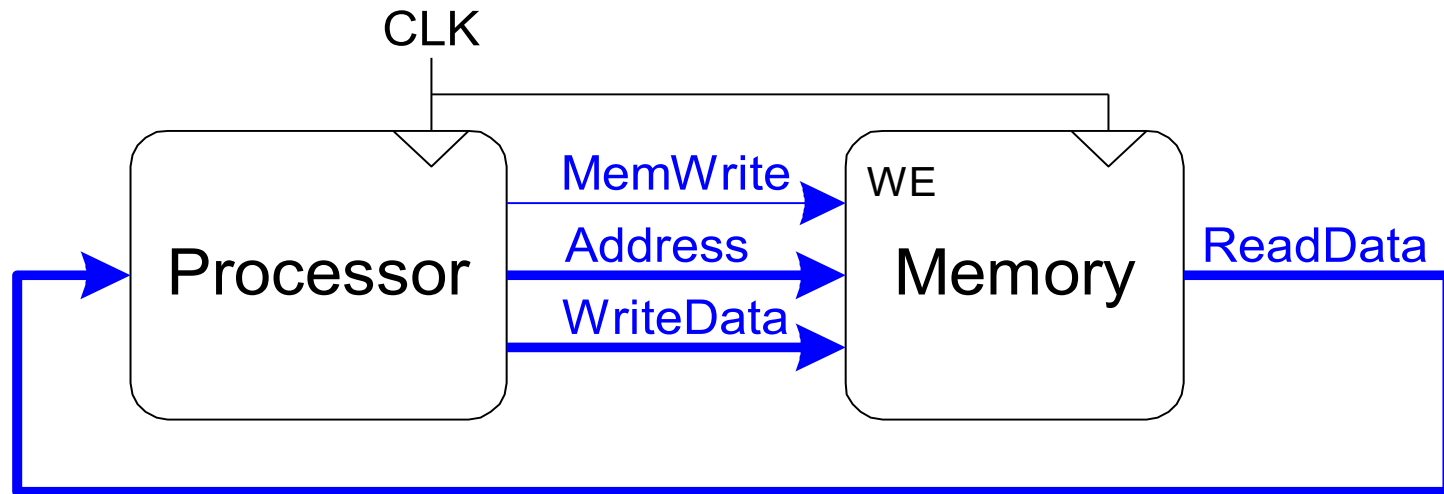


Input/Output (cont.)

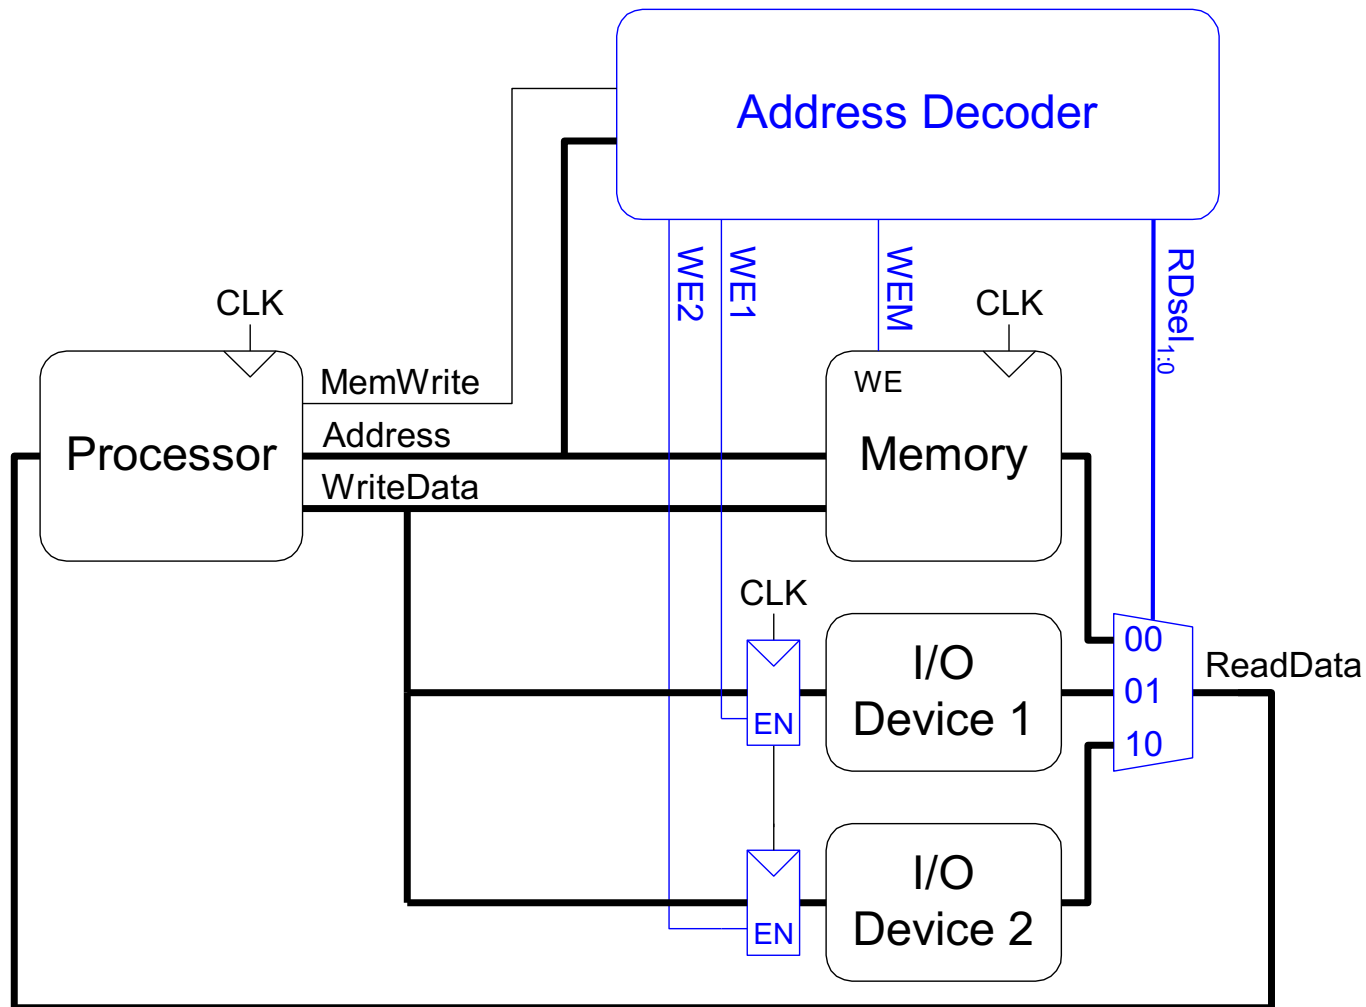
- How to Access I/O?
 - From ISA perspective
- I/O Configuration
 - I/O-Mapped I/O
 - Memory-Mapped I/O



Memory Interface



Memory-Mapped I/O Hardware



Memory-Mapped I/O Hardware (cont.)

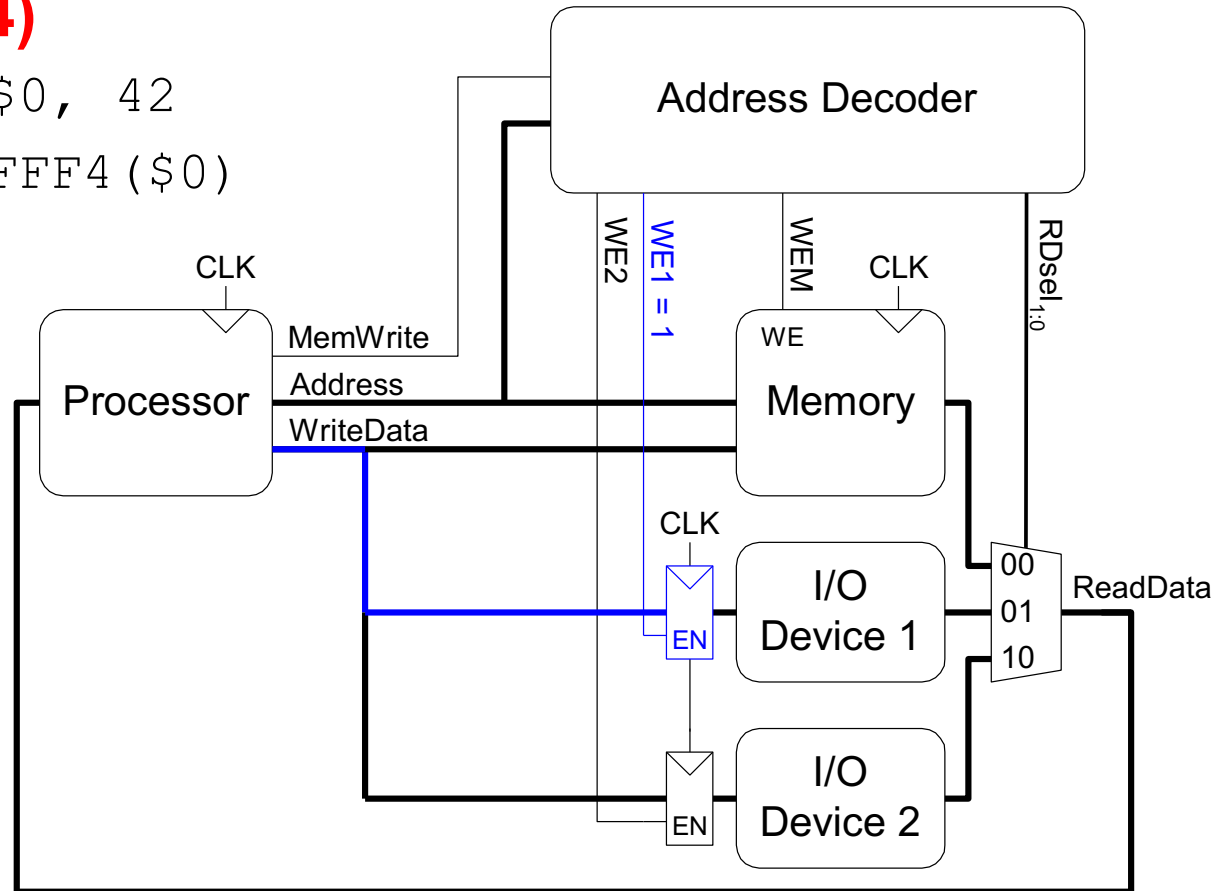
- Suppose I/O Device 1 is assigned the address 0xFFFFFFFF4
 - Write the value 42 to I/O Device 1
 - Read value from I/O Device 1 and place in `r1`



Memory-Mapped I/O Hardware (cont.)

- Write the value 42 to I/O Device 1 (0xFFFFFFF4)

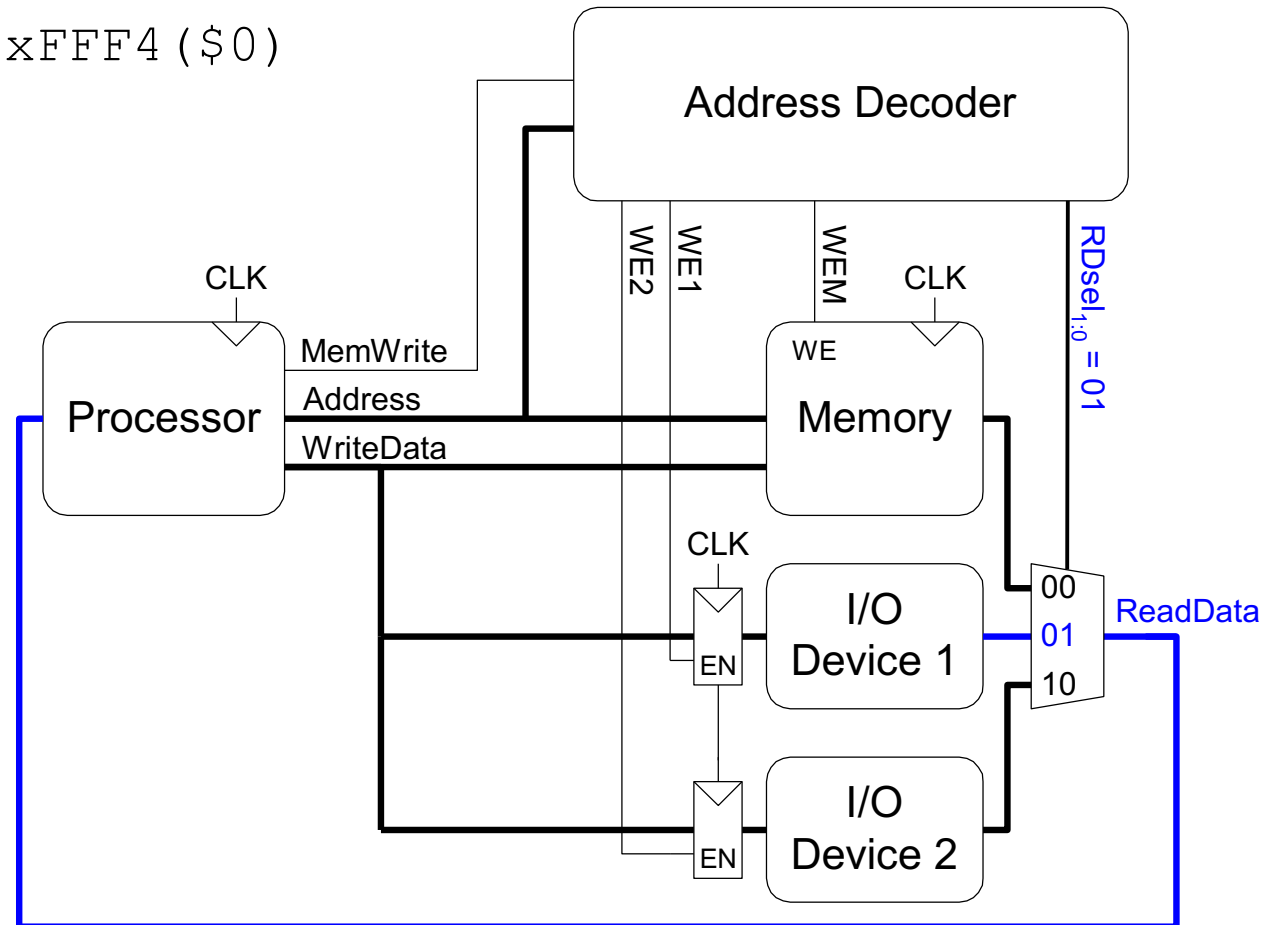
```
addi r1, $0, 42  
sw r1, 0xFFFF4($0)
```



Memory-Mapped I/O Hardware (cont.)

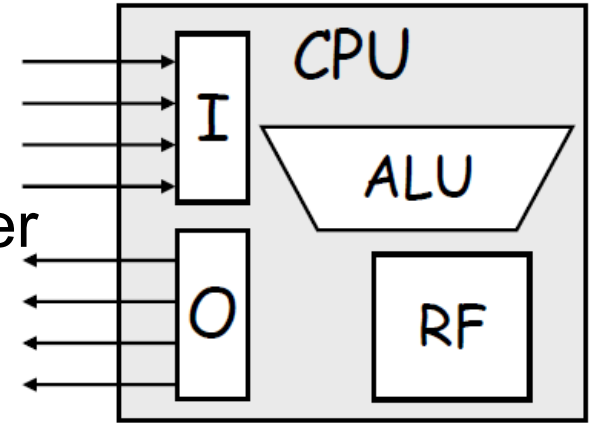
- Read value from I/O Device 1 and place in `r1`

```
lw r1, 0xFFF4($0)
```



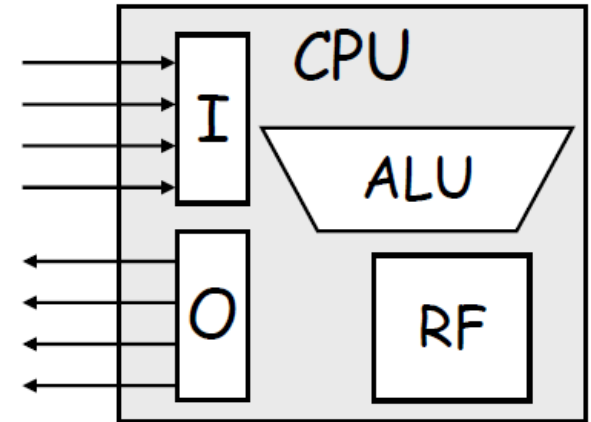
Input/Output (cont.)

- I/O Mapped I/O
 - Dedicated I/O instructions
 - Part of ISA
 - Output
 - Values written to an output register
 - On output pins of an external port
 - Input
 - Values read from an input register
 - At input pins of an external port
- Example: Intel microprocessors: IN / OUT



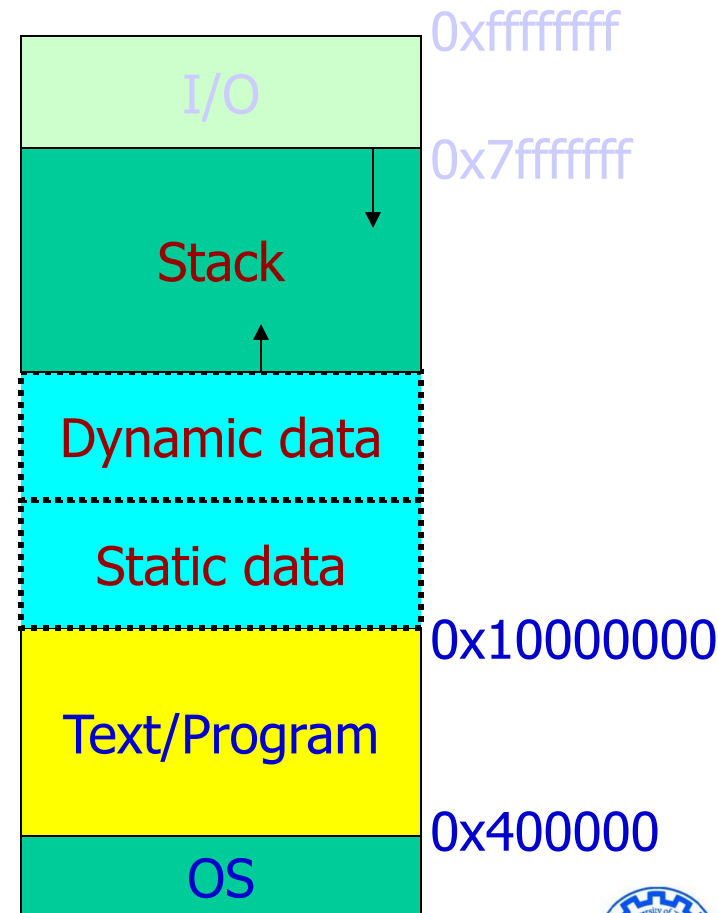
I/O-Mapped I/O

- Pros
 - Low latency
 - Simpler decoding
- Cons
 - Not implemented by all CPUs
 - Limited addressing
 - Predetermined # of I/O signals



Input/Output (cont.)

- Memory-Mapped I/O
 - Performed by simple **load/store** instructions
 - A subset of unused memory addresses mapped to registers of external devices
 - Memory & devices on bus programmed to respond only to their own address ranges



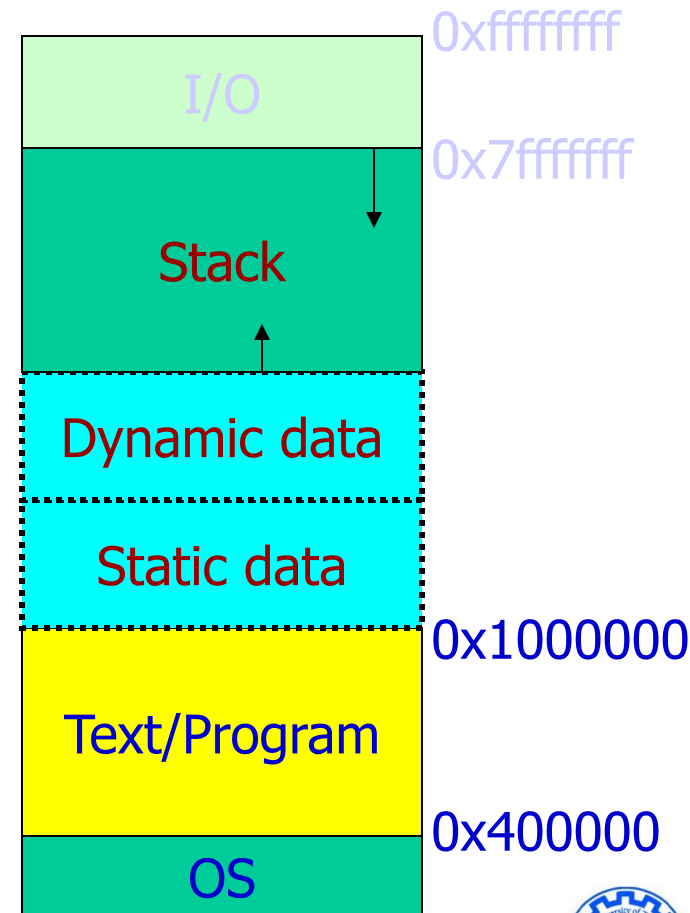
Memory-Mapped I/O

- Pros

- Easy to handle memory locations (SW perspective)
- Wide range of addresses

- Cons

- Caching may cause to memory incoherency issue
- Slow & consumes CPU cycles
 - → Use DMA





Input/Output (cont.)

- Standard I/O Protocols
 - PCIe
 - Serial ATA (SATA)
 - Serial Attached SCSI (SAS)
 - USB 2.0
 - Parallel ATA (PATA)

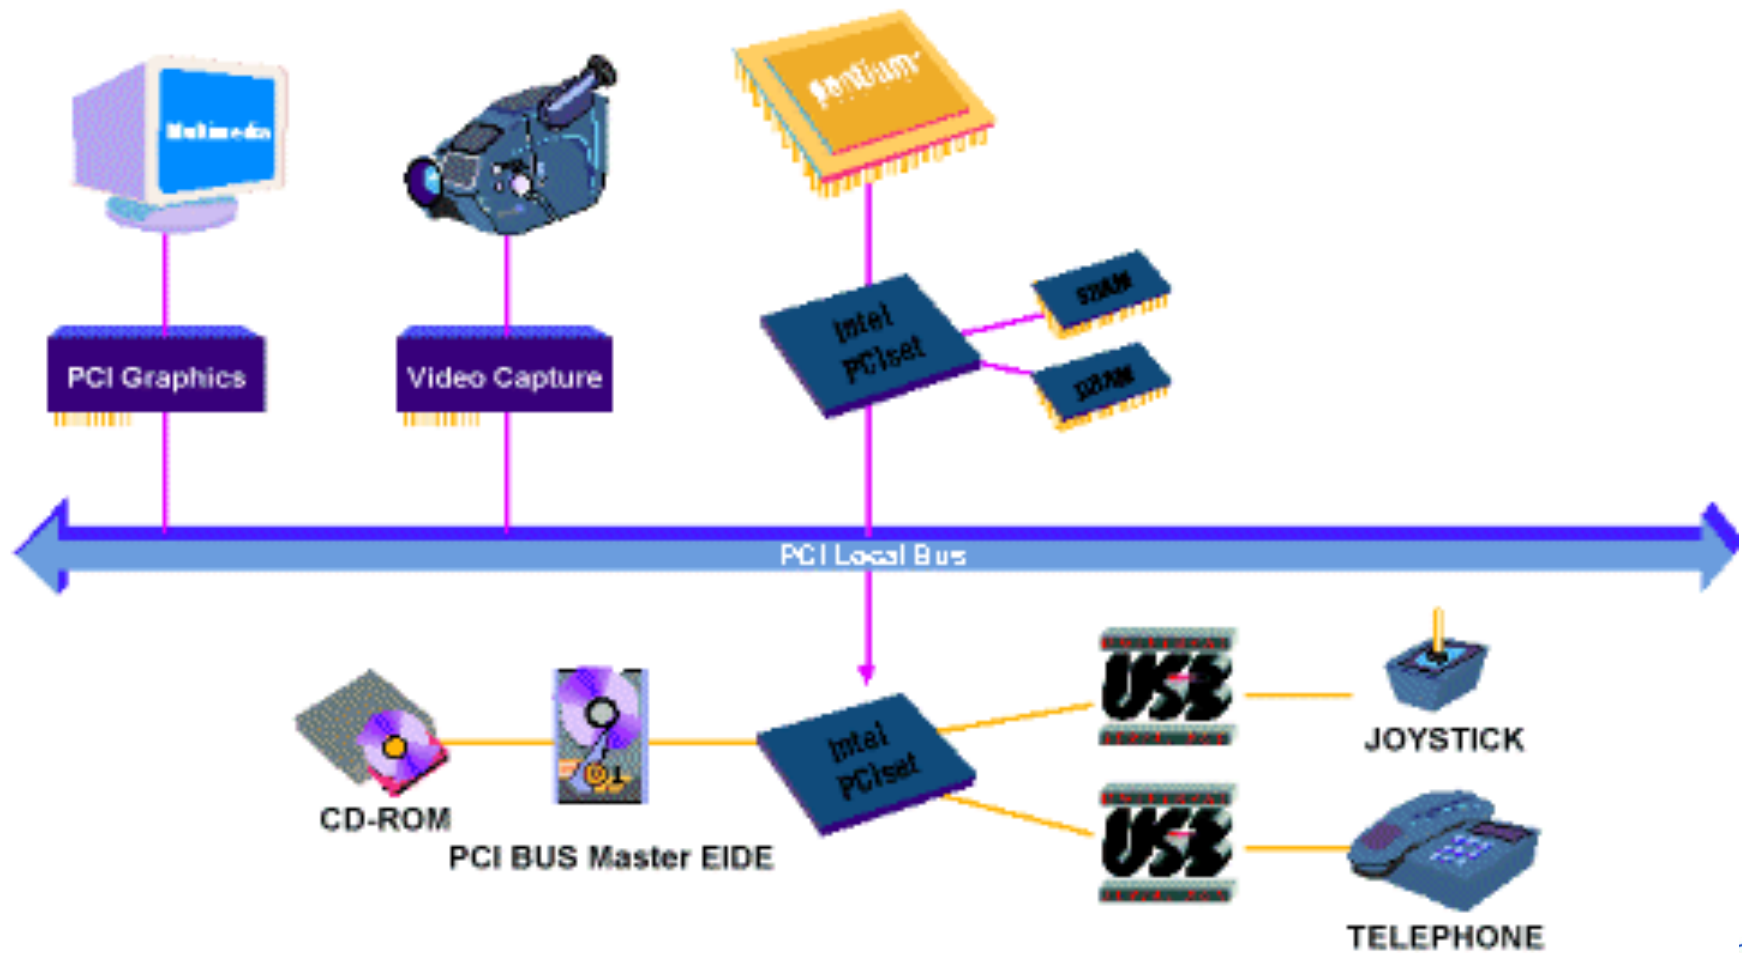


What is a Bus?


- Bus
 - Shared communication link
 - Single set of wires used to connect multiple subsystems



Buses




Advantages of Buses



- Versatility
 - New devices can be added easily
 - Peripherals can be moved between computer systems that use same bus standard
- Low Cost
 - A single set of wires shared in multiple ways


Disadvantage of Buses



- Creates a communication bottleneck
 - Bus bandwidth limits maximum I/O throughput
- Maximum bus speed largely limited by:
 - **Length** of bus
 - **Number** of devices on bus
 - **Slowest** device on bus



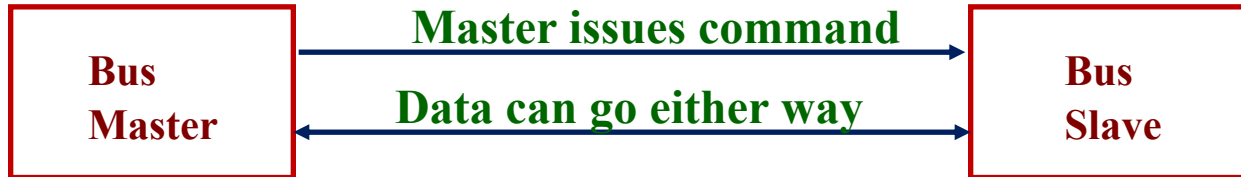
General Organization of Buses



- **Control Lines**
 - Signal requests and acknowledgments
 - Indicate what type of information is on data lines
- **Data Lines**
 - Carry information between source and destination
 - Data and Addresses
 - Complex commands



Master versus Slave



- A **Bus Transaction** Includes Two Parts:
 - Issuing command (and address)
 - Transferring data
 - request
 - action
- Master: one who starts bus transaction by:
 - Issuing command (and address)
- Slave: one who responds to address by:
 - Sending data to master (M) if M asks for data
 - Receiving data from M if M wants to send data



Types of Buses

- Processor-Memory Bus
- I/O Bus
- Backplane Bus



Types of Buses

- Processor-Memory Bus
 - Design specific
 - Short and high speed
 - Only need to match memory system
 - Maximize memory-to-processor bandwidth
 - Connects directly to processor
 - Optimized for cache block transfers



Types of Buses

- I/O Bus
 - Industry standard
 - Usually is lengthy and slower
 - Need to match a wide range of I/O devices
 - Connects to processor-memory bus or backplane bus




Types of Buses

- Backplane Bus
 - Standard or proprietary
 - A single bus used for:
 - Processor to memory communication
 - Communication between I/O devices and memory
 - Backplane is an interconnection structure within chassis
 - Allow processors, memory, and I/O devices to coexist




Types of Buses

- Backplane Bus
 - Advantages: Simple and low cost
 - Disadvantages: slow and bus can become a major bottleneck
 - Example: IBM PC - AT



Two-Bus System

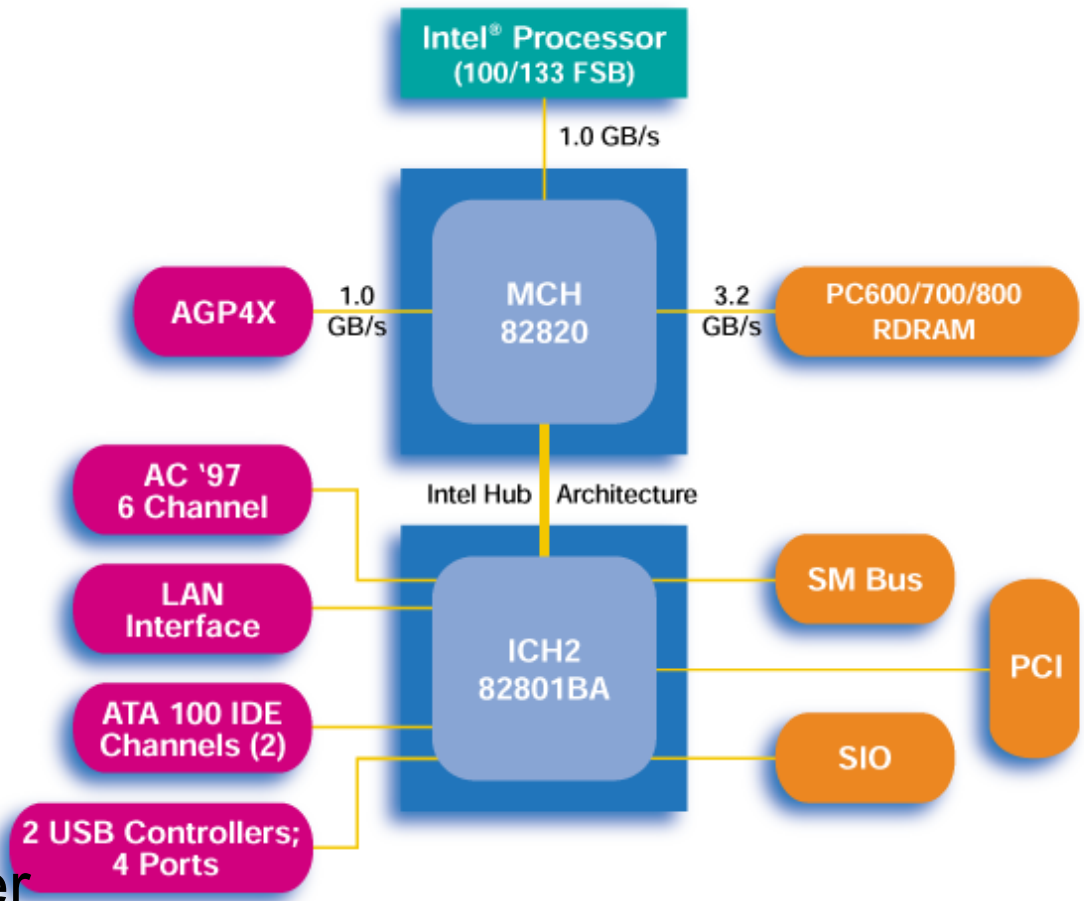


- I/O buses tap into processor-memory bus via bus adaptors:
- I/O buses provide expansion slots for I/O devices
- Example: Apple Macintosh-II



Main components of Intel Chipset: Pentium II/III


- Northbridge:
 - Handles memory
 - Graphics
- Southbridge: I/O
 - PCI bus
 - Disk controllers
 - USB controllers
 - Audio
 - Serial I/O
 - Interrupt controller
 - Timers



Delegating I/O Responsibility from CPU: DMA

- Direct Memory Access (DMA):
 - External to CPU
 - Act as a maser on bus
 - Transfer blocks of data to or from memory without CPU intervention

CPU sends a starting address, direction, and length count to DMAC. Then issues "start".



DMAC provides handshake signals for Peripheral Controller, and Memory Addresses and handshake signals for Memory.



Bus Transactions

- Transaction Phases
 - Master **requests** ownership from arbiter
 - Arbiter **grants** ownership to master
 - Master drives **address** for all to see
 - A slave **claims** transaction
 - Master (or slave) drives **data** (depending on read or write) for all to see
 - Master **terminates** transaction and bus ownership



Basic Bus Signals

- CLK
 - All devices synchronized by a clock signal
- Private Signals to/from Arbiter per Master
 - **REQ** (output): assert to request ownership; de-assert to signal end of transaction
 - **GNT** (input): ownership is granted



Basic Bus Signals

- “Broadcast” signals shared by all devices
 - **AD** (address/data bus, bi-directional): master drives address during the address phase, master/slave drives data during the data phase
 - Why not have separate address bus and data bus?
- R/W (bi-directional):
 - Bus commands, e.g. read vs. write

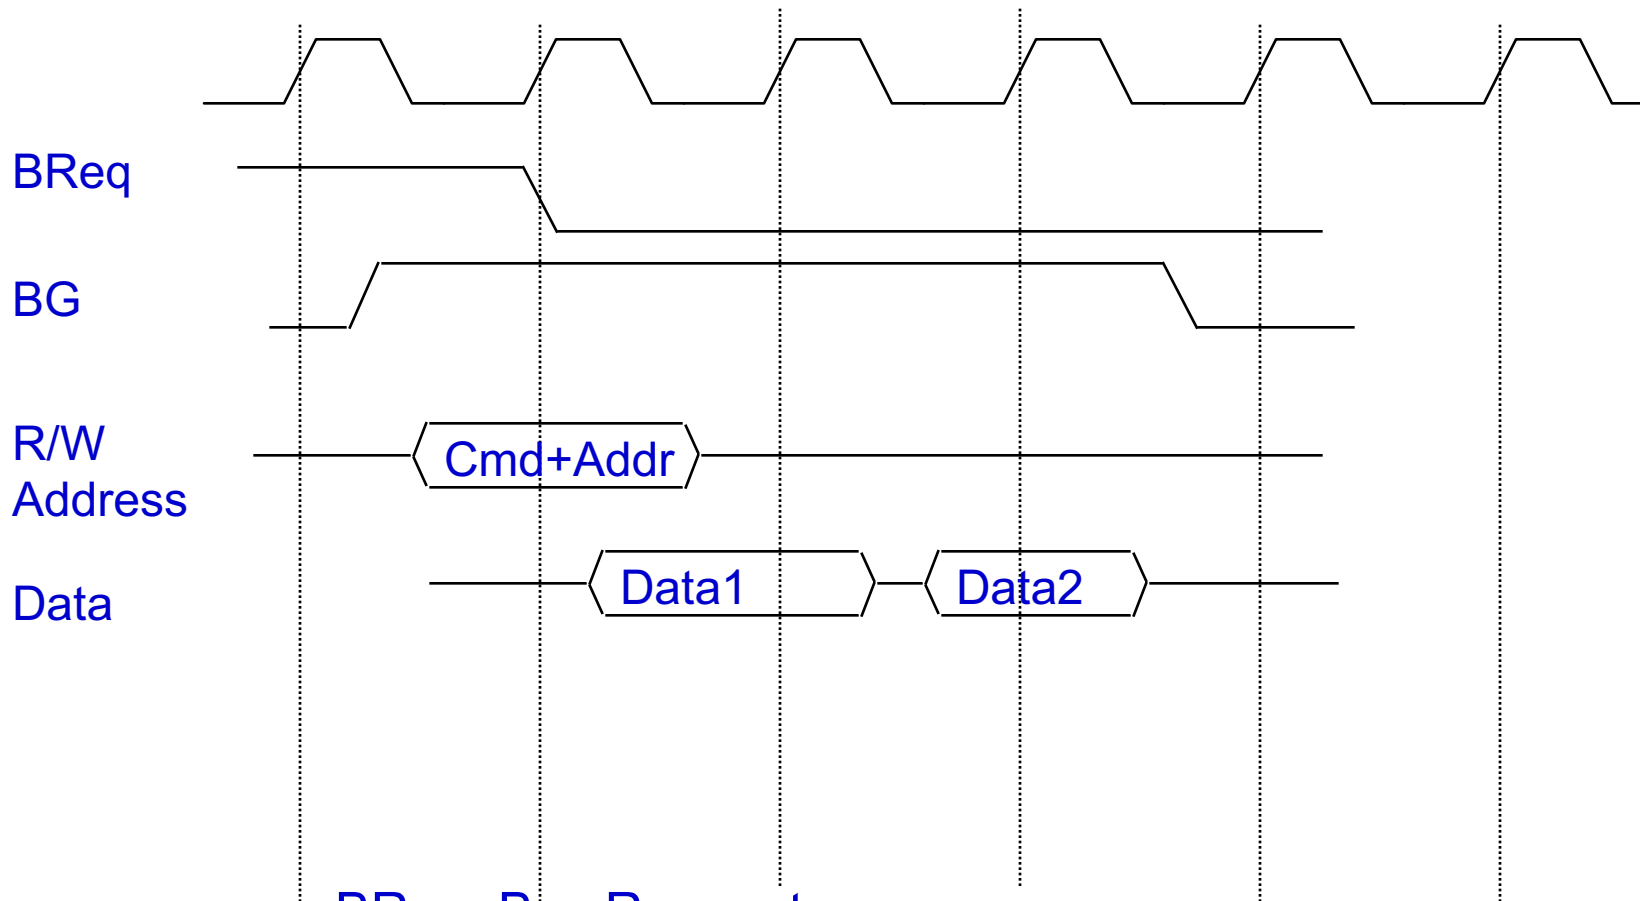


Synchronous Bus

- **Synchronous** Bus:
 - Includes a clock in control lines
 - A fixed protocol relative to the clock
 - Advantage: little logic and very fast
 - Disadvantages:
 - Every device on the bus must run at the same clock rate
 - To avoid clock skew, they cannot be long if they are fast



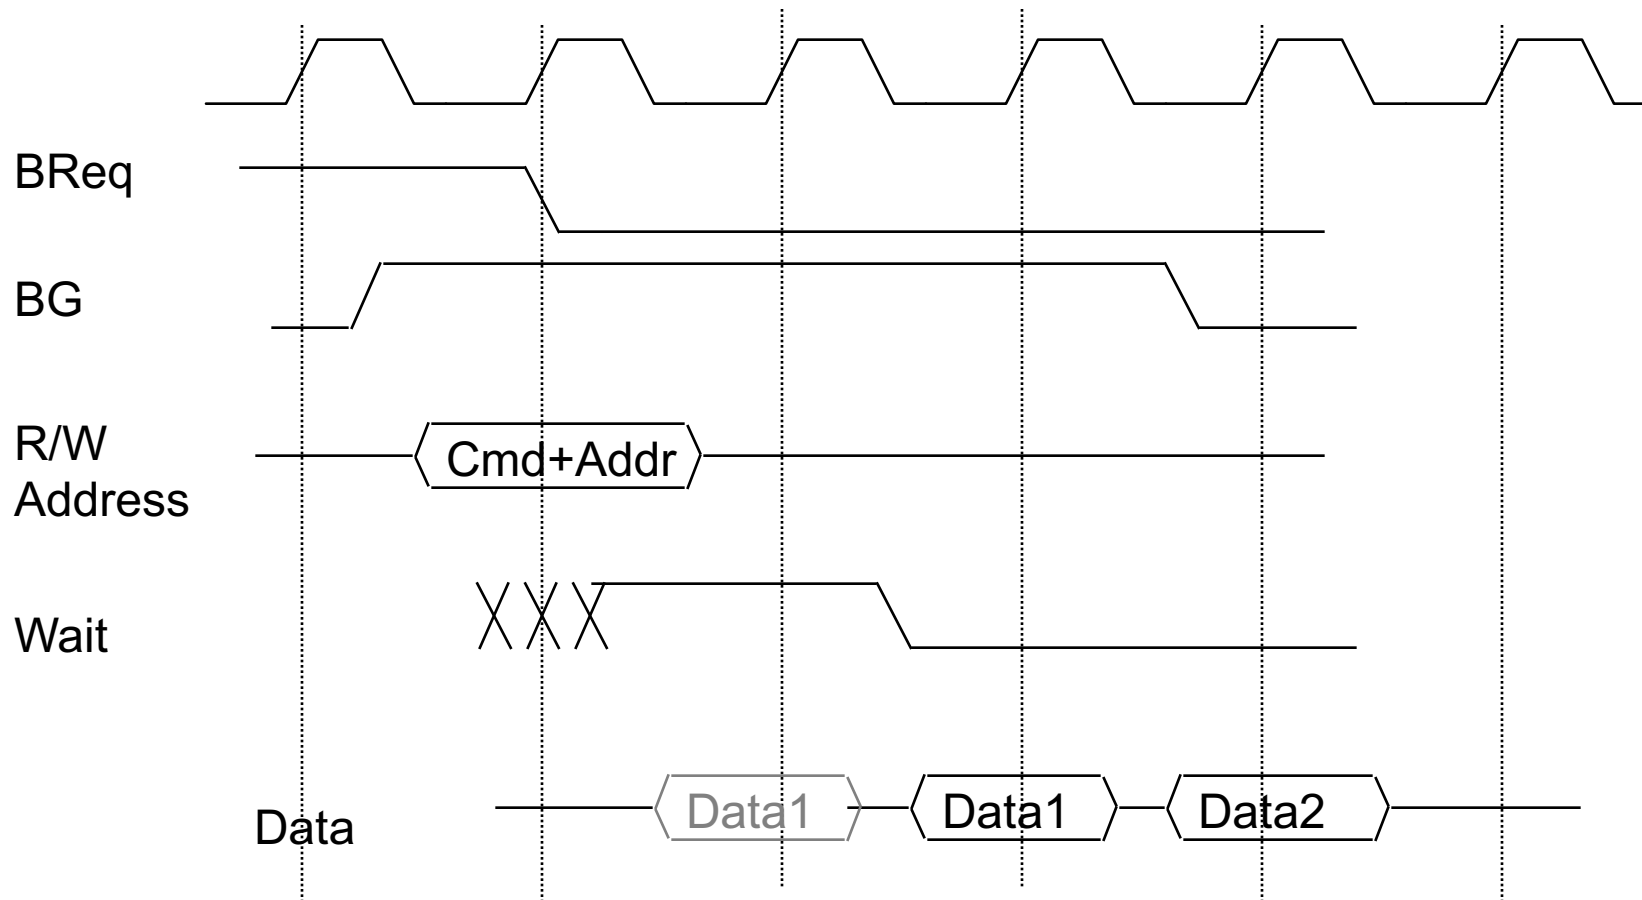
Simple Synchronous Protocol



BReq: Bus Request
BG: Bus Grant



Typical Synchronous Protocol



- Slave indicates when it is prepared for data xfer
- Actual transfer goes at bus rate




Asynchronous Bus


- **Asynchronous** Bus:
 - It is not clocked
 - It can accommodate a wide range of devices
 - It can be lengthened without worrying about clock skew
 - It requires a handshaking protocol



Data Transfer Initiated by Source Unit



(a) Block Diagram




(b) Timing Diagram


Source-Initiated strobe for Data Transfer



Data Transfer Initiated by Destination Unit



(a) Block Diagram





(b) Timing Diagram


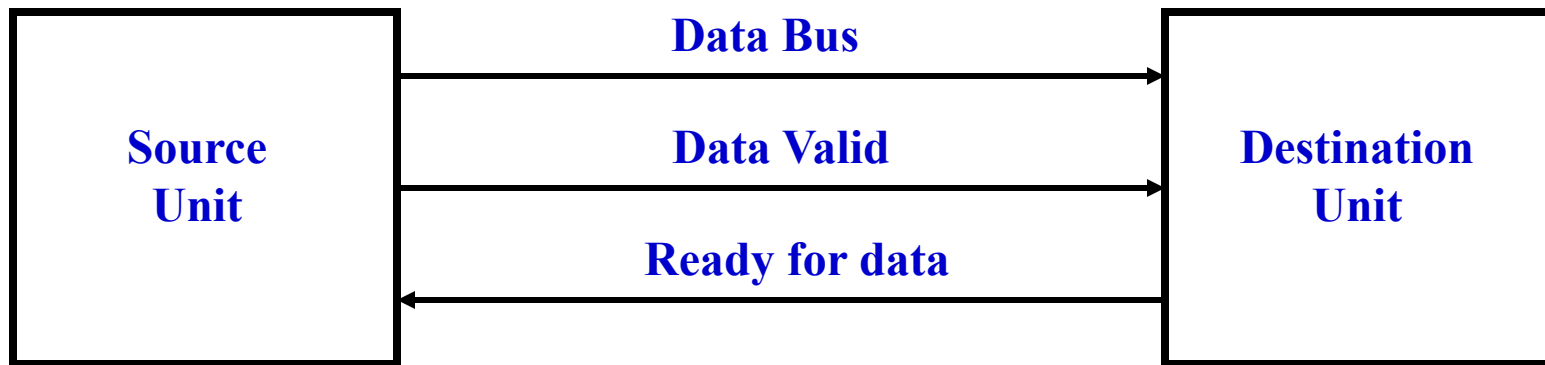
Destination-Initiated strobe for Data Transfer



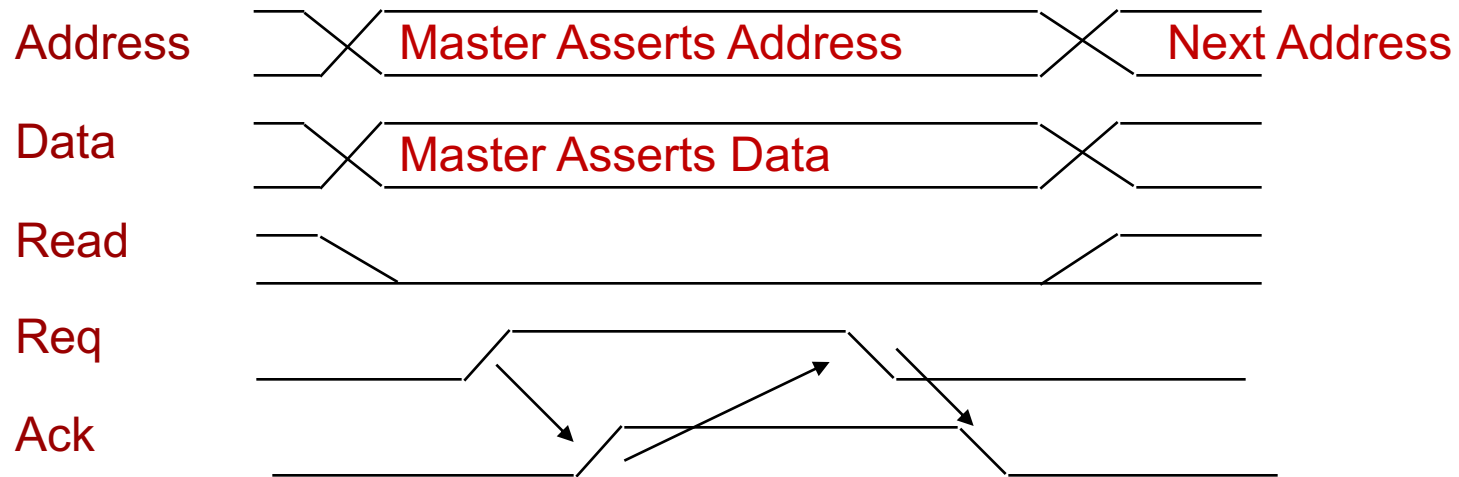
Source-Initiated Transfer Using Handshake



Destination-Initiated Transfer Using Handshake



Asynchronous Write Transaction



t0 t1 t2 t3 t4 t5

t0 : Master has obtained control and asserts address, direction, data
 Waits a specified amount of time for slaves to decode target.

t1: Master asserts request line


t2: Slave asserts ack, indicating data received

t3: Master releases req

t4: Slave releases ack




Asynchronous Read Transaction



- t0 : Master has obtained control and asserts address, direction, data
Waits a specified amount of time for slaves to decode target.
- t1: Master asserts request line
- t2: Slave asserts ack, indicating ready to transmit data
- t3: Master releases req, data received
- t4: Slave releases ack



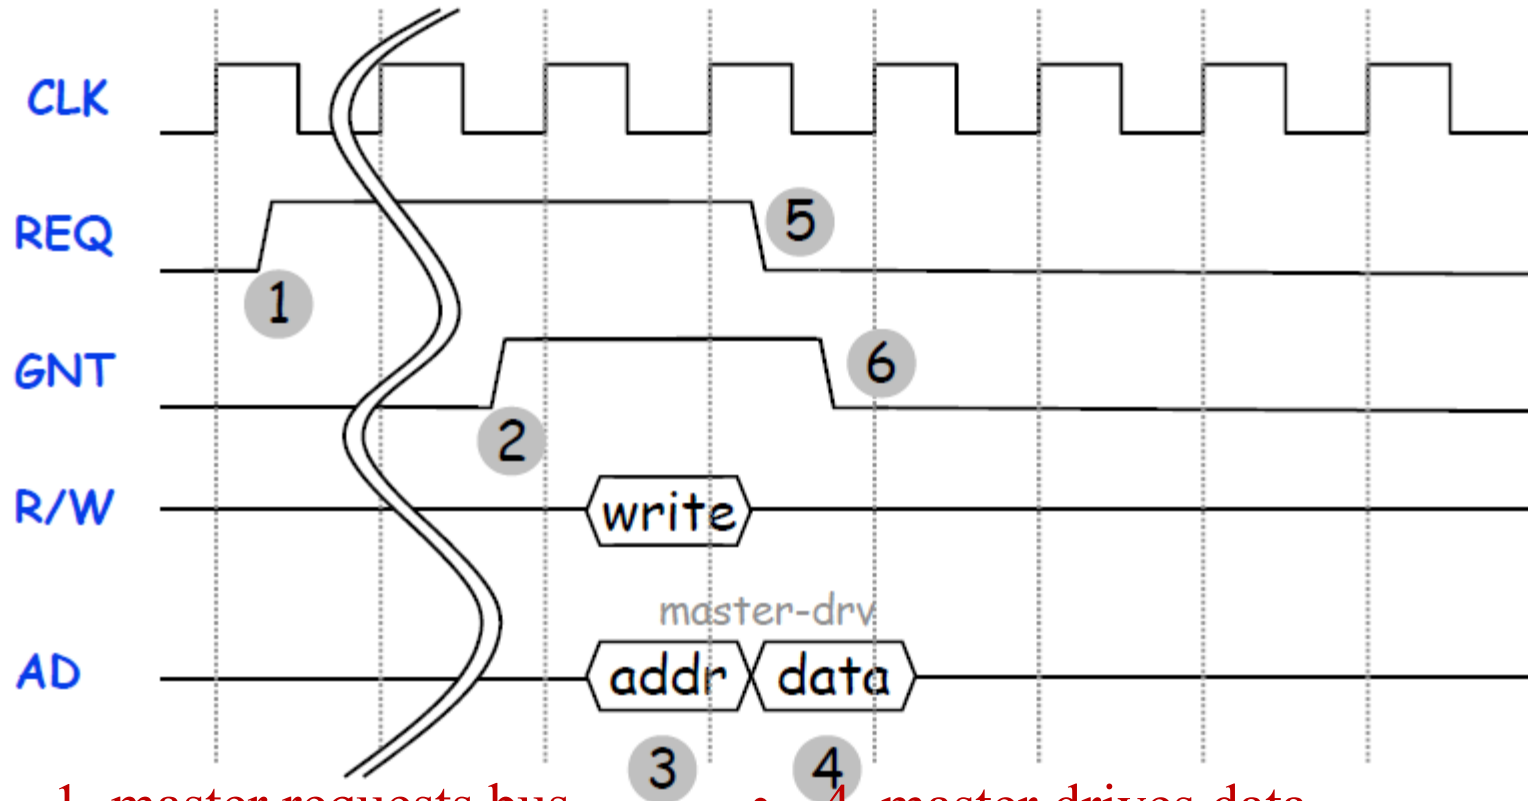
Simple Read Transaction



- 1. master requests bus
- 2. arbiter grants bus
- 3. master drives address /command
- 4. bus-turnaround cycle
- 5. slave drives data
- 6. master signals final cycle
- 7. arbiter acknowledges




Simple Write Transaction



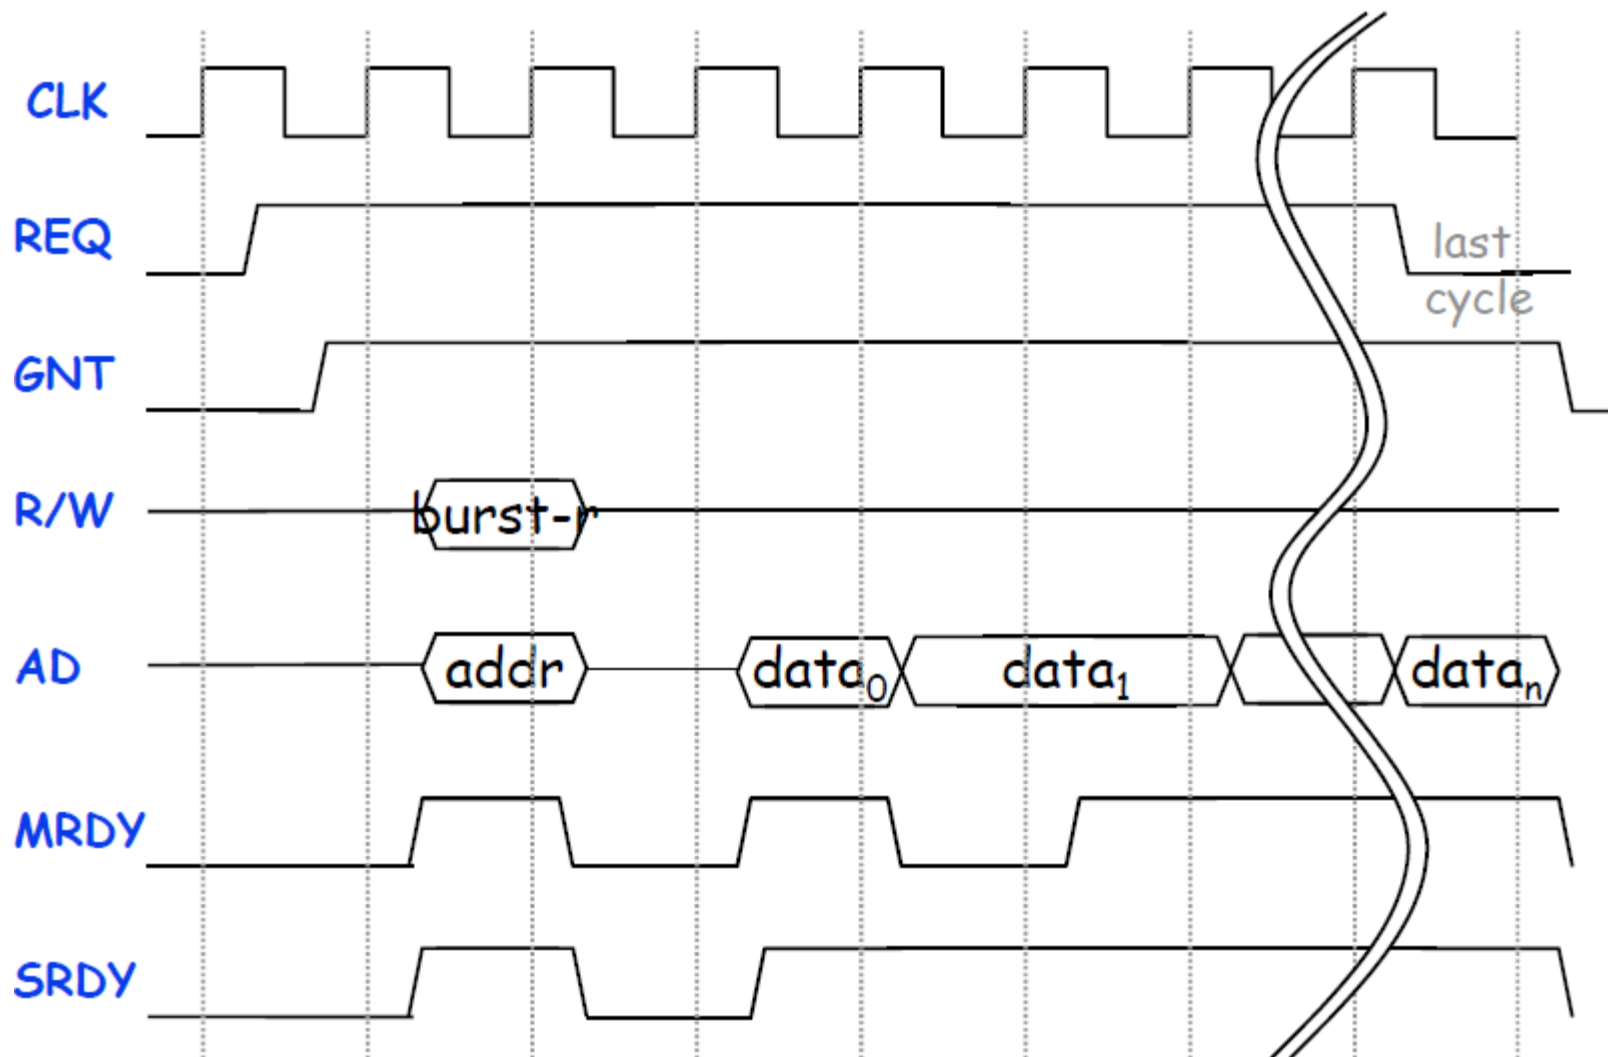
- 1. master requests bus
- 2. arbiter grants bus
- 3. master drives address /command
- 4. master drives data
- 5. master signals final cycle
- 6. arbiter acknowledges



Asynch Read Transaction



Burst Read Transaction




Bus Arbitration

- Arbitration scheme:
 - A bus master wants to use the bus; asserts the bus request
 - A bus master cannot use the bus until its request is granted
 - A bus master must signal to the arbiter after finish using the bus
- Arbitration schemes balance two factors:
 - Bus **priority**: the highest priority device should be serviced first
 - **Fairness**: Even the lowest priority device should never be completely locked out from bus

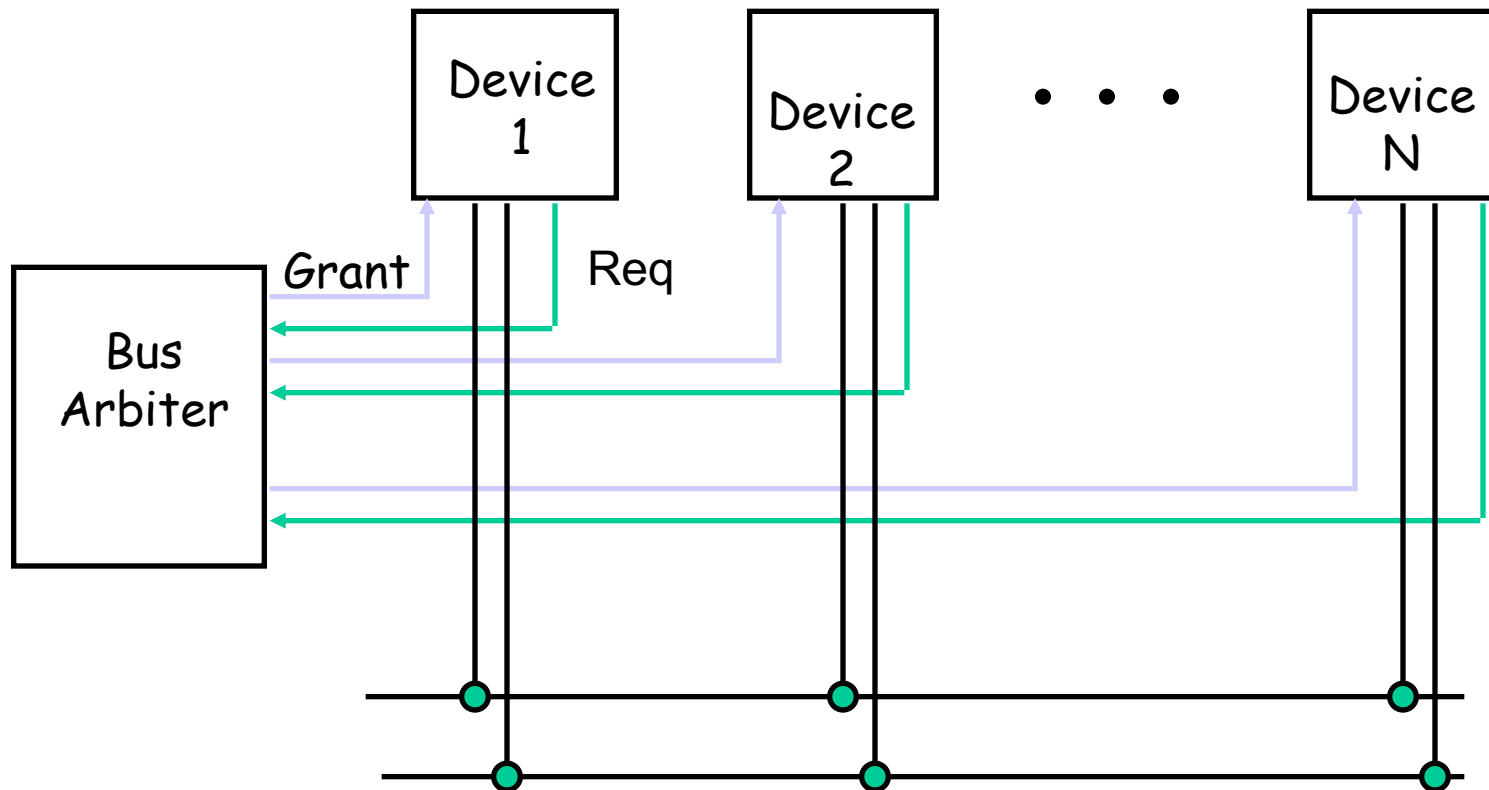


Simple Bus Arbitration

- Master-slave arrangement:
 - Only bus master can control access to bus
 - It initiates and controls all bus requests
 - Slave responds to read/write requests
 - Example:
 - Processor is only bus master
 - All bus requests controlled by processor
 - Major drawback ?



Centralized Parallel Arbitration

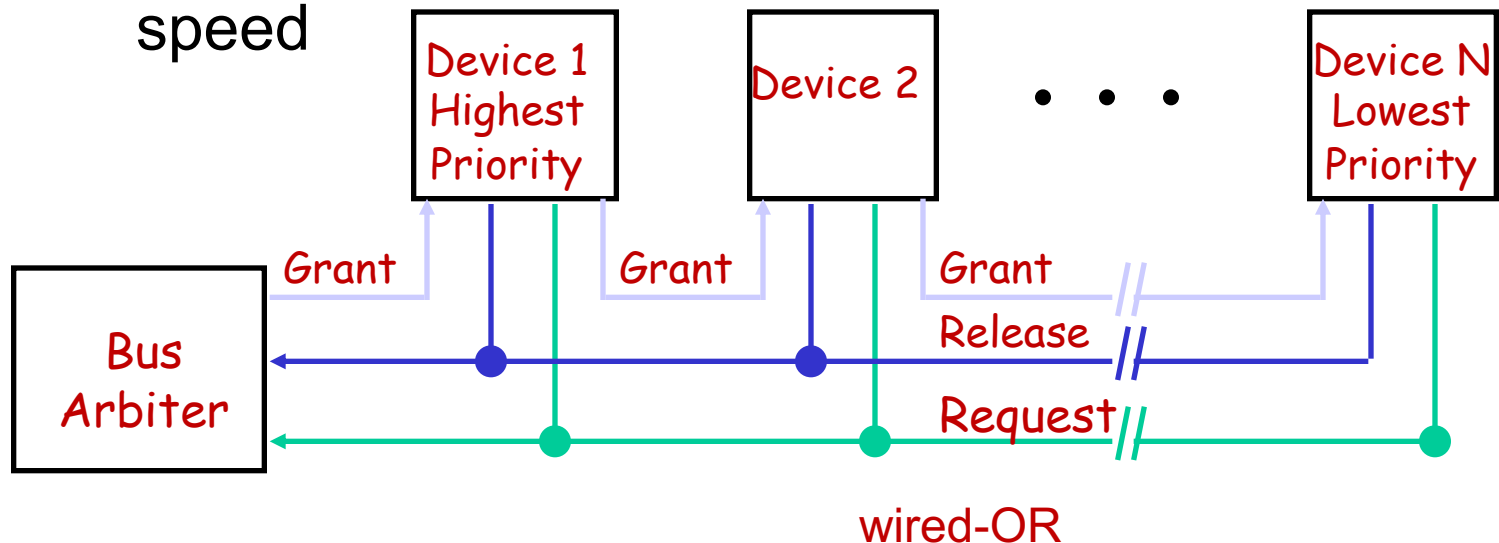


- Used in essentially all processor-memory busses and in high-speed I/O busses



Daisy Chain Bus Arbitrations Scheme

- Advantage: simple
- Disadvantages:
 - Cannot assure fairness:
 - A low-priority device may be locked out indefinitely
 - Use of daisy chain grant signal also limits bus speed



Increasing Bus Bandwidth

- Separate Address and Data Lines
- Data Bus Width
- Block Transfers



Increasing Bus Bandwidth

- Separate Address and Data Lines
 - vs. multiplexed address/data lines
 - Address and data can be transmitted in one bus cycle if separate address and data lines are available
 - Cost:
 - More bus lines
 - Increased complexity



Increasing Bus Bandwidth

- Data Bus Width
 - By increasing width of data bus, transfers of multiple words require fewer bus cycles
 - Example: SPARCstation 20's memory bus is 128 bit wide
 - Cost: more bus lines



Increasing Bus Bandwidth

- Block Transfers (Bursts Transfer)
 - Allow bus to transfer multiple words in back-to-back bus cycles
 - Only one address sent at beginning
 - Bus is not released until the last word is transferred
 - Cost:
 - Increased complexity
 - Decreased response time for request



I/O System Characteristics

- Dependability
 - Particularly for storage devices
- Performance Measures
 - Latency (response time)
 - Throughput (bandwidth)
 - Desktops & embedded systems
 - Mainly interested in response time & diversity of devices
 - Servers
 - Mainly interested in throughput & expandability of devices



I/O Device Examples

<u>Device</u>	<u>Behavior</u>	<u>Partner</u>	<u>Data Rate: KB/sec</u>
Keyboard	Input	Human	0.01
Mouse	Input	Human	0.02
Line Printer	Output	Human	1.00
Floppy disk	Storage	Machine	50.00
Laser Printer	Output	Human	100.00
Optical Disk	Storage	Machine	500.00
Magnetic Disk	Storage	Machine	5,000.00
Network-LAN	Input/Output	Machine	20 – 1,000.00
Graphics Display	Output	Human	30,000.00





Thanks for Your Attention!

